

# Amazon Fine Food Reviews Analysis Using RF and GBDT

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

EDA: <https://nycdatasience.com/blog/student-works/amazon-fine-foods-visualization/>  
(<https://nycdatasience.com/blog/student-works/amazon-fine-foods-visualization/>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

## Objective:

Given a review, determine whether the review is positive (Rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

## Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score id above 3, then the recommendation wil be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn import cross_validation
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.calibration import CalibratedClassifierCV
from sklearn.svm import SVC
from sklearn.linear_model import SGDClassifier
from sklearn.svm import LinearSVC
from sklearn.tree import DecisionTreeClassifier

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
```

```
D:\Anaconda3\lib\site-packages\gensim\utils.py:1209: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
  warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")
D:\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
  "This module will be removed in 0.20.", DeprecationWarning)
```

## [1]. Reading Data

```
In [2]: # using SQLite Table to read data.
con = sqlite3.connect('D:\\TGM\\ML\\AmazonFineFoodReviews\\database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 100000""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (100000, 10)

Out[2]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulne
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1

```
In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [4]: print(display.shape)
display.head()
```

```
(80668, 7)
```

```
Out[4]:
```

	UserId	ProductId	ProfileName	Time	Score	Text	Cou
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out[5]:
```

	UserId	ProductId	ProfileName	Time	Score	Text
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...

```
In [6]: display['COUNT(*)'].sum()
```

```
Out[6]: 393063
```

## Exploratory Data Analysis

### [2] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:



```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

Out[7]:

	<b>Id</b>	<b>ProductId</b>	<b>UserId</b>	<b>ProfileName</b>	<b>HelpfulnessNumerator</b>	<b>Helpful</b>
<b>0</b>	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2
<b>1</b>	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2
<b>2</b>	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2
<b>3</b>	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2
<b>4</b>	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2

As can be seen above the same user has multiple reviews of the with the same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

```
Out[9]: (87775, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 87.775
```

**Observation:-** It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[11]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulDenominator
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing Lets see the number of entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
print(final['Score'].value_counts())

(87773, 10)
1    73592
0    14181
Name: Score, dtype: int64
```

## [3]. Text Preprocessing.

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont take any chances till they know what is going on with the china imports.

=====

The Candy Blocks were a nice visual for the Lego Birthday party but the candy has little taste to it. Very little of the 2 lbs that I bought were eaten and I threw the rest away. I would not buy the candy again.

=====

was way to hot for my blood, took a bite and did a jig lol

=====

My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid of the fishy smell, don't get it. But I think my dog likes it because of the smell. These treats are really small in size. They are great for training. You can give your dog several of these without worrying about him over eating. Amazon's price was much more reasonable than any other retailer. You can buy a 1 pound bag on Amazon for almost the same price as a 6 ounce bag at other retailers. It's definitely worth it to buy a big bag if your dog eats them a lot.

=====

In [15]: *# https://stackoverflow.com/a/47091490/4084039*

```
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"'re", " are", phrase)
    phrase = re.sub(r"'s", " is", phrase)
    phrase = re.sub(r"'d", " would", phrase)
    phrase = re.sub(r"'ll", " will", phrase)
    phrase = re.sub(r"'t", " not", phrase)
    phrase = re.sub(r"'ve", " have", phrase)
    phrase = re.sub(r"'m", " am", phrase)
    return phrase
```

In [16]: `sent_4900 = decontracted(sent_4900)`  
`print(sent_4900)`  
`print("="*50)`

My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid of the fishy smell, do not get it. But I think my dog likes it because of the smell. These treats are really small in size. They are great for training. You can give your dog several of these without worrying about him over eating. Amazon is price was much more reasonable than any other retailer. You can buy a 1 pound bag on Amazon for almost the same price as a 6 ounce bag at other retailers. It is definitely worth it to buy a big bag if your dog eats them a lot.

=====

In [17]: *#remove words with numbers python: https://stackoverflow.com/a/18082370/4084039*

```
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its very hard to find any chicken products made in the USA but they are out there, but this one isnt. Its too bad too because its a good product but I wont take any chances till they know what is going on with the china imports.

In [18]: *#remove spacial character: https://stackoverflow.com/a/5843547/4084039*

```
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

was way to hot for my blood took a bite and did a jig lol



```
In [21]: preprocessed_reviews[1500]
```

```
Out[21]: 'way hot blood took bite jig lol'
```

```
In [22]: final['cleaned_text']=preprocessed_reviews
```

```
In [23]: final.shape
```

```
Out[23]: (87773, 11)
```

```
In [24]: final["Score"].value_counts()
```

```
Out[24]: 1    73592  
         0    14181  
         Name: Score, dtype: int64
```



```
In [25]: #Sorted the data based on time and took 100k data points
final["Time"] = pd.to_datetime(final["Time"], unit = "s")
final = final.sort_values(by = "Time")
final.head()
```

Out[25]:

	<b>Id</b>	<b>ProductId</b>	<b>UserId</b>	<b>ProfileName</b>	<b>HelpfulnessNumerator</b>	<b>He</b>
<b>70688</b>	76882	B00002N8SM	A32DW342WBJ6BX	Buttersugar	0	0
<b>1146</b>	1245	B00002Z754	A29Z5PI9BW2PU3	Robbie	7	7
<b>1145</b>	1244	B00002Z754	A3B8RCEI0FXFI6	B G Chase	10	10
<b>28086</b>	30629	B00008RCMI	A19E94CF5O1LY7	Andrew Arnold	0	0
<b>28087</b>	30630	B00008RCMI	A284C7M23F0APC	A. Mendoza	0	0

```
In [26]: Y = final['Score'].values
X = final['cleaned_text'].values
print(Y.shape)
print(type(Y))
print(X.shape)
print(type(X))
```

```
(87773,)
<class 'numpy.ndarray'>
(87773,)
<class 'numpy.ndarray'>
```

```
In [27]: # split the data set into train and test
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X,Y,test_size=0.3, random_
state=0)

# split the train data set into cross validation train and cross validation te
st
X_tr, X_cv, Y_tr, Y_cv = train_test_split(X,Y, test_size=0.3, random_state=0)

print('='*100)
print("After splitting")
print("X_Train Shape:",X_Train.shape, "Y_Train Shape:",Y_Train.shape)
print("X_cv Shape:",X_cv.shape, "Y_cv Shape:",Y_cv.shape)
print("X_Test Shape",X_Test.shape, "Y_Test Shape",Y_Test.shape)
```

```
=====
=====
After splitting
X_Train Shape: (61441,) Y_Train Shape: (61441,)
X_cv Shape: (26332,) Y_cv Shape (26332,)
X_Test Shape (26332,) Y_Test Shape (26332,)
```

## [4] Featurization

### <font color = Orange>[4.1] Applying RF on BoW Vectorizer</font>

```
In [28]: #Bow
count_vect = CountVectorizer(ngram_range=(1,2)) #in scikit-learn
count_vect.fit(X_Train)
print("some feature names ", count_vect.get_feature_names()[:10])
X_Train_Bow = count_vect.transform(X_Train)
X_Test_Bow = count_vect.transform(X_Test)
X_CV_Bow = count_vect.transform(X_cv)

print('='*50)

#final_counts = count_vect.transform(X_Test)

print("the type of X Train : ",type(X_Train_Bow))
print("the shape of Train BOW vectorizer ",X_Train_Bow.get_shape())
print("the shape of Test BOW vectorizer ",X_Test_Bow.get_shape())
print("the shape of CV BOW vectorizer ",X_CV_Bow.get_shape())
#print("the number of unique words ", final_counts.get_shape()[1])

some feature names  ['aa', 'aa caffene', 'aa coffee', 'aa cups', 'aa dark',
'aa extra', 'aa favorite', 'aa kona', 'aa may', 'aa not']
=====
the type of X Train :  <class 'scipy.sparse.csr.csr_matrix'>
the shape of Train BOW vectorizer  (61441, 1076376)
the shape of Test BOW vectorizer  (26332, 1076376)
the shape of CV BOW vectorizer  (26332, 1076376)
```

```
In [29]: import warnings
warnings.filterwarnings('ignore')
scalar = StandardScaler(with_mean=False)
X_Train_Bow = scalar.fit_transform(X_Train_Bow)
X_Test_Bow = scalar.transform(X_Test_Bow)
X_CV_Bow = scalar.transform(X_CV_Bow)

print("the type of X Train : ",type(X_Train_Bow))
print("the shape of Train BOW vectorizer ",X_Train_Bow.get_shape())
print("the shape of Test BOW vectorizer ",X_Test_Bow.get_shape())
print("the shape of CV BOW vectorizer ",X_CV_Bow.get_shape())

the type of X Train :  <class 'scipy.sparse.csr.csr_matrix'>
the shape of Train BOW vectorizer  (61441, 1076376)
the shape of Test BOW vectorizer  (26332, 1076376)
the shape of CV BOW vectorizer  (26332, 1076376)
```

```
In [42]: from sklearn.ensemble import RandomForestClassifier

Depth = [1,5,10]
BaseLearners = [10,20,40,60,80,100,120,150]

def Optimal_Values(X_Train,Y_Train):
    param_grid = {
        'max_depth':Depth,
        'n_estimators':BaseLearners
    }
    rf = RandomForestClassifier(class_weight = 'balanced', max_features='sqrt'
    )
    grid_search=GridSearchCV(rf,param_grid,scoring='roc_auc',n_jobs=-1,cv=3, verbose=2)
    grid_search.fit(X_Train,Y_Train)

    print("Best hyper paramters:",grid_search.best_params_)
    print("Best accuracy value: ",grid_search.best_score_ )
```

### <font color = blue>[4.1.1] Finding Best Hyper Parameters </font>

```
In [43]: Optimal_Values(X_Train_Bow, Y_Train)

Fitting 3 folds for each of 24 candidates, totalling 72 fits

[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 4.6min
[Parallel(n_jobs=-1)]: Done 72 out of 72 | elapsed: 14.3min finished


Best hyper paramters: {'max_depth': 10, 'n_estimators': 150}
Best accuracy value: 0.9021594429853547
```

### <font color = blue>[4.1.2] HeatMap for Train Data</font>

```


In [44]: X=[]
Y=[]
cv_data=[]
train_data=[]
for a in BaseLearners:
    for b in Depth:
        rf=RandomForestClassifier(class_weight = 'balanced', max_features='sqrt',max_depth=b,n_estimators=a)
        rf.fit(X_Train_Bow,Y_Train)
        cv_pred=rf.predict_proba(X_CV_Bow)[:,-1]
        train_pred=rf.predict_proba(X_Train_Bow)[:,-1]
        X.append(a)
        Y.append(b)
        cv_data.append(roc_auc_score(Y_cv,cv_pred))
        train_data.append(roc_auc_score(Y_Train,train_pred))
text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': train_data})
text_pivoted = text.pivot("max_depth","n_estimators", "AUC")
ax = sns.heatmap(text_pivoted,annot=True)
plt.title('Accuracy plot for Train Data')
plt.show()

```



### <font color = blue>[4.1.3] HeatMap for Cross Validation Data </font>

```
In [45]: text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': cv_data})
text_pivoted = text.pivot("max_depth", "n_estimators", "AUC")
ax = sns.heatmap(text_pivoted, annot=True)
plt.title('Accuracy plot for Cross Validation Data')
plt.show()
```




#### [4.1.4] ROC Curve of Random Forest

```
In [46]: rf = RandomForestClassifier(class_weight = 'balanced', max_depth= 5, max_features='sqrt', n_estimators=150)
rf.fit(X_Train_Bow,Y_Train)
prediction = rf.predict_proba(X_Test_Bow)[:,:1]
print(prediction)
print(rf)
```

```
Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, rf.predict_proba(X_Train_Bow)[:,:1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, rf.predict_proba(X_Test_Bow)[:,:1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)
```

```
plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random Forest')
plt.show()
```

```
[0.51309303 0.5155679 0.4974664 ... 0.50019621 0.49406763 0.48826417]
RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=5, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=150, n_jobs=1, oob_score=False, random_state=None,
                        verbose=0, warm_start=False)
```



### <font color = blue>[4.1.5]Train and Test Accuracy</font>

```
In [47]: Training_Accuracy_Bow = rf.score(X_Train_Bow, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Bow)
Training_Error_Bow = 1 - Training_Accuracy_Bow
print('Training_Error=%0.3f'%Training_Error_Bow)

Test_Accuracy_Bow = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Bow)
Test_Error_Bow = 1 - Test_Accuracy_Bow
print('Test_Error=%0.3f'%Test_Error_Bow)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_alpha_bow, Test_Accuracy_Bow))


Training_Accuracy=0.855
Training_Error=0.145
Test_Accuracy=0.845
Test_Error=0.155
```

### <font color = blue>[4.1.6] Confusion Matrix for Train Data </font>



```
In [48]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, rf.predict(X_Train_Bow))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=np.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, fmt='d')
```


Out[48]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fd96593550>



**[4.1.7] Confusion Matrix for Test Data**

```
In [49]: #With the reference of below link:
#https://www.kaggle.com/agungor2/various-confusion-matrix-plots
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, rf.predict(X_Test_Bow))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[49]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fd95eb2fd0>



## <font color = blue>[4.1.8] Classification Report</font>

```
In [50]: from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

	precision	recall	f1-score	support
0	0.51	0.76	0.61	4234
1	0.95	0.86	0.90	22098
avg / total	0.88	0.85	0.86	26332

## <font color = Blue>[4.1.9] Feature Importance with WordCloud </font>

```
In [51]: from wordcloud import WordCloud
Imp_features = count_vect.get_feature_names()
text = ''
feat=rf.feature_importances_
features=np.argsort(feat)[::-1][:30]
for i in features[0:20]:
    text += Imp_features[i]

plt.figure(figsize=(12,6))
wc = WordCloud(background_color='white',
               width=600,
               height=300,
               max_font_size=50,
               max_words=20).generate(text)
plt.imshow(wc, interpolation='bilinear')
plt.title("Imp Features", fontsize=20)
plt.axis("off")
plt.show()
```

Imp Features

recommendhighlystaleeasywastesnacktastykeepplasticcanceledhopingworkstastes

buying

highly

evenlunchlovenot

liketreatsticknot

## [4.2] Applying RF on TF-IDF Vectorizer

```
In [52]: #TF-IDF
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=5)
tf_idf_vect.fit_transform(X_Train)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names()[0:10])
print('='*50)

X_Train_TfIdf = tf_idf_vect.transform(X_Train)
X_Test_TfIdf = tf_idf_vect.transform(X_Test)
X_CV_TfIdf = tf_idf_vect.transform(X_cv)

#final_tf_idf = tf_idf_vect.transform(X_Test)
print("the type of count vectorizer ",type(X_Train_TfIdf))
print("the shape of out text TFIDF vectorizer ",X_Train_TfIdf.get_shape())
print("the shape of out text TFIDF vectorizer ",X_Test_TfIdf.get_shape())
print("the shape of out text TFIDF vectorizer ",X_CV_TfIdf.get_shape())
#print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[1])
```

```
some sample features(unique words in the corpus) ['aa', 'aaa', 'aafco', 'aback', 'abandon', 'abandoned', 'abdominal', 'abilities', 'ability', 'ability make']
```

```
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (61441, 80521)
the shape of out text TFIDF vectorizer (26332, 80521)
the shape of out text TFIDF vectorizer (26332, 80521)
```

```
In [53]: scalar = StandardScaler(with_mean=False)
X_Train_TfIdf = scalar.fit_transform(X_Train_TfIdf)
X_Test_TfIdf = scalar.transform(X_Test_TfIdf)
X_CV_TfIdf = scalar.transform(X_CV_TfIdf)

print("the type of count vectorizer ",type(X_Train_TfIdf))
print("the shape of out text TFIDF vectorizer ",X_Train_TfIdf.get_shape())
print("the shape of out text TFIDF vectorizer ",X_Test_TfIdf.get_shape())
print("the shape of out text TFIDF vectorizer ",X_CV_TfIdf.get_shape())
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (61441, 80521)
the shape of out text TFIDF vectorizer (26332, 80521)
the shape of out text TFIDF vectorizer (26332, 80521)
```

### <font color = blue>[4.2.1] Finding Best Hyper Parameters </font>

In [54]: `Optimal_Values(X_Train_Tfidf, Y_Train)`

Fitting 3 folds for each of 24 candidates, totalling 72 fits


```
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 72 out of 72 | elapsed: 2.6min finished
```

Best hyper paramters: {'max\_depth': 10, 'n\_estimators': 150}

Best accuracy value: 0.9129617321121051


## <font color = blue>[4.2.2] HeatMap for Train Data</font>

```
In [55]: X=[]
Y=[]
cv_data=[]
train_data=[]
for a in BaseLearners:
    for b in Depth:
        rf=RandomForestClassifier(class_weight = 'balanced', max_features='sqrt',max_depth=b,n_estimators=a)
        rf.fit(X_Train_Tfidf,Y_Train)
        cv_pred=rf.predict_proba(X_CV_Tfidf)[:,-1]
        train_pred=rf.predict_proba(X_Train_Tfidf)[:,-1]
        X.append(a)
        Y.append(b)
        cv_data.append(roc_auc_score(Y_cv,cv_pred))
        train_data.append(roc_auc_score(Y_Train,train_pred))
text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': train_data})
text_pivoted = text.pivot("max_depth","n_estimators", "AUC")
ax = sns.heatmap(text_pivoted,annot=True)
plt.title('Accuracy plot for Train Data')
plt.show()
```



## <font color = blue>[4.2.3] HeatMap for Cross Validation Data </font>

```
In [56]: text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': cv_data})
text_pivoted = text.pivot("max_depth", "n_estimators", "AUC")
ax = sns.heatmap(text_pivoted, annot=True)
plt.title('Accuracy plot for Cross Validation Data')
plt.show()
```



#### <font color = blue>[4.2.4] ROC Curve of Random Forest </font>

```

In [57]: #Testing with test data
rf = RandomForestClassifier(class_weight = 'balanced', max_depth= 5, max_features='sqrt', n_estimators=150)
rf.fit(X_Train_TfIdf,Y_Train)
prediction = rf.predict_proba(X_Test_TfIdf)[:,:1]
print(prediction)
print(rf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, rf.predict_proba(X_Train_TfIdf)[:,:1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, rf.predict_proba(X_Test_TfIdf)[:,:1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random Forest')
plt.show()


```

```
[0.51723181 0.49873784 0.49730541 ... 0.50323789 0.50620822 0.45049272]
```

```

RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=5, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=150, n_jobs=1, oob_score=False, random_state=None,
                        verbose=0, warm_start=False)

```



<font color = blue>[4.2.5]Train and Test Accuracy</font>

```
In [58]: Training_Accuracy_Tfidf = rf.score(X_Train_Tfidf, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Tfidf)
Training_Error_Tfidf = 1 - Training_Accuracy_Tfidf
print('Training_Error=%0.3f'%Training_Error_Tfidf)

Test_Accuracy_Tfidf = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Tfidf)
Test_Error_Tfidf = 1 - Test_Accuracy_Tfidf
print('Test_Error=%0.3f'%Test_Error_Tfidf)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_alpha_bow, Test_Accuracy_Bow))


Training_Accuracy=0.855
Training_Error=0.145
Test_Accuracy=0.846
Test_Error=0.154
```

## <font color = blue>[4.2.6] Confusion Matrix for Train Data </font>



```
In [59]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, rf.predict(X_Train_Tfidf))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=np.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, fmt='d')
```


Out[59]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fd8dd26fd0>



**[4.2.7] Confusion Matrix for Test Data**

```
In [60]: #With the reference of below link:
#https://www.kaggle.com/agungor2/various-confusion-matrix-plots
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, rf.predict(X_Test_Tfidf))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[60]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fd8e06f128>



## <font color = blue>[4.2.8] Classification Report</font>

```
In [61]: print(classification_report(Y_Test, prediction.round()))
```

	precision	recall	f1-score	support
0	0.52	0.72	0.60	4234
1	0.94	0.87	0.90	22098
avg / total	0.87	0.85	0.86	26332

## <font color = Blue>[4.2.9] Feature Importance with WordCloud </font>

```
In [62]: Imp_features = tf_idf_vect.get_feature_names()
text = ''
feat=rf.feature_importances_
features=np.argsort(feat)[::-1][:30]
for i in features[0:20]:
    text += Imp_features[i]
    text += " "
plt.figure(figsize=(12,6))
wc = WordCloud(background_color='white',
               width=600,
               height=300,
               max_font_size=50,
               max_words=20).generate(text)
plt.imshow(wc, interpolation='bilinear')
plt.title("Imp Features", fontsize=20)
plt.axis("off")
plt.show()
```



## <font color = Orange>[4.3]Word2Vec </font>

```
In [63]: i=0
list_of_sentence_train=[]
for sentence in X_Train:
    list_of_sentence_train.append(sentence.split())

w2v_model=Word2Vec(list_of_sentence_train,min_count=5,size=50, workers=4)
w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

number of words that occurred minimum 5 times 14786
sample words ['weekend', 'week', 'long', 'fast', 'using', 'rice', 'green',
'tea', 'works', 'wonders', 'one', 'energy', 'level', 'tasty', 'even', 'bit',
'salt', 'makes', 'much', 'pleasant', 'family', 'favorite', 'flavor', 'hans
n', 'diet', 'sodas', 'clean', 'crisp', 'taste', 'enjoyable', 'meals', 'calm
s', 'upset', 'tummy', 'love', 'compared', 'ones', 'used', 'eat', 'like', 'nis
sin', 'maruchan', 'really', 'tell', 'difference', 'big', 'tub', 'spice', 'dro
ps', 'better']
```

### <font color = blue>[4.3.1] Computing avg w2v for train</font>

```
In [64]: %%time
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in t
his list
for sent in list_of_sentence_train: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might
need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_train.append(sent_vec)
sent_vectors_train = np.array(sent_vectors_train)
print(sent_vectors_train.shape)
print(sent_vectors_train[0])

(61441, 50)
[ 0.65592798  0.18707622 -0.49871344 -0.40255735 -0.44068012  0.40905059
 -1.1956997  -0.49727267  0.54399973  0.0255063  -0.37244177 -0.11311698
 -0.98367127 -0.15229551 -0.11185685  0.01850094  0.16695645 -0.19574464
  0.6338068  0.56661695  0.49751364 -0.02130096 -0.59853306  0.08581292
  0.08077065 -0.42216597  0.12347004  0.17604289 -0.77206999  0.83425402
 -0.32205746  0.27287036 -0.33797387 -0.21264058 -0.36737667 -0.25136496
  0.1382337  0.12414662 -0.3995532  0.74111378 -0.95913343  0.0579938
 -0.35465602  1.02291975 -0.21281362  0.3895559  -0.17693804 -0.15483701
 -0.33596356  1.33765634]
Wall time: 2min 28s
```

**<font color = blue>[4.3.2] Computing avg w2v for CV</font>**

```

In [65]: %%time
i=0
list_of_sentence_cv=[]
for sentence in X_cv:
    list_of_sentence_cv.append(sentence.split())
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sentence_cv: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might
    need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_cv.append(sent_vec)
sent_vectors_cv = np.array(sent_vectors_cv)
print(sent_vectors_cv.shape)
print(sent_vectors_cv[0])

```

```

(26332, 50)
[ 0.91863407  0.32927822 -1.15930034 -0.97737705 -0.29546582  0.79953468
  0.04454586  1.23157543 -0.46114961 -0.17512483  0.59920082  0.12326036
 -0.14115333  0.12126782 -0.46732131  0.17765068  0.62152439 -0.63263151
 -0.0381849  -0.53950762  1.24986861 -0.99977002  1.05252012  0.73728212
 -0.63139761 -1.4498087   0.36142013 -0.8750116  -0.62328856  0.31623498
 -0.49659881  0.9203474  -0.52036034 -0.05922971 -0.28708383 -0.64734941
 -0.46971358  0.45888514  0.22286932  0.88516332 -0.66994867  0.27606364
 -0.29864415  0.21530222  1.02035046 -0.23230878 -0.47036943 -0.02433476
 -0.36756549  1.56408059]
Wall time: 1min 12s

```

**<font color = blue>[4.3.3] Computing avg w2v for test</font>**

```
In [66]: %%time

i=0
list_of_sentence_test=[]
for sentence in X_Test:
    list_of_sentence_test.append(sentence.split())

# average Word2Vec
# compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sentence_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_test.append(sent_vec)
sent_vectors_test = np.array(sent_vectors_test)
print(sent_vectors_test.shape)
print(sent_vectors_test[0])
```

```
(26332, 50)
[ 0.91863407  0.32927822 -1.15930034 -0.97737705 -0.29546582  0.79953468
  0.04454586  1.23157543 -0.46114961 -0.17512483  0.59920082  0.12326036
 -0.14115333  0.12126782 -0.46732131  0.17765068  0.62152439 -0.63263151
 -0.0381849  -0.53950762  1.24986861 -0.99977002  1.05252012  0.73728212
 -0.63139761 -1.4498087   0.36142013 -0.8750116  -0.62328856  0.31623498
 -0.49659881  0.9203474  -0.52036034 -0.05922971 -0.28708383 -0.64734941
 -0.46971358  0.45888514  0.22286932  0.88516332 -0.66994867  0.27606364
 -0.29864415  0.21530222  1.02035046 -0.23230878 -0.47036943 -0.02433476
 -0.36756549  1.56408059]
Wall time: 1min 25s
```

#### <font color = blue>[4.3.4] Finding Best Hyper Parameters </font>

```
In [67]: Optimal_Values(sent_vectors_train, Y_Train)

Fitting 3 folds for each of 24 candidates, totalling 72 fits

[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 72 out of 72 | elapsed: 7.8min finished


Best hyper paramters: {'max_depth': 10, 'n_estimators': 150}
Best accuracy value: 0.8882572868608006
```

#### <font color = blue>[4.3.5] HeatMap for Train Data</font>

```


In [68]: X=[]
Y=[]
cv_data=[]
train_data=[]
for a in BaseLearners:
    for b in Depth:
        rf=RandomForestClassifier(class_weight = 'balanced', max_features='sqrt',max_depth=b,n_estimators=a)
        rf.fit(sent_vectors_train,Y_Train)
        cv_pred=rf.predict_proba(sent_vectors_cv)[:,:1]
        train_pred=rf.predict_proba(sent_vectors_train)[:,:1]
        X.append(a)
        Y.append(b)
        cv_data.append(roc_auc_score(Y_cv,cv_pred))
        train_data.append(roc_auc_score(Y_Train,train_pred))
text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': train_data})
text_pivoted = text.pivot("max_depth","n_estimators", "AUC")
ax = sns.heatmap(text_pivoted,annot=True)
plt.title('Accuracy plot for Train Data')
plt.show()

```



### <font color = blue>[4.3.6] HeatMap for Cross Validation Data </font>

```
In [69]: text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': cv_data})
text_pivoted = text.pivot("max_depth", "n_estimators", "AUC")
ax = sns.heatmap(text_pivoted, annot=True)
plt.title('Accuracy plot for Cross Validation Data')
plt.show()
```



### <font color = blue>[4.3.7] ROC Curve of Random Forest </font>




```
In [70]: #Testing with test data
rf = RandomForestClassifier(class_weight = 'balanced', max_depth= 5, max_features='sqrt', n_estimators=150)
rf.fit(sent_vectors_train,Y_Train)
prediction = rf.predict_proba(sent_vectors_test)[:,:1]
print(prediction)
print(rf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, rf.predict_proba(sent_vectors_train)[:,:1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, rf.predict_proba(sent_vectors_test)[:,:1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random Forest')
plt.show()
```

```
[0.71710563 0.61763942 0.38896586 ... 0.49951279 0.58292034 0.17845168]
RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=5, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=150, n_jobs=1, oob_score=False, random_state=None,
                        verbose=0, warm_start=False)
```



**<font color = blue>[4.3.8]Train and Test Accuracy</font>**

```
In [71]: Training_Accuracy_w2v = rf.score(sent_vectors_train, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_w2v)
Training_Error_w2v = 1 - Training_Accuracy_w2v
print('Training_Error=%0.3f'%Training_Error_w2v)


Test_Accuracy_w2v = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_w2v)
Test_Error_w2v = 1 - Test_Accuracy_w2v
print('Test_Error=%0.3f'%Test_Error_w2v)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_alpha_bow, Test_Accuracy_Bow))

Training_Accuracy=0.784
Training_Error=0.216
Test_Accuracy=0.775
Test_Error=0.225
```

### <font color = blue>[4.3.9] Confusion Matrix for Train Data </font>

```
In [72]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, rf.predict(sent_vectors_train))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=np.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, fmt='d')
```


Out[72]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fd901aa828>



**[4.3.10] Confusion Matrix for Test Data**

```
In [73]: #With the reference of below link:
#https://www.kaggle.com/agungor2/various-confusion-matrix-plots
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, rf.predict(sent_vectors_test))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[73]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdc3117eb8>



<font color = blue>[4.3.11] Classification Report</font>

```
In [74]: print(classification_report(Y_Test, prediction.round()))
```

	precision	recall	f1-score	support
0	0.40	0.79	0.53	4234
1	0.95	0.77	0.85	22098
avg / total	0.86	0.78	0.80	26332

## <font color = Orange>[4.4] Applying RF on TFIDF weighted W2v Vectorizer </font>

```
In [75]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
model.fit(X_Train)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

### <font color = blue>[4.4.1] Compute TF-IDF weighted Word2Vec for Train</font>

```

In [76]: i=0
list_of_sentence_train=[]
for sentence in X_Train:
    list_of_sentence_train.append(sentence.split())
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val
= tfidf

tfidf_sent_vectors_train = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sentence_train: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#
            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_train.append(sent_vec)
    row += 1

```

#### <font color = blue>[4.4.2] Compute TF-IDF weighted Word2Vec for Test</font>

```

In [77]: i=0
list_of_sentence_test=[]
for sentence in X_Test:
    list_of_sentence_test.append(sentence.split())
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val
= tfidf

tfidf_sent_vectors_test = []; # the tfidf-w2v for each sentence/review is stor
ed in this list
row=0;
for sent in list_of_sentence_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#
            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_test.append(sent_vec)
    row += 1

```

### <font color = blue>[4.4.3] Compute TF-IDF weighted Word2Vec for CV</font>

```

In [78]: i=0
list_of_sentence_cv=[]
for sentence in X_cv:
    list_of_sentence_cv.append(sentence.split())
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val
= tfidf

tfidf_sent_vectors_cv = []; # the tfidf-w2v for each sentence/review is stored
in this list
row=0;
for sent in list_of_sentence_cv: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#
            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_cv.append(sent_vec)
    row += 1

```

#### <font color = blue>[4.4.4] Finding Best Hyper Parameters </font>

```

In [79]: Optimal_Values(tfidf_sent_vectors_train, Y_Train)

```

Fitting 3 folds for each of 24 candidates, totalling 72 fits

```
[Parallel(n_jobs=-1)]: Done 33 tasks | elapsed: 1.8min
```

```
[Parallel(n_jobs=-1)]: Done 72 out of 72 | elapsed: 8.0min finished
```

Best hyper paramters: {'max\_depth': 10, 'n\_estimators': 150}

Best accuracy value: 0.8663797891047916


#### <font color = blue>[4.4.5] HeatMap for Train Data</font>



```


In [80]: X=[]
Y=[]
cv_data=[]
train_data=[]
for a in [10,20,40,60,80,100,120,150]:
    for b in [1,5]:
        rf=RandomForestClassifier(class_weight = 'balanced', max_features='sqrt',max_depth=b,n_estimators=a)
        rf.fit(tfidf_sent_vectors_train,Y_Train)
        cv_pred=rf.predict_proba(tfidf_sent_vectors_cv)[: ,1]
        train_pred=rf.predict_proba(tfidf_sent_vectors_train)[: ,1]
        X.append(a)
        Y.append(b)
        cv_data.append(roc_auc_score(Y_cv,cv_pred))
        train_data.append(roc_auc_score(Y_Train,train_pred))
text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': train_data})
text_pivoted = text.pivot("max_depth","n_estimators", "AUC")
ax = sns.heatmap(text_pivoted,annot=True)
plt.title('Accuracy plot for Train Data')
plt.show()

```



#### <font color = blue>[4.4.6] HeatMap for Cross Validation Data </font>

```
In [81]: text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': cv_data})
text_pivoted = text.pivot("max_depth", "n_estimators", "AUC")
ax = sns.heatmap(text_pivoted, annot=True)
plt.title('Accuracy plot for Cross Validation Data')
plt.show()
```




**[4.4.7] ROC Curve of Random Forest**

```
In [82]: #Testing with test data
rf = RandomForestClassifier(class_weight = 'balanced', max_depth= 5, max_features='sqrt', n_estimators=150)
rf.fit(tfidf_sent_vectors_train,Y_Train)
prediction = rf.predict_proba(tfidf_sent_vectors_test)[: ,1]
print(prediction)
print(rf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, rf.predict_proba(tfidf_sent_vectors_train)[: ,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, rf.predict_proba(tfidf_sent_vectors_test)[: ,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random Forest')
plt.show()
```

```
[0.646715  0.60636767 0.35076168 ... 0.48921192 0.52550631 0.20799956]
RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=5, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=150, n_jobs=1, oob_score=False, random_state=None,
                        verbose=0, warm_start=False)
```



**<font color = blue>[4.4.8]Train and Test Accuracy</font>**

```
In [83]: Training_Accuracy_tfidf2v = rf.score(tfidf_sent_vectors_train, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_tfidf2v)
Training_Error_tfidf2v = 1 - Training_Accuracy_tfidf2v
print('Training_Error=%0.3f'%Training_Error_tfidf2v)


Test_Accuracy_tfidf2v = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_tfidf2v)
Test_Error_tfidf2v = 1 - Test_Accuracy_tfidf2v
print('Test_Error=%0.3f'%Test_Error_tfidf2v)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_alpha_bow, Test_Accuracy_Bow))

Training_Accuracy=0.775
Training_Error=0.225
Test_Accuracy=0.769
Test_Error=0.231
```

#### <font color = blue>[4.4.9] Confusion Matrix for Train Data</font>

```
In [84]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, rf.predict(tfidf_sent_vectors_train))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=np.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, fmt='d')
```

Out[84]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdc3185978>




**[4.4.10] Confusion Matrix for Test Data**

```

In [85]: #With the reference of below link:
#https://www.kaggle.com/agungor2/various-confusion-matrix-plots
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, rf.predict(tfidf_sent_vectors_test))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')

```

Out[85]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdc54b59b0>



**<font color = blue>[4.4.11] Classification Report</font>**

```
In [86]: print(classification_report(Y_Test, prediction.round()))
```

	precision	recall	f1-score	support
0	0.39	0.74	0.51	4234
1	0.94	0.77	0.85	22098
avg / total	0.85	0.77	0.79	26332

## <font color = Blue>[4.4.12] Feature Importance with WordCloud</font>

```
In [87]: Imp_features = model.get_feature_names()
text = ''
feat=rf.feature_importances_
features=np.argsort(feat)[::-1][:30]
for i in features[0:20]:
    text += Imp_features[i]

plt.figure(figsize=(12,6))
wc = WordCloud(background_color='white',
               width=600,
               height=300,
               max_font_size=50,
               max_words=20).generate(text)
plt.imshow(wc, interpolation='bilinear')
plt.title("Imp Features", fontsize=20)
plt.axis("off")
plt.show()
```

### Imp Features

aaaaahabandonabcessedabdaashabbeyaaaaaaHHHHhhabidingabcaberrationabbyaahaaaandabideabettaraafcoamazonaahsaadpabibreviation

## <font color = Orange>[5] Applying GBDT Using XGBoost</font>

### <font color = Orange>[5.1] Applying GBDT using XGBoost on BoW Vectorizer</font>

#### <font color = blue>[5.1.1] Finding Best Hyper Parameters</font>

```
In [89]: import xgboost as xgb
        from xgboost.sklearn import XGBClassifier

        Depth = [1,5]
        BaseLearners = [10,20,40,60,80,100,120,150]

        param_grid = {
            'max_depth':Depth,
            'n_estimators':BaseLearners
        }
        xgb = XGBClassifier(class_weight = 'balanced', booster='gbtree')
        grid_search=GridSearchCV(xgb,param_grid,scoring='roc_auc',n_jobs=-1,cv=3, verbose=2)
        grid_search.fit(X_Train_Bow,Y_Train)
        print("Best hyper paramters:",grid_search.best_params_)
        print("Best accuracy value: ",grid_search.best_score_ )
```

Fitting 3 folds for each of 16 candidates, totalling 48 fits

[Parallel(n\_jobs=-1)]: Done 33 tasks | elapsed: 29.4min

[Parallel(n\_jobs=-1)]: Done 48 out of 48 | elapsed: 100.0min finished

Best hyper paramters: {'max\_depth': 5, 'n\_estimators': 150}

Best accuracy value: 0.9214042036468455


#### <font color = blue>[5.1.2] HeatMap for Train Data</font>



```


In [90]: X=[]
Y=[]
cv_data=[]
train_data=[]
for a in BaseLearners:
    for b in Depth:
        xgb=XGBClassifier(class_weight = 'balanced', booster='gbtree', max_depth=b,n_estimators=a)
        xgb.fit(X_Train_Bow,Y_Train)
        cv_pred=xgb.predict_proba(X_CV_Bow)[: ,1]
        train_pred=xgb.predict_proba(X_Train_Bow)[: ,1]
        X.append(a)
        Y.append(b)
        cv_data.append(roc_auc_score(Y_cv,cv_pred))
        train_data.append(roc_auc_score(Y_Train,train_pred))
text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': train_data})
text_pivoted = text.pivot("max_depth","n_estimators", "AUC")
ax = sns.heatmap(text_pivoted,annot=True)
plt.title('Accuracy plot for Train Data')
plt.show()

```



### <font color = blue>[5.1.3] HeatMap for Cross Validation Data </font>

```
In [91]: text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': cv_data})
text_pivoted = text.pivot("max_depth", "n_estimators", "AUC")
ax = sns.heatmap(text_pivoted, annot=True)
plt.title('Accuracy plot for Cross Validation Data')
plt.show()
```




#### <font color = blue>[5.1.4] ROC Curve of Random Forest </font>

```
In [92]: #Testing with test data
xgb =XGBClassifier(class_weight = 'balanced', booster = 'gbtree', max_depth =
5, n_estimators = 150)
xgb.fit(X_Train_Bow,Y_Train)
prediction = xgb.predict_proba(X_Test_Bow)[:,:1]
print(prediction)
print(xgb)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, xgb.predict_proba(X_Train
_Bow)[:,:1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, xgb.predict_proba(X_Test_Bow
)[:,:1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random Forest')
plt.show()
```

```
[0.9463006 0.89967996 0.72777814 ... 0.83416975 0.7347367 0.29207373]
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
  colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
  gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=5,
  min_child_weight=1, missing=None, n_estimators=150, n_jobs=1,
  nthread=None, objective='binary:logistic', random_state=0,
  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
  silent=None, subsample=1, verbosity=1)
```



### <font color = blue>[5.1.6]Train and Test Accuracy</font>

```
In [93]: Training_Accuracy_Bow = xgb.score(X_Train_Bow, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Bow)
Training_Error_Bow = 1 - Training_Accuracy_Bow
print('Training_Error=%0.3f'%Training_Error_Bow)


Test_Accuracy_Bow = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Bow)
Test_Error_Bow = 1 - Test_Accuracy_Bow
print('Test_Error=%0.3f'%Test_Error_Bow)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))

Training_Accuracy=0.906
Training_Error=0.094
Test_Accuracy=0.898
Test_Error=0.102
```

### <font color = blue>[5.1.7] Confusion Matrix for Train Data </font>

```
In [94]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, xgb.predict(X_Train_Bow))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=np.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, fmt='d')
```

Out[94]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdd6414898>




**<font color = blue>[5.1.8] Confusion Matrix for Test Data </font>**

```

In [95]: #With the reference of below link:
#https://www.kaggle.com/agungor2/various-confusion-matrix-plots
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, xgb.predict(X_Test_Bow))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')

```

Out[95]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdd6414710>



### <font color = blue>[5.1.9] Classification Report</font>

```
In [96]: print(classification_report(Y_Test, prediction.round()))
```

	precision	recall	f1-score	support
0	0.89	0.41	0.57	4234
1	0.90	0.99	0.94	22098
avg / total	0.90	0.90	0.88	26332

## <font color = Blue>[5.1.10] Feature Importance with WordCloud </font>

```
In [97]: from wordcloud import WordCloud
Imp_features = count_vect.get_feature_names()
text = ''
feat=xgb.feature_importances_
features=np.argsort(feat)[::-1]
for i in features[0:20]:
    text += Imp_features[i]

plt.figure(figsize=(12,6))
wc = WordCloud(background_color='white',
               width=600,
               height=300,
               max_font_size=50,
               max_words=20).generate(text)
plt.imshow(wc, interpolation='bilinear')
plt.title("Imp Features", fontsize=20)
plt.axis("off")
plt.show()
```

Imp Features



## <font color = Orange>[5.2] Apply GBDT using XGBoost on TFIDF Vectorizer </font>

### <font color = blue>[5.2.1] Finding Best Hyper Parameters </font>

```
In [98]: param_grid = {
          'max_depth':Depth,
          'n_estimators':BaseLearners
        }
xgb = XGBClassifier(class_weight = 'balanced', booster='gbtree')
grid_search=GridSearchCV(xgb,param_grid,scoring='roc_auc',n_jobs=-1,cv=3, verbose=2)
grid_search.fit(X_Train_TfIdf,Y_Train)
print("Best hyper paramters:",grid_search.best_params_)
print("Best accuracy value: ",grid_search.best_score_ )
```

Fitting 3 folds for each of 16 candidates, totalling 48 fits

[Parallel(n\_jobs=-1)]: Done 33 tasks | elapsed: 9.8min

[Parallel(n\_jobs=-1)]: Done 48 out of 48 | elapsed: 766.7min finished

Best hyper paramters: {'max\_depth': 5, 'n\_estimators': 150}

Best accuracy value: 0.9214538753224479

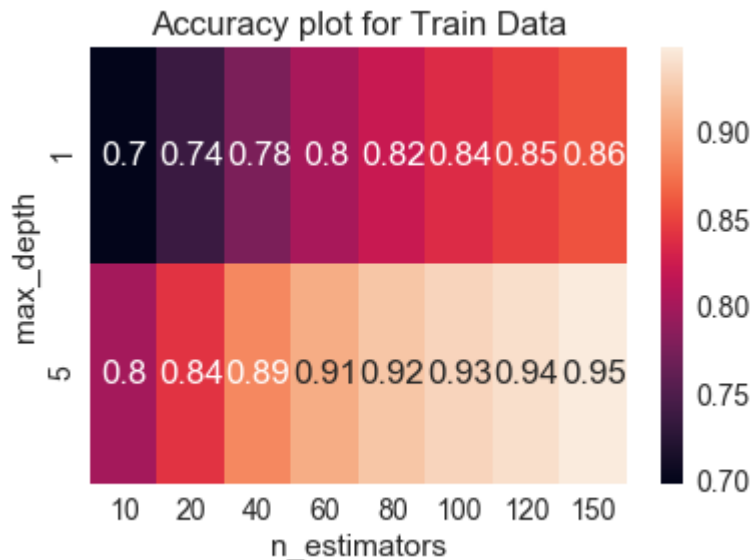
### <font color = blue>[5.2.2] HeatMap for Train Data</font>



```


In [99]: X=[]
Y=[]
cv_data=[]
train_data=[]
for a in BaseLearners:
    for b in Depth:
        xgb=XGBClassifier(class_weight = 'balanced', booster='gbtree', max_depth=b,n_estimators=a)
        xgb.fit(X_Train_TfIdf,Y_Train)
        cv_pred=xgb.predict_proba(X_CV_TfIdf)[: ,1]
        train_pred=xgb.predict_proba(X_Train_TfIdf)[: ,1]
        X.append(a)
        Y.append(b)
        cv_data.append(roc_auc_score(Y_cv,cv_pred))
        train_data.append(roc_auc_score(Y_Train,train_pred))
text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': train_data})
text_pivoted = text.pivot("max_depth","n_estimators", "AUC")
ax = sns.heatmap(text_pivoted,annot=True)
plt.title('Accuracy plot for Train Data')
plt.show()

```



### <font color = blue>[5.2.3] HeatMap for Cross Validation Data </font>

```
In [100]: text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': cv_data})
text_pivoted = text.pivot("max_depth", "n_estimators", "AUC")
ax = sns.heatmap(text_pivoted, annot=True)
plt.title('Accuracy plot for Cross Validation Data')
plt.show()
```




**[5.2.5] ROC Curve of Random Forest**

```
In [101]: #Testing with test data
xgb =XGBClassifier(class_weight = 'balanced', booster = 'gbtree', max_depth =
5, n_estimators = 150)
xgb.fit(X_Train_TfIdf,Y_Train)
prediction = xgb.predict_proba(X_Test_TfIdf)[:,:1]
print(prediction)
print(xgb)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, xgb.predict_proba(X_Train
_TfIdf)[:,:1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, xgb.predict_proba(X_Test_TfI
df)[:,:1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random Forest')
plt.show()
```

```
[0.961208  0.90644634 0.66455096 ... 0.8303739  0.6179137  0.3207916 ]
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
  colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
  gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=5,
  min_child_weight=1, missing=None, n_estimators=150, n_jobs=1,
  nthread=None, objective='binary:logistic', random_state=0,
  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
  silent=None, subsample=1, verbosity=1)
```



## <font color = blue>[5.2.6]Train and Test Accuracy</font>

```
In [103]: Training_Accuracy_Tfidf = xgb.score(X_Train_Tfidf, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Tfidf)
Training_Error_Tfidf = 1 - Training_Accuracy_Tfidf
print('Training_Error=%0.3f'%Training_Error_Tfidf)


Test_Accuracy_Tfidf = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Tfidf)
Test_Error_Tfidf = 1 - Test_Accuracy_Tfidf
print('Test_Error=%0.3f'%Test_Error_Tfidf)

Training_Accuracy=0.910
Training_Error=0.090
Test_Accuracy=0.900
Test_Error=0.100
```

### <font color = blue>[5.2.7] Confusion Matrix for Train Data </font>

```
In [104]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, xgb.predict(X_Train_Tfidf))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=np.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':12}, fmt='d')
```


Out[104]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fddd5a35c0>



**<font color = blue>[5.2.8] Confusion Matrix for Test Data </font>**

```
In [105]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, xgb.predict(X_Test_Tfidf))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':12}, f
mt='d')
```

Out[105]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdc5a15e80>



## <font color = blue>[5.2.9] Classification Report</font>

```
In [106]: print(classification_report(Y_Test, prediction.round()))
```

	precision	recall	f1-score	support
0	0.89	0.43	0.58	4234
1	0.90	0.99	0.94	22098
avg / total	0.90	0.90	0.88	26332

## <font color = Blue>[5.2.10] Feature Importance with WordCloud </font>

```
In [107]: Imp_features = tf_idf_vect.get_feature_names()
text = ''
feat=xgb.feature_importances_
features=np.argsort(feat)[::-1][:30]
for i in features[0:20]:
    text += Imp_features[i]

plt.figure(figsize=(12,6))
wc = WordCloud(background_color='white',
               width=600,
               height=300,
               max_font_size=50,
               max_words=20).generate(text)
plt.imshow(wc, interpolation='bilinear')
plt.title("Imp Features", fontsize=20)
plt.axis("off")
plt.show()
```

Imp Features



returnworsthorriblenot  
buyterriblegreatdisappointingbesthighlynot  
disappointeddisgustingdisappointeddeliciousbadlovesawfulmoneyeasyperfectrefund

## <font color = Orange>[5.3] Applying GBDT using XGBoost on Word2Vec Vectorizer </font>

### <font color = blue>[5.3.1] Finding Best Hyper Parameters </font>

```
In [108]: param_grid = {
            'max_depth':Depth,
            'n_estimators':BaseLearners
          }
xgb = XGBClassifier(class_weight = 'balanced', booster='gbtree')
grid_search=GridSearchCV(xgb,param_grid,scoring='roc_auc',n_jobs=-1,cv=3, verbose=2)
grid_search.fit(sent_vectors_train,Y_Train)
print("Best hyper paramters:",grid_search.best_params_)
print("Best accuracy value: ",grid_search.best_score_ )
```

Fitting 3 folds for each of 16 candidates, totalling 48 fits

```
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 2.8min
[Parallel(n_jobs=-1)]: Done 48 out of 48 | elapsed: 8.5min finished
```

```
Best hyper paramters: {'max_depth': 5, 'n_estimators': 150}
Best accuracy value: 0.9073752621051433
```

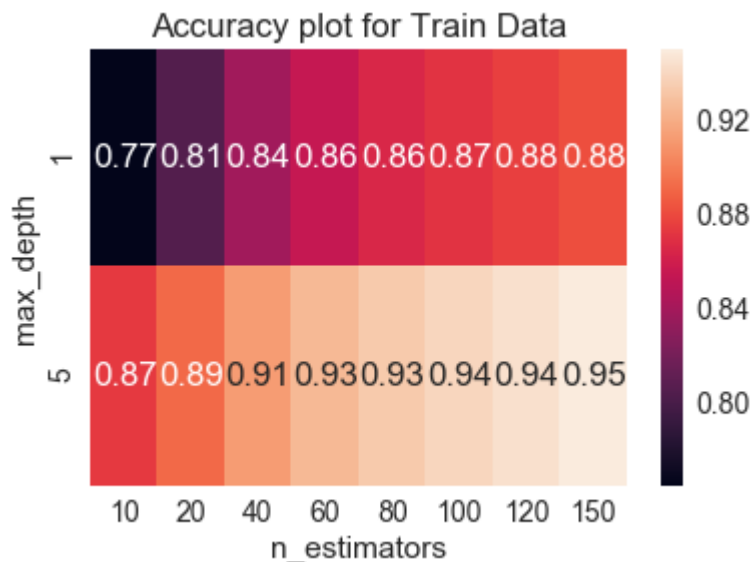
### <font color = blue>[5.3.2] HeatMap for Train Data</font>



```


In [109]: X=[]
Y=[]
cv_data=[]
train_data=[]
for a in BaseLearners:
    for b in Depth:
        xgb=XGBClassifier(class_weight = 'balanced', booster='gbtree', max_depth=b,n_estimators=a)
        xgb.fit(sent_vectors_train,Y_Train)
        cv_pred=xgb.predict_proba(sent_vectors_cv)[: ,1]
        train_pred=xgb.predict_proba(sent_vectors_train)[: ,1]
        X.append(a)
        Y.append(b)
        cv_data.append(roc_auc_score(Y_cv,cv_pred))
        train_data.append(roc_auc_score(Y_Train,train_pred))
text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': train_data})
text_pivoted = text.pivot("max_depth","n_estimators", "AUC")
ax = sns.heatmap(text_pivoted,annot=True)
plt.title('Accuracy plot for Train Data')
plt.show()

```



### <font color = blue>[5.3.2] HeatMap for Cross Validation Data </font>

```
In [110]: text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': cv_data})
text_pivoted = text.pivot("max_depth", "n_estimators", "AUC")
ax = sns.heatmap(text_pivoted, annot=True)
plt.title('Accuracy plot for Cross Validation Data')
plt.show()
```




### <font color = blue>[5.3.3] ROC Curve of Random Forest </font>

```
In [111]: xgb =XGBClassifier(class_weight = 'balanced', booster = 'gbtree', max_depth =
5, n_estimators = 150)
xgb.fit(sent_vectors_train,Y_Train)
prediction = xgb.predict_proba(sent_vectors_test)[:,:1]
print(prediction)
print(xgb)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, xgb.predict_proba(sent_vectors_train)[:,:1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, xgb.predict_proba(sent_vectors_test)[:,:1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random Forest')
plt.show()
```

```
[0.9613765 0.9606293 0.27951214 ... 0.72093636 0.90785134 0.19023645]
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
  colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
  gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=5,
  min_child_weight=1, missing=None, n_estimators=150, n_jobs=1,
  nthread=None, objective='binary:logistic', random_state=0,
  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
  silent=None, subsample=1, verbosity=1)
```



### <font color = blue>[5.3.4]Train and Test Accuracy</font>

```
In [113]: Training_Accuracy_w2v = xgb.score(sent_vectors_train, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_w2v)
Training_Error_w2v = 1 - Training_Accuracy_w2v
print('Training_Error=%0.3f'%Training_Error_w2v)


Test_Accuracy_w2v = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_w2v)
Test_Error_w2v = 1 - Test_Accuracy_w2v
print('Test_Error=%0.3f'%Test_Error_w2v)

Training_Accuracy=0.912
Training_Error=0.088
Test_Accuracy=0.890
Test_Error=0.110
```

### <font color = blue>[5.3.5] Confusion Matrix for Train Data </font>

```
In [114]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, xgb.predict(sent_vectors_train))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=np.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':12}, fmt='d')
```


Out[114]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdd90cdb70>



**<font color = blue>[5.3.6] Confusion Matrix for Test Data </font>**

```
In [115]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, xgb.predict(sent_vectors_test))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':12}, f
mt='d')
```

Out[115]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdb3f6710>



### <font color = blue>[5.3.7] Classification Report</font>

```
In [116]: print(classification_report(Y_Test, prediction.round()))
```

	precision	recall	f1-score	support
0	0.74	0.49	0.59	4234
1	0.91	0.97	0.94	22098
avg / total	0.88	0.89	0.88	26332

## <font color = Orange>[5.4] Applying GBDT using XGBoost on TFIDF weighted W2V Vectorizer </font>

### <font color = blue>[5.4.1] Finding Best Hyper Parameters </font>

```
In [117]: tfidf_sent_vectors_train=np.array(tfidf_sent_vectors_train)
          tfidf_sent_vectors_cv=np.array(tfidf_sent_vectors_cv)
          tfidf_sent_vectors_test=np.array(tfidf_sent_vectors_test)

          param_grid = {
                        'max_depth':Depth,
                        'n_estimators':BaseLearners
                      }
          xgb = XGBClassifier(class_weight = 'balanced', booster='gbtree')
          grid_search=GridSearchCV(xgb,param_grid,scoring='roc_auc',n_jobs=-1,cv=3, verbose=2)
          grid_search.fit(tfidf_sent_vectors_train,Y_Train)
          print("Best hyper paramters:",grid_search.best_params_)
          print("Best accuracy value: ",grid_search.best_score_ )
```

Fitting 3 folds for each of 16 candidates, totalling 48 fits

[Parallel(n\_jobs=-1)]: Done 33 tasks | elapsed: 3.6min

[Parallel(n\_jobs=-1)]: Done 48 out of 48 | elapsed: 9.3min finished

Best hyper paramters: {'max\_depth': 5, 'n\_estimators': 150}


Best accuracy value: 0.8859234430056635

### <font color = blue>[5.4.2] HeatMap for Train Data</font>

```

In [118]: X=[]
Y=[]
cv_data=[]
train_data=[]
for a in BaseLearners:
    for b in Depth:
        xgb=XGBClassifier(class_weight = 'balanced', booster='gbtree', max_depth=b,n_estimators=a)
        xgb.fit(tfidf_sent_vectors_train,Y_Train)
        cv_pred=xgb.predict_proba(tfidf_sent_vectors_cv)[: ,1]
        train_pred=xgb.predict_proba(tfidf_sent_vectors_train)[: ,1]
        X.append(a)
        Y.append(b)
        cv_data.append(roc_auc_score(Y_cv,cv_pred))
        train_data.append(roc_auc_score(Y_Train,train_pred))
text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': train_data})
text_pivoted = text.pivot("max_depth","n_estimators", "AUC")
ax = sns.heatmap(text_pivoted,annot=True)
plt.title('Accuracy plot for Train Data')
plt.show()


```



### <font color = blue>[5.4.3] HeatMap for Cross Validation Data </font>



```
In [119]: text = pd.DataFrame({'max_depth': Y, 'n_estimators': X, 'AUC': cv_data})
text_pivoted = text.pivot("max_depth", "n_estimators", "AUC")
ax = sns.heatmap(text_pivoted, annot=True)
plt.title('Accuracy plot for Cross Validation Data')
plt.show()
```




#### <font color = blue>[5.4.4] ROC Curve of Random Forest </font>

```
In [120]: xgb =XGBClassifier(class_weight = 'balanced', booster = 'gbtree', max_depth =
5, n_estimators = 150)
xgb.fit(tfidf_sent_vectors_train,Y_Train)
prediction = xgb.predict_proba(tfidf_sent_vectors_test)[:,-1]
print(prediction)
print(xgb)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, xgb.predict_proba(tfidf_s
ent_vectors_train)[:,-1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, xgb.predict_proba(tfidf_sent
_vectors_test)[:,-1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of Random Forest')
plt.show()
```

```
[0.9467274  0.937574  0.52574444 ... 0.8006779  0.7914084  0.3287601 ]
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
  colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
  gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=5,
  min_child_weight=1, missing=None, n_estimators=150, n_jobs=1,
  nthread=None, objective='binary:logistic', random_state=0,
  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
  silent=None, subsample=1, verbosity=1)
```



### <font color = blue>[5.4.5]Train and Test Accuracy</font>

```
In [121]: Training_Accuracy_Bow = xgb.score(tfidf_sent_vectors_train, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Bow)
Training_Error_Bow = 1 - Training_Accuracy_Bow
print('Training_Error=%0.3f'%Training_Error_Bow)


Test_Accuracy_Bow = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Bow)
Test_Error_Bow = 1 - Test_Accuracy_Bow
print('Test_Error=%0.3f'%Test_Error_Bow)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_alpha_bow, Test_Accuracy_Bow))

Training_Accuracy=0.902
Training_Error=0.098
Test_Accuracy=0.879
Test_Error=0.121
```

### <font color = blue>[5.4.6] Confusion Matrix for Train Data </font>

```
In [122]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, xgb.predict(tfidf_sent_vectors_train))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=np.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, fmt='d')
```


Out[122]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdc30ac240>



**<font color = blue>[5.4.7] Confusion Matrix for Test Data </font>**

```
In [123]: from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, xgb.predict(tfidf_sent_vectors_test))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[123]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1fdc30ace10>



### <font color = blue>[5.4.8] Classification Report</font>

```
In [124]: print(classification_report(Y_Test, prediction.round()))
```

	precision	recall	f1-score	support
0	0.72	0.41	0.52	4234
1	0.90	0.97	0.93	22098
avg / total	0.87	0.88	0.86	26332

## <font color = Blue>[5.4.9] Feature Importance with WordCloud </font>

```
In [125]: from wordcloud import WordCloud
Imp_features = model.get_feature_names()
text = ''
feat=rf.feature_importances_
features=np.argsort(feat)[::-1]
for i in features[0:20]:
    text += Imp_features[i]

plt.figure(figsize=(12,6))
wc = WordCloud(background_color='white',
               width=600,
               height=300,
               max_font_size=50,
               max_words=20).generate(text)
plt.imshow(wc, interpolation='bilinear')
plt.title("Imp Features", fontsize=20)
plt.axis("off")
plt.show()
```

Imp Features

aaaaabandonabcessedabdasahabbeyaaaaaaahhhhhhabidingabcaberrationabbyaahaaaandabideabettaraafcoanazonaahsaadpabbreviation

## <font color = Green>Pretty Table</font>

```
In [127]: from prettytable import PrettyTable
comparision = PrettyTable()
comparision.field_names = ["Vectorizer", "Algorithm", "Best Depth", "Best Esti
mator","Accuracy", "Training Error", "Test Error"]
comparision.add_row(["BOW", "RF", 10,150,0.90,0.145,0.155])
comparision.add_row(["TF-IDF", "RF", 10,150, 0.91,0.145,0.846])
comparision.add_row(["Avg W2V", "RF", 10,150,0.88,0.216,0.225])
comparision.add_row(["TF-IDFWeighted W2V", "RF", 10,150,0.86,0.225,0.231])
comparision.add_row(["BOW", "GBDT", 5,150,0.92,0.094,0.102])
comparision.add_row(["TF-IDF", "GBDT", 5,150,0.92,0.090,0.100])
comparision.add_row(["Avg W2V", "GBDT", 5,150,0.907,0.088,0.110])
comparision.add_row(["TF-IDFWeighted W2V", "GBDT", 5,150,0.88,0.098,0.121])
print(comparision)
```

```
+-----+-----+-----+-----+-----+
+-----+
|      Vectorizer      | Algorithm | Best Depth | Best Estimator | Accuracy | T
raining Error | Test Error |
+-----+-----+-----+-----+-----+
+-----+
|      BOW      | RF | 10 | 150 | 0.9 |
0.145 | 0.155 |
|      TF-IDF      | RF | 10 | 150 | 0.91 |
0.145 | 0.846 |
|      Avg W2V      | RF | 10 | 150 | 0.88 |
0.216 | 0.225 |
| TF-IDFWeighted W2V | RF | 10 | 150 | 0.86 |
0.225 | 0.231 |
|      BOW      | GBDT | 5 | 150 | 0.92 |
0.094 | 0.102 |
|      TF-IDF      | GBDT | 5 | 150 | 0.92 |
0.09 | 0.1 |
|      Avg W2V      | GBDT | 5 | 150 | 0.907 |
0.088 | 0.11 |
| TF-IDFWeighted W2V | GBDT | 5 | 150 | 0.88 |
0.098 | 0.121 |
+-----+-----+-----+-----+-----+
+-----+
```

<font color = Green>Conclusion</font>

1. Applied Random Forest and GBDT using XGBosst on all the 4 vectorizers(BOW, TFIDF, AVG-W2V, TFIDF-AVG\_W2V).
2. Sorted the data based on Time and Considered 100 K data points for Training set 70K, Test set: 30K.
3. Plotted Seaborn Heatmap for the performance of the modal on both train and cross validation data for hyper parameter n\_estimators as X-axis and max\_depth as Y-axis.
4. Used accuracy and n\_estimators, max\_depth as a metric for hyperparameter tuning. And took the best depth in the range of [1,5,10] and the best n\_estimators in range [5, 10, 20, 40, 60, 80, 100, 120, 150] for GBDT took the best depth in the range of [1,5].
5. Found the top 20 features for featurizations BOW & TFIDF & TFIDFW2V and printed the important features using WordCloud.
6. Plotted the ROC curve and Confusion Matix for Train and Test Data for each featurizations.
7. With reference to the pretty table, here is my understanding: Accuracy imporved while using XGBoost alogorithm.