Amazon Fine Food Reviews Analysis Using SVM

Data Source: https://www.kaggle.com/snap/amazon-fine-food-reviews (https://www.kaggle.com/snap/amazon-fine-food-reviews)

EDA: https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/ (https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454
Number of users: 256,059
Number of products: 74,258
Timespan: Oct 1999 - Oct 2012
Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unqiue identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (Rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use the Score/Rating. A rating of 4 or 5 could be cosnidered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is nuetral and ignored. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score id above 3, then the recommendation wil be set to "positive". Otherwise, it will be set to "negative".

In [76]:
```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")



import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn import cross_validation
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.calibration import CalibratedClassifierCV
from sklearn.svm import SVC
from sklearn.linear_model import SGDClassifier
from sklearn.svm import LinearSVC

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
```

[1]. Reading Data

In [77]:
```
# using SQLite Table to read data.
con = sqlite3.connect('D:\\TGM\\ML\\AmazonFineFoodReviews\\database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 da
ta points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3
LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 L
IMIT 100000""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a
negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

```
Number of data points in our data (100000, 10)
```

Out[77]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulne
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1

In [78]:

```
display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

In [79]:
```
print(display.shape)
display.head()
```

(80668, 7)

Out[79]:

	UserId	ProductId	ProfileName	Time	Score	Text	COL
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

In [80]:
```
display[display['UserId']=='AZY10LLTJ71NX']
```

Out[80]:

	UserId	ProductId	ProfileName	Time	Score	Text
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...

In [81]:
```
display['COUNT(*)'].sum()
```

Out[81]: 393063

Exploratory Data Analysis

[2] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [82]:
```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

Out[82]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpful
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2

As can be seen above the same user has multiple reviews of the with the same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delelte the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

In [83]:
```
#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inp
lace=False, kind='quicksort', na_position='last')
```

In [84]:
```
#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"
}, keep='first', inplace=False)
final.shape
```

Out[84]: (87775, 10)

In [85]:
```
#Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[85]: 87.775

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calcualtions

In [86]:
```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[86]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfuln
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2

In [87]:
```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

In [88]:
```
#Before starting the next phase of preprocessing lets see the number of entrie
s left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
print(final['Score'].value_counts())
```

```
(87773, 10)
1    73592
0    14181
Name: Score, dtype: int64
```

[3]. Text Preprocessing.

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was obsereved to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

In [89]:
```
# printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

```
My dogs loves this chicken but its a product from China, so we wont be buying
it anymore.  Its very hard to find any chicken products made in the USA but t
hey are out there, but this one isnt.  Its too bad too because its a good pro
duct but I wont take any chances till they know what is going on with the chi
na imports.
==================================================
The Candy Blocks were a nice visual for the Lego Birthday party but the candy
has little taste to it.  Very little of the 2 lbs that I bought were eaten an
d I threw the rest away.  I would not buy the candy again.
==================================================
was way to hot for my blood, took a bite and did a jig  lol
==================================================
My dog LOVES these treats. They tend to have a very strong fish oil smell. So
if you are afraid of the fishy smell, don't get it. But I think my dog likes
it because of the smell. These treats are really small in size. They are grea
t for training. You can give your dog several of these without worrying about
him over eating. Amazon's price was much more reasonable than any other retai
ler. You can buy a 1 pound bag on Amazon for almost the same price as a 6 oun
ce bag at other retailers. It's definitely worth it to buy a big bag if your
dog eats them a lot.
==================================================
```

In [90]:
```
# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can\'t", "can not", phrase)

    # general
    phrase = re.sub(r"n\'t", " not", phrase)
    phrase = re.sub(r"\'re", " are", phrase)
    phrase = re.sub(r"\'s", " is", phrase)
    phrase = re.sub(r"\'d", " would", phrase)
    phrase = re.sub(r"\'ll", " will", phrase)
    phrase = re.sub(r"\'t", " not", phrase)
    phrase = re.sub(r"\'ve", " have", phrase)
    phrase = re.sub(r"\'m", " am", phrase)
    return phrase
```

In [91]:
```
sent_4900 = decontracted(sent_4900)
print(sent_4900)
print("="*50)
```

```
My dog LOVES these treats. They tend to have a very strong fish oil smell. So
if you are afraid of the fishy smell, do not get it. But I think my dog likes
it because of the smell. These treats are really small in size. They are grea
t for training. You can give your dog several of these without worrying about
him over eating. Amazon is price was much more reasonable than any other reta
iler. You can buy a 1 pound bag on Amazon for almost the same price as a 6 ou
nce bag at other retailers. It is definitely worth it to buy a big bag if you
r dog eats them a lot.
==================================================
```

In [92]:
```
#remove words with numbers python: https://stackoverflow.com/a/18082370/408403
9
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

```
My dogs loves this chicken but its a product from China, so we wont be buying
it anymore.  Its very hard to find any chicken products made in the USA but t
hey are out there, but this one isnt.  Its too bad too because its a good pro
duct but I wont take any chances till they know what is going on with the chi
na imports.
```

In [93]:
```
#remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

```
was way to hot for my blood took a bite and did a jig lol
```

In [94]:
```
# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in the 1st
step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours',
'ourselves', 'you', "you're", "you've",\
            "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he'
, 'him', 'his', 'himself', \
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'it
self', 'they', 'them', 'their',\
            'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 't
hat', "that'll", 'these', 'those', \
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have',
'has', 'had', 'having', 'do', 'does', \
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'becau
se', 'as', 'until', 'while', 'of', \
            'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into',
'through', 'during', 'before', 'after',\
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on',
'off', 'over', 'under', 'again', 'further',\
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'a
ll', 'any', 'both', 'each', 'few', 'more',\
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'tha
n', 'too', 'very', \
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "shoul
d've", 'now', 'd', 'll', 'm', 'o', 're', \
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn',
"didn't", 'doesn', "doesn't", 'hadn',\
            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'm
a', 'mightn', "mightn't", 'mustn',\
            "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shoul
dn't", 'wasn', "wasn't", 'weren', "weren't", \
            'won', "won't", 'wouldn', "wouldn't"])
```

In [95]:
```
# Combining all the above stundents
from bs4 import BeautifulSoup
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentance in tqdm(final['Text'].values):
    sentance = re.sub(r"http\S+", "", sentance)
    sentance = BeautifulSoup(sentance, 'lxml').get_text()
    sentance = decontracted(sentance)
    sentance = re.sub("\S*\d\S*", "", sentance).strip()
    sentance = re.sub('[^A-Za-z]+', ' ', sentance)
    # https://gist.github.com/sebleier/554280
    sentance = ' '.join(e.lower() for e in sentance.split() if e.lower() not i
n stopwords)
    preprocessed_reviews.append(sentance.strip())
```

```
100%|
| 87773/87773 [00:56<00:00, 1566.45it/s]
```

In [96]:
```
preprocessed_reviews[1500]
```
Out[96]:
```
'way hot blood took bite jig lol'
```

In [97]:
```
final['cleaned_text']=preprocessed_reviews
```

In [98]:
```
final.shape
```
Out[98]:
```
(87773, 11)
```

In [99]:
```
final["Score"].value_counts()
```
Out[99]:
```
1    73592
0    14181
Name: Score, dtype: int64
```

In [100]:
```
#Sorted the data based on time and took 100k data points
final["Time"] = pd.to_datetime(final["Time"], unit = "s")
final = final.sort_values(by = "Time")
final.head()
```

Out[100]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	He
70688	76882	B00002N8SM	A32DW342WBJ6BX	Buttersugar	0	0
1146	1245	B00002Z754	A29Z5PI9BW2PU3	Robbie	7	7
1145	1244	B00002Z754	A3B8RCEI0FXFI6	B G Chase	10	10
28086	30629	B00008RCMI	A19E94CF5O1LY7	Andrew Arnold	0	0
28087	30630	B00008RCMI	A284C7M23F0APC	A. Mendoza	0	0

In [101]:
```
Y = final['Score'].values
X = final['cleaned_text'].values
#X = preprocessed_reviews
#Y = np.array(final['Score'])
print(Y.shape)
print(type(Y))
print(X.shape)
print(type(X))
```

```
(87773,)
<class 'numpy.ndarray'>
(87773,)
<class 'numpy.ndarray'>
```

In [102]:
```
# split the data set into train and test
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X,Y,test_size=0.3, random_
state=0)

# split the train data set into cross validation train and cross validation te
st
X_tr, X_cv, Y_tr, Y_cv = train_test_split(X,Y, test_size=0.3, random_state=0)

print('='*100)
print("After splitting")
print("X_Train Shape:",X_Train.shape, "Y_Train Shape:",Y_Train.shape)
print("X_cv Shape:",X_cv.shape,       "Y_cv Shape",Y_cv.shape)
print("X_Test Shape",X_Test.shape,    "Y_Test Shape",Y_Test.shape)
```

```
=============================================================================
=======================
After splitting
X_Train Shape: (61441,) Y_Train Shape: (61441,)
X_cv Shape: (26332,) Y_cv Shape (26332,)
X_Test Shape (26332,) Y_Test Shape (26332,)
```

[3.2] Preprocess Summary

[4] Featurization

[4.1] BAG OF WORDS

In [103]:
```
#BoW
count_vect = CountVectorizer(ngram_range=(1,2)) #in scikit-learn
count_vect.fit(X_Train)
print("some feature names ", count_vect.get_feature_names()[:10])
X_Train_Bow = count_vect.transform(X_Train)
X_Test_Bow = count_vect.transform(X_Test)
X_CV_Bow = count_vect.transform(X_cv)

print('='*50)

#final_counts = count_vect.transform(X_Test)

print("the type of X Train : ",type(X_Train_Bow))
print("the shape of Train BOW vectorizer ",X_Train_Bow.get_shape())
print("the shape of Test BOW vectorizer ",X_Test_Bow.get_shape())
print("the shape of CV BOW vectorizer ",X_CV_Bow.get_shape())
#print("the number of unique words ", final_counts.get_shape()[1])
```

```
some feature names  ['aa', 'aa caffene', 'aa coffee', 'aa cups', 'aa dark',
'aa extra', 'aa favorite', 'aa kona', 'aa may', 'aa not']
==================================================
the type of X Train :  <class 'scipy.sparse.csr.csr_matrix'>
the shape of Train BOW vectorizer  (61441, 1076376)
the shape of Test BOW vectorizer  (26332, 1076376)
the shape of CV BOW vectorizer  (26332, 1076376)
```

In [104]:
```
import warnings
warnings.filterwarnings('ignore')
scalar = StandardScaler(with_mean=False)
X_Train_Bow = scalar.fit_transform(X_Train_Bow)
X_Test_Bow = scalar.transform(X_Test_Bow)
X_CV_Bow = scalar.transform(X_CV_Bow)

print("the type of X Train : ",type(X_Train_Bow))
print("the shape of Train BOW vectorizer ",X_Train_Bow.get_shape())
print("the shape of Test BOW vectorizer ",X_Test_Bow.get_shape())
print("the shape of CV BOW vectorizer ",X_CV_Bow.get_shape())
```

```
the type of X Train :  <class 'scipy.sparse.csr.csr_matrix'>
the shape of Train BOW vectorizer  (61441, 1076376)
the shape of Test BOW vectorizer  (26332, 1076376)
the shape of CV BOW vectorizer  (26332, 1076376)
```

[4.1] Linear Kernel

In [109]:
```
import math
def Optimal_Lamda_L1(X_Train,Y_Train,X_CV,Y_CV):
    train_AUC_L1 = []
    CV_AUC_L1 = []
    tuned_parameters=[10**-4, 10**-3, 10**-2, 10**-1, 1,10**1, 10**2, 10**3, 1
0**4]
    for j in  tqdm(tuned_parameters):
        clf = SGDClassifier(alpha=j, class_weight='balanced', penalty='l1', lo
ss='hinge')
        calibrated_clf = CalibratedClassifierCV(clf, cv=5)
        calibrated_clf.fit(X_Train, Y_Train)
        y_train_pred =  calibrated_clf.predict_proba(X_Train)[:,1]
        y_cv_pred =  calibrated_clf.predict_proba(X_CV)[:,1]
        train_AUC_L1.append(roc_auc_score(Y_Train,y_train_pred))
        CV_AUC_L1.append(roc_auc_score(Y_CV, y_cv_pred))

    #Error plots with penaly L1
    plt.plot(np.log(tuned_parameters), train_AUC_L1, label='Train AUC-L1')
    plt.plot(np.log(tuned_parameters), CV_AUC_L1, label='CV AUC-L1')
    plt.legend()
    plt.xlabel("Hyperparameter (Lambda)")
    plt.ylabel("AUC")
    plt.title("AUC PLOT")
    plt.show()

    #Cv auc scores with penalty L1
    print("CV AUS Scores with Penalty=? Cv auc scores with penalty L1")
    print(CV_AUC_L1)
    print("Maximun AUC value :",max(CV_AUC_L1))
    print("Index",CV_AUC_L1.index(max(CV_AUC_L1)))
```

In [110]:
```
import math
def Optimal_Lamda_L2(X_Train,Y_Train,X_CV,Y_CV):
    train_AUC_L2 = []
    CV_AUC_L2 = []
    cv_scores = []
    tuned_parameters=[10**-4, 10**-3, 10**-2, 10**-1, 1,10**1, 10**2, 10**3, 1
0**4]
    for j in  tqdm(tuned_parameters):
        clf = SGDClassifier(alpha=j, class_weight='balanced', penalty='l2', lo
ss='hinge')
        calibrated_clf = CalibratedClassifierCV(clf, cv=5)
        calibrated_clf.fit(X_Train, Y_Train)
        y_train_pred =  calibrated_clf.predict_proba(X_Train)[:,1]
        y_cv_pred =  calibrated_clf.predict_proba(X_CV)[:,1]
        train_AUC_L2.append(roc_auc_score(Y_Train,y_train_pred))
        CV_AUC_L2.append(roc_auc_score(Y_CV, y_cv_pred))

    #Error plots with penaly L2
    plt.plot(np.log(tuned_parameters), train_AUC_L2, label='Train AUC-L2')
    plt.plot(np.log(tuned_parameters), CV_AUC_L2, label='CV AUC-L2')
    plt.legend()
    plt.xlabel("Hyperparameter (Lambda)")
    plt.ylabel("AUC")
    plt.title("AUC PLOT")
    plt.show()

    #Cv auc scores with penalty L2
    print("CV AUS Scores with Penalty=? Cv auc scores with penalty L2")
    print(CV_AUC_L2)
    print("Maximun AUC value :",max(CV_AUC_L2))
    print("Index",CV_AUC_L2.index(max(CV_AUC_L2)))
```

In [107]:
```
Optimal_Lamda_L1(X_Train_Bow, Y_Train, X_CV_Bow,Y_cv)
```

```
100%|
| 9/9 [00:28<00:00,  3.32s/it]
```

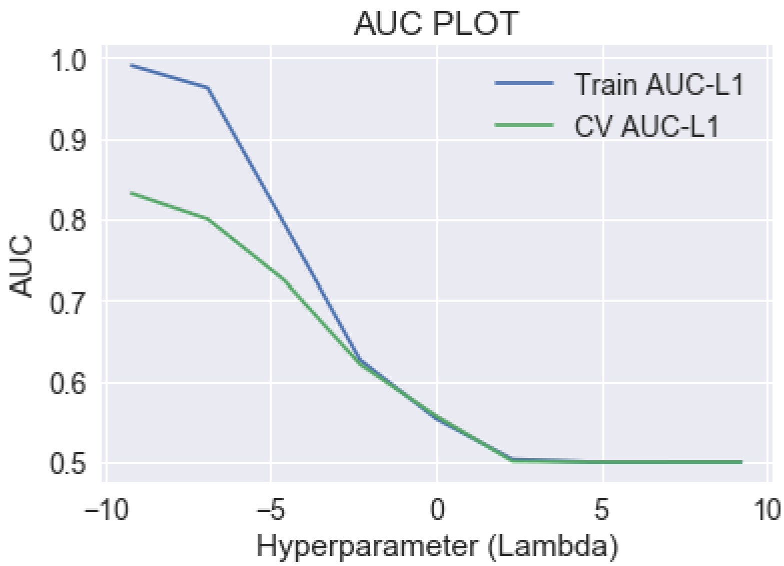


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L1
[0.831609221053483, 0.7997954681454404, 0.7246442854099527, 0.62108846161426
4, 0.5570247894753875, 0.5010185390513413, 0.5000226264820345, 0.5, 0.5]
Maximun AUC value : 0.831609221053483
Index 0
```

In [112]:
```
Optimal_Lamda_L2(X_Train_Bow, Y_Train, X_CV_Bow,Y_cv)
```

```
100%|██████████████████████████████████████████████████
████████| 9/9 [00:32<00:00,  3.81s/it]
```

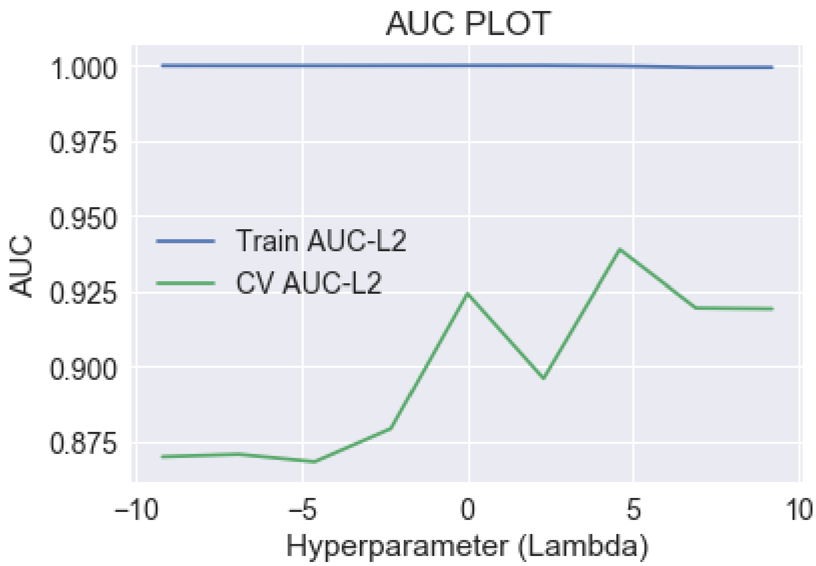


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L2
[0.8699834246323106, 0.8707944295717452, 0.8683134256630607, 0.87934106212062
7, 0.924285089740454, 0.895990551044296, 0.9388157908518728, 0.91944666719080
58, 0.9191993897754295]
Maximun AUC value : 0.9388157908518728
Index 6
```

[4.1.3] ROC Curve of SVM

In [114]:
```
#Testing with test data
clf = SGDClassifier(alpha=0.1, class_weight='balanced', loss='hinge')
calibrated_clf = CalibratedClassifierCV(clf, cv=5)
calibrated_clf.fit(X_Train_Bow,Y_Train)
prediction = calibrated_clf.predict_proba(X_Test_Bow)[:,1]
print(prediction)
print(clf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, calibrated_clf.predict_pr
oba(X_Train_Bow)[:,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, calibrated_clf.predict_proba
(X_Test_Bow)[:,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of MNB')
plt.show()
```

```
[0.90236714 0.90391762 0.7174323  ... 0.86114097 0.88722453 0.32113633]
SGDClassifier(alpha=0.1, average=False, class_weight='balanced', epsilon=0.1,
       eta0=0.0, fit_intercept=True, l1_ratio=0.15,
       learning_rate='optimal', loss='hinge', max_iter=None, n_iter=None,
       n_jobs=1, penalty='l2', power_t=0.5, random_state=None,
       shuffle=True, tol=None, verbose=0, warm_start=False)
```

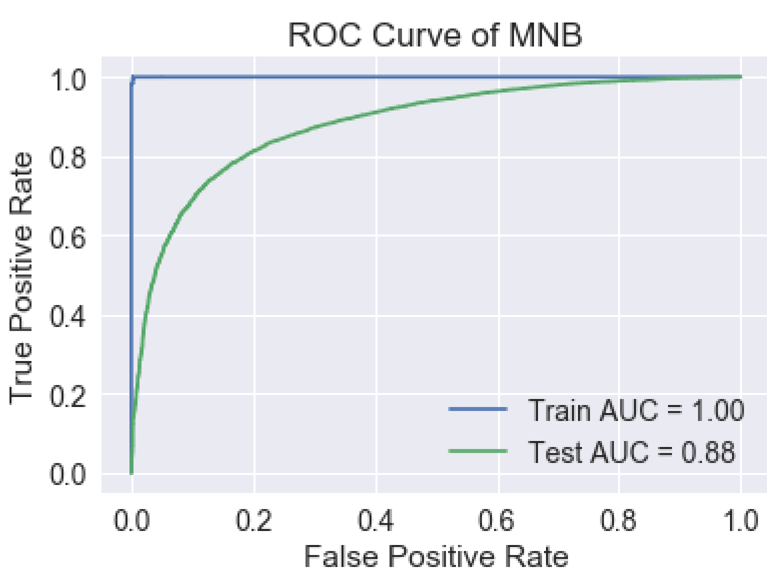

[4.1.4]Train and Test Accuracy

In [115]:
```
Training_Accuracy_Bow = calibrated_clf.score(X_Train_Bow, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Bow)
Training_Error_Bow = 1 - Training_Accuracy_Bow
print('Training_Error=%0.3f'%Training_Error_Bow)

Test_Accuracy_Bow = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Bow)
Test_Error_Bow = 1 - Test_Accuracy_Bow
print('Test_Error=%0.3f'%Test_Error_Bow)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))
```

```
Training_Accuracy=0.978
Training_Error=0.022
Test_Accuracy=0.873
Test_Error=0.127
```

[4.1.5] Confusion Matrix

In [116]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, calibrated_clf.predict(X_Train_Bow))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=n
p.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[116]: <matplotlib.axes._subplots.AxesSubplot at 0x23d131b2c50>

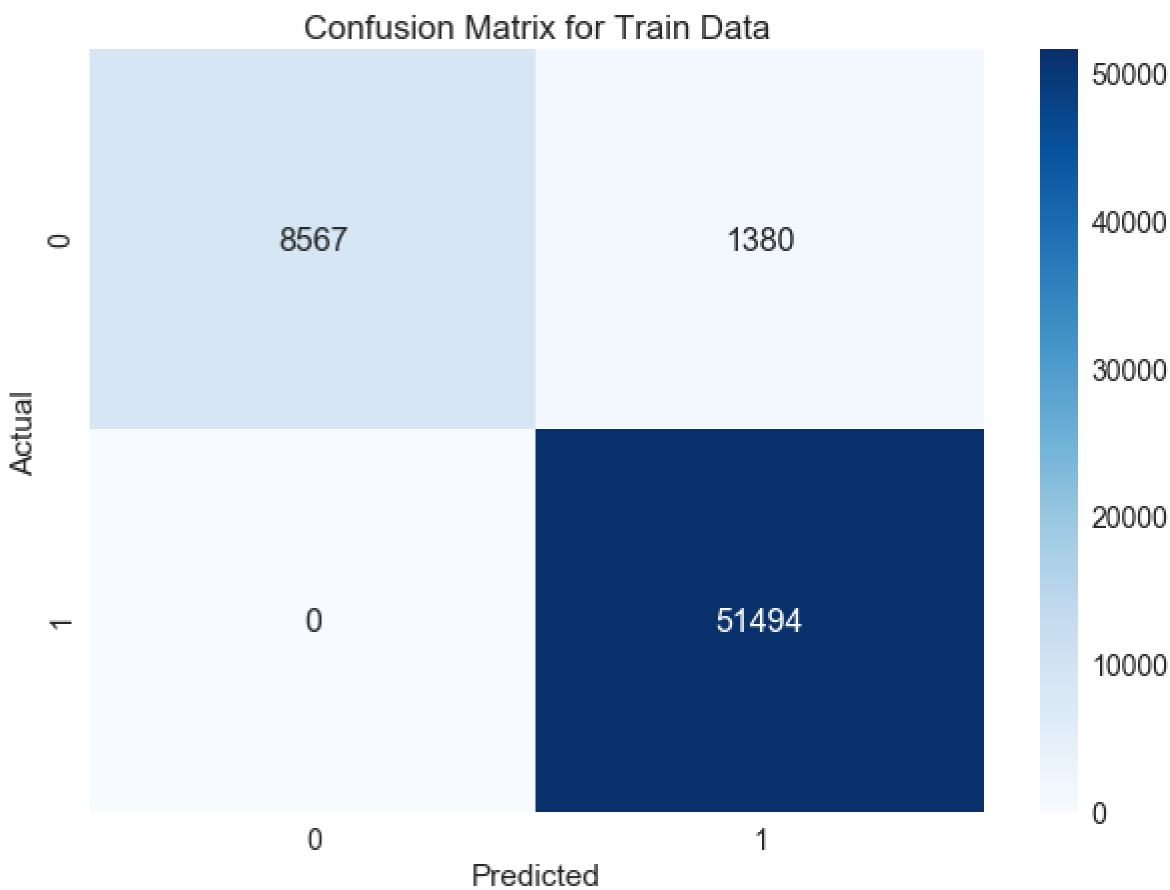

In [117]:

```
#With the reference of below link:
#https://www.kaggle.com/agungor2/various-confusion-matrix-plots
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, calibrated_clf.predict(X_Test_Bow))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[117]: <matplotlib.axes._subplots.AxesSubplot at 0x23d493c18d0>

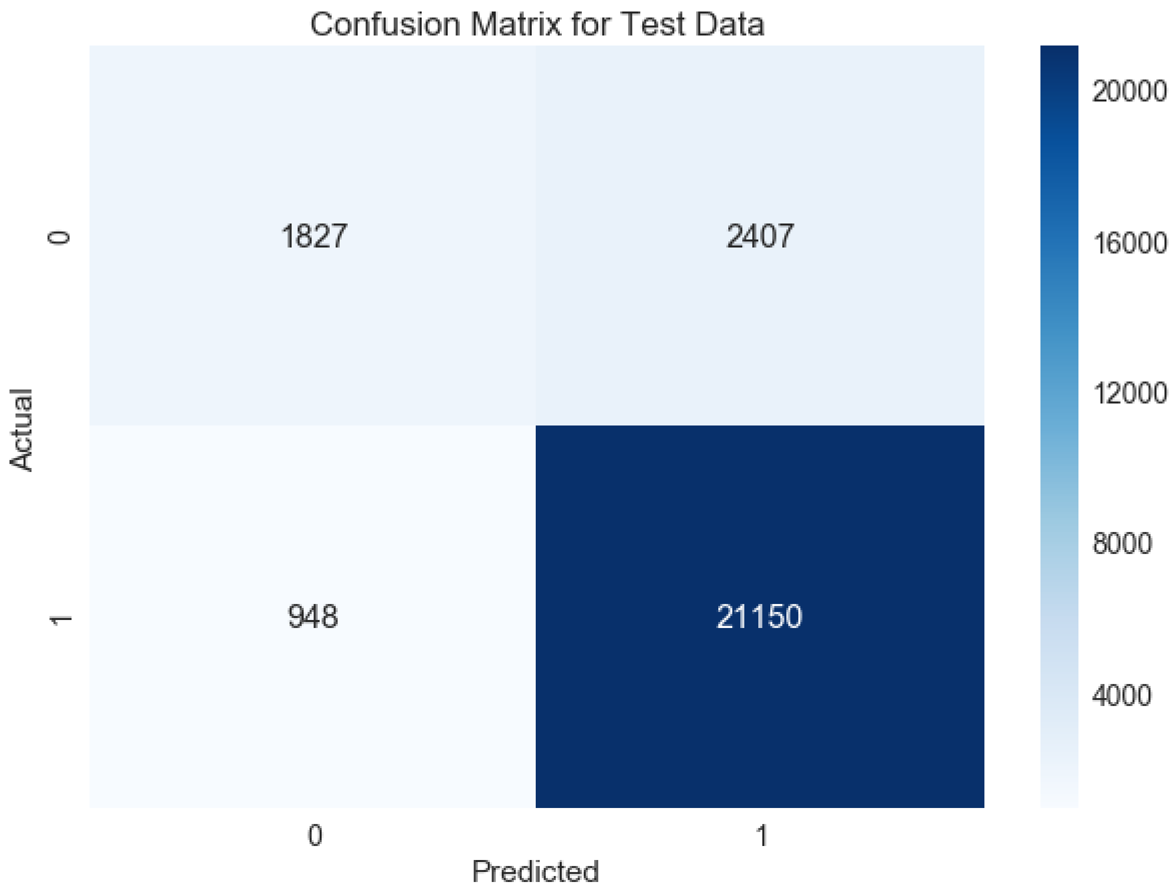

[4.1.6] Classification Report

In [118]:
```
from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

```
             precision    recall  f1-score   support

          0       0.66      0.43      0.52      4234
          1       0.90      0.96      0.93     22098

avg / total       0.86      0.87      0.86     26332
```

[4.1.7] Feature Importance

In [119]:
```
#https://medium.com/@aneesha/visualising-top-features-in-linear-svm-with-sciki
t-learn-and-matplotlib-3454ab18a14d
def plot_coefficients(classifier, feature_names, top_features=20):
    coef = classifier.coef_.ravel()
    top_positive_coefficients = np.argsort(coef)[-top_features:]
    top_negative_coefficients = np.argsort(coef)[:top_features]
    top_coefficients = np.hstack([top_negative_coefficients, top_positive_coef
ficients])

    # create plot
    plt.figure(figsize=(15, 5))
    colors = ['red' if c < 0 else 'blue' for c in coef[top_coefficients]]
    plt.bar(np.arange(2 * top_features), coef[top_coefficients], color=colors)
    feature_names = np.array(feature_names)
    plt.xticks(np.arange(1, 1 + 2 * top_features), feature_names[top_coefficie
nts], rotation=60, ha='right')
    plt.show()
```

Feature Importance for Positive and Negetive Class

In [121]:
```
clf = SGDClassifier(alpha=0.1, class_weight='balanced')
clf.fit(X_Train_Bow, Y_Train)
plot_coefficients(clf, count_vect.get_feature_names())
```

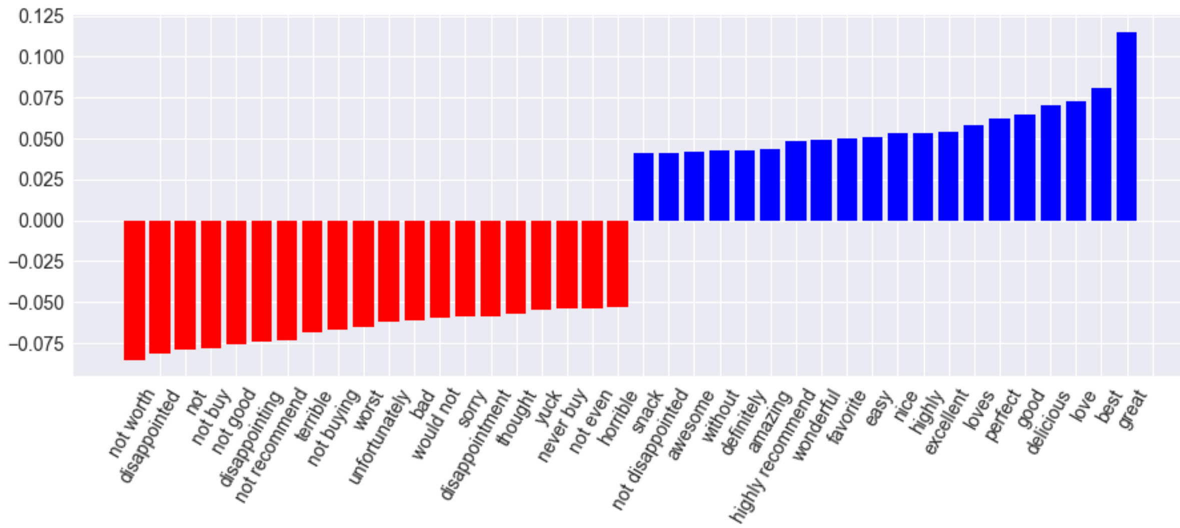

[4.2] TF-IDF

In [122]:
```
#TF-IDF
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=5)
tf_idf_vect.fit_transform(X_Train)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_featu
re_names()[0:10])
print('='*50)

X_Train_TfIdf = tf_idf_vect.transform(X_Train)
X_Test_TfIdf = tf_idf_vect.transform(X_Test)
X_CV_TfIdf = tf_idf_vect.transform(X_cv)

#final_tf_idf = tf_idf_vect.transform(X_Test)
print("the type of count vectorizer ",type(X_Train_TfIdf))
print("the shape of out text TFIDF vectorizer ",X_Train_TfIdf.get_shape())
print("the shape of out text TFIDF vectorizer ",X_Test_TfIdf.get_shape())
print("the shape of out text TFIDF vectorizer ",X_CV_TfIdf.get_shape())
#print("the number of unique words including both unigrams and bigrams ", fina
l_tf_idf.get_shape()[1])
```

```
some sample features(unique words in the corpus) ['aa', 'aaa', 'aafco', 'abac
k', 'abandon', 'abandoned', 'abdominal', 'abilities', 'ability', 'ability mak
e']
==================================================
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer  (61441, 80521)
the shape of out text TFIDF vectorizer  (26332, 80521)
the shape of out text TFIDF vectorizer  (26332, 80521)
```

In [123]:
```
scalar = StandardScaler(with_mean=False)
X_Train_TfIdf = scalar.fit_transform(X_Train_TfIdf)
X_Test_TfIdf = scalar.transform(X_Test_TfIdf)
X_CV_TfIdf = scalar.transform(X_CV_TfIdf)

print("the type of count vectorizer ",type(X_Train_TfIdf))
print("the shape of out text TFIDF vectorizer ",X_Train_TfIdf.get_shape())
print("the shape of out text TFIDF vectorizer ",X_Test_TfIdf.get_shape())
print("the shape of out text TFIDF vectorizer ",X_CV_TfIdf.get_shape())
```

```
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer  (61441, 80521)
the shape of out text TFIDF vectorizer  (26332, 80521)
the shape of out text TFIDF vectorizer  (26332, 80521)
```

[4.2.1] Hyperameter tuning with L1 Regulizer and AUC Plot

In [124]:
```
Optimal_Lamda_L1(X_Train_TfIdf, Y_Train, X_CV_TfIdf,Y_cv)
```

```
100%|██████████████████████████████████████████████████
██████████| 9/9 [00:18<00:00,  1.94s/it]
```

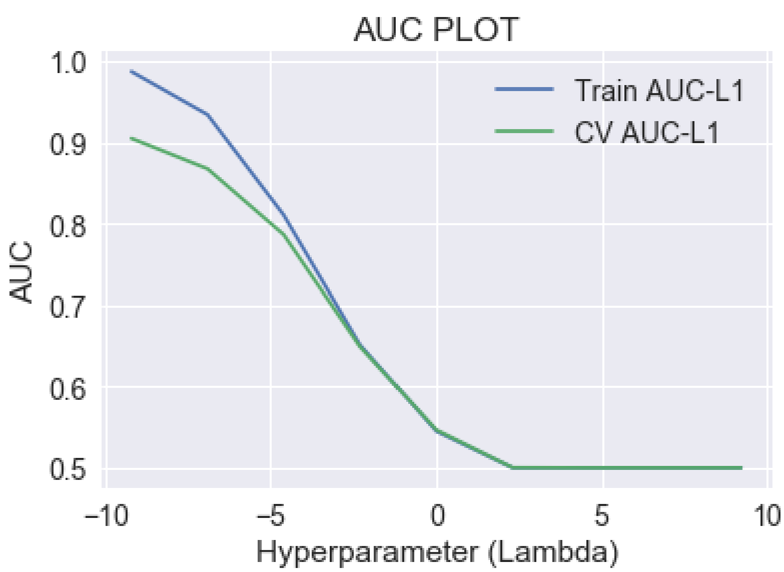


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L1
[0.9048317446913698, 0.8674258733148721, 0.7864371437184119, 0.64921860293989
08, 0.546532236719559, 0.5002657622999672, 0.5, 0.5, 0.5]
Maximun AUC value : 0.9048317446913698
Index 0
```

[4.2.2] Hyperameter tuning with L2 Regulizer and AUC Plot

In [125]:
```
Optimal_Lamda_L2(X_Train_TfIdf, Y_Train, X_CV_TfIdf,Y_cv)
```

```
100%|
| 9/9 [00:13<00:00,  1.37s/it]
```

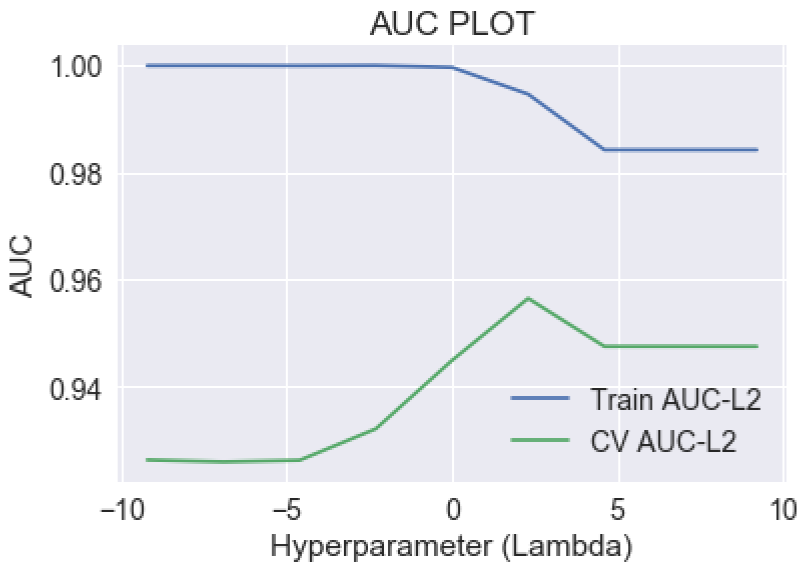


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L2
[0.9262989535214652, 0.9259391529115398, 0.9262368669677858, 0.93211943165697
29, 0.9449295154623841, 0.9565408553036794, 0.9475433070011101, 0.94754360626
49256, 0.9475434031930507]
Maximun AUC value : 0.9565408553036794
Index 5
```

[4.2.3] ROC Curve of SVM

In [126]:
```
#Testing with test data
clf = SGDClassifier(alpha=0.1, class_weight='balanced', loss='hinge')
calibrated_clf = CalibratedClassifierCV(clf, cv=3)
calibrated_clf.fit(X_Train_TfIdf,Y_Train)
prediction = calibrated_clf.predict_proba(X_Test_TfIdf)[:,1]
print(prediction)
print(clf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, calibrated_clf.predict_pr
oba(X_Train_TfIdf)[:,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, calibrated_clf.predict_proba
(X_Test_TfIdf)[:,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of MNB')
plt.show()
```

```
[0.92766173 0.99322665 0.13431146 ... 0.25847924 0.23463799 0.35375614]
SGDClassifier(alpha=0.1, average=False, class_weight='balanced', epsilon=0.1,
       eta0=0.0, fit_intercept=True, l1_ratio=0.15,
       learning_rate='optimal', loss='hinge', max_iter=None, n_iter=None,
       n_jobs=1, penalty='l2', power_t=0.5, random_state=None,
       shuffle=True, tol=None, verbose=0, warm_start=False)
```

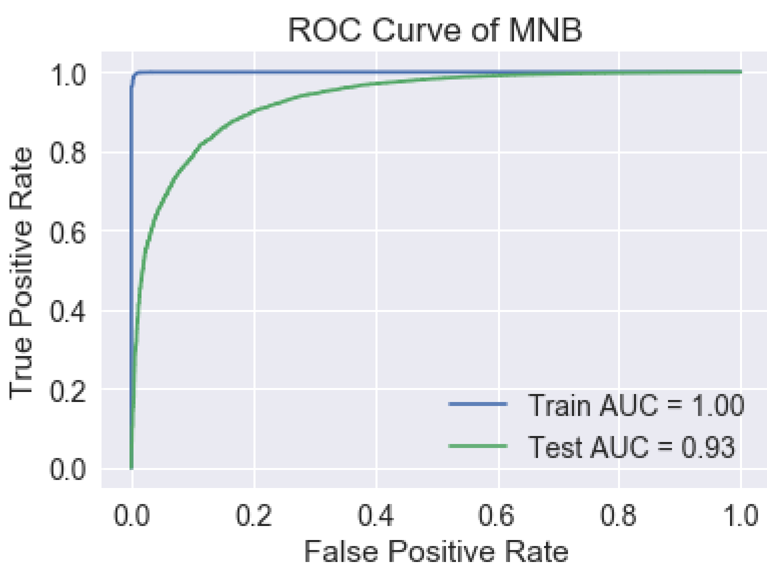

[4.2.4]Train and Test Accuracy

In [129]:
```
Training_Accuracy_Tfidf = calibrated_clf.score(X_Train_TfIdf, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Tfidf)
Training_Error_Tfidf = 1 - Training_Accuracy_Tfidf
print('Training_Error=%0.3f'%Training_Error_Tfidf)

Test_Accuracy_Tfidf = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Tfidf)
Test_Error_Tfidf = 1 - Test_Accuracy_Tfidf
print('Test_Error=%0.3f'%Test_Error_Tfidf)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))
```

```
Training_Accuracy=0.989
Training_Error=0.011
Test_Accuracy=0.908
Test_Error=0.092
```

[4.2.5] Confusion Matrix

In [130]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, calibrated_clf.predict(X_Train_TfIdf))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=n
p.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[130]: <matplotlib.axes._subplots.AxesSubplot at 0x23d48f79550>

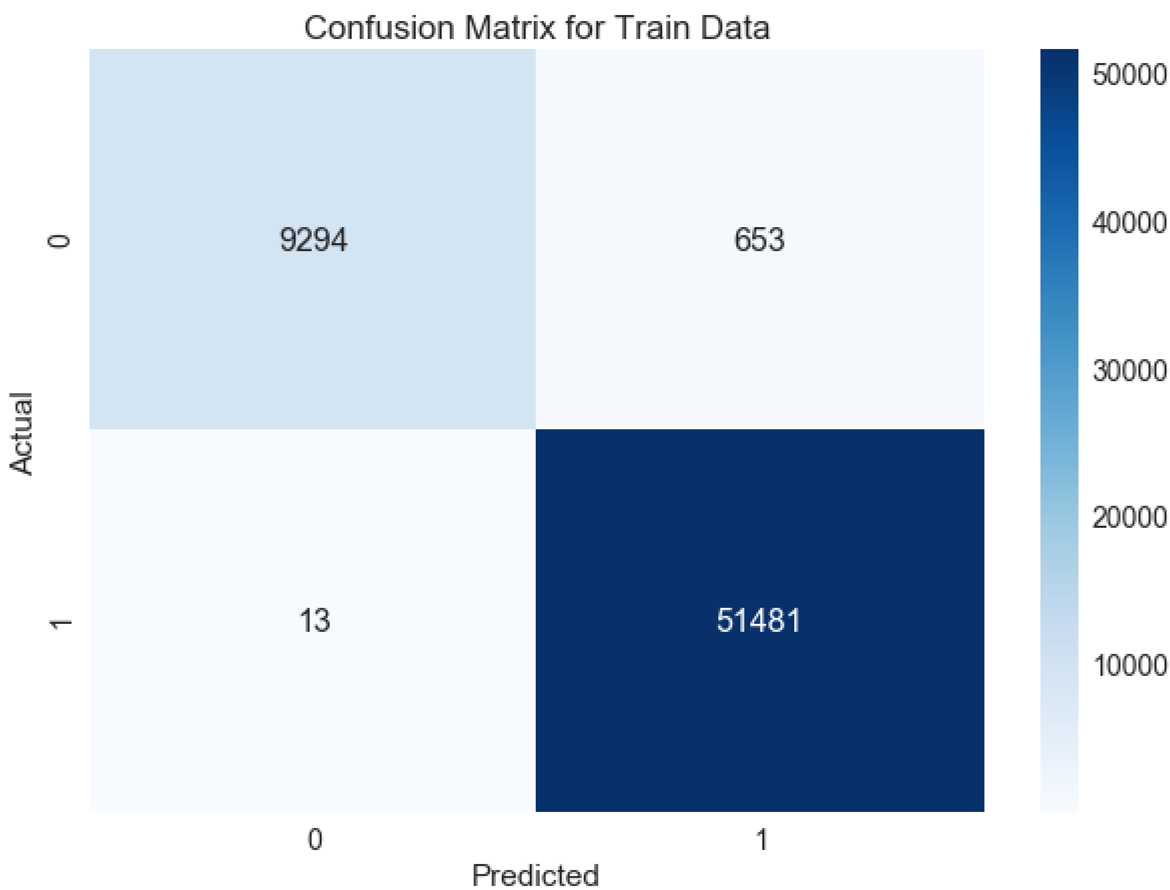

In [131]:
```
#With the reference of below link:
#https://www.kaggle.com/agungor2/various-confusion-matrix-plots
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, calibrated_clf.predict(X_Test_TfIdf))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[131]: <matplotlib.axes._subplots.AxesSubplot at 0x23d37c3ee80>

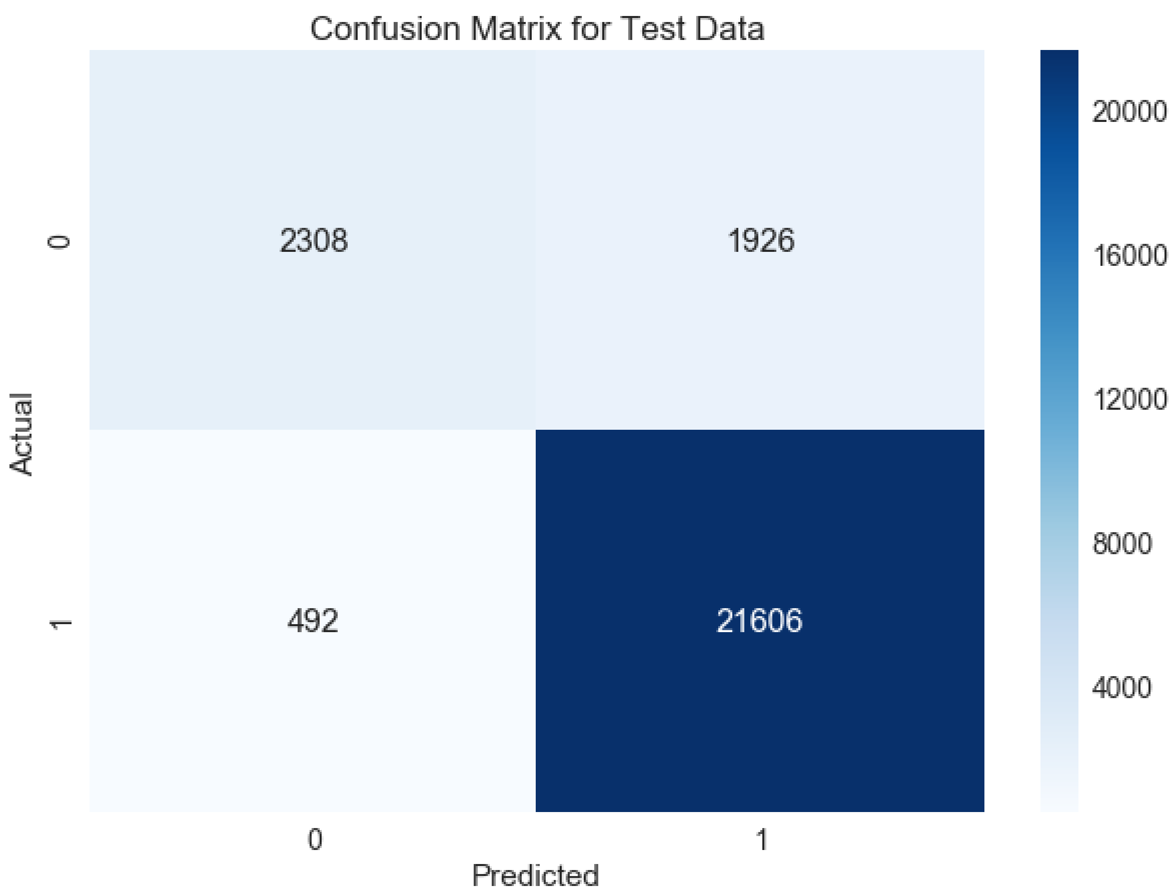

[4.2.6] Classification Report

In [132]:
```
from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

```
             precision    recall  f1-score   support

          0       0.82      0.55      0.66      4234
          1       0.92      0.98      0.95     22098

avg / total       0.90      0.91      0.90     26332
```

[4.2.7] Feature Importance

Feature Importance for Positive and Negetive Class

In [133]:
```
clf = SGDClassifier(alpha=0.1, class_weight='balanced')
clf.fit(X_Train_TfIdf, Y_Train)
plot_coefficients(clf, tf_idf_vect.get_feature_names())
```

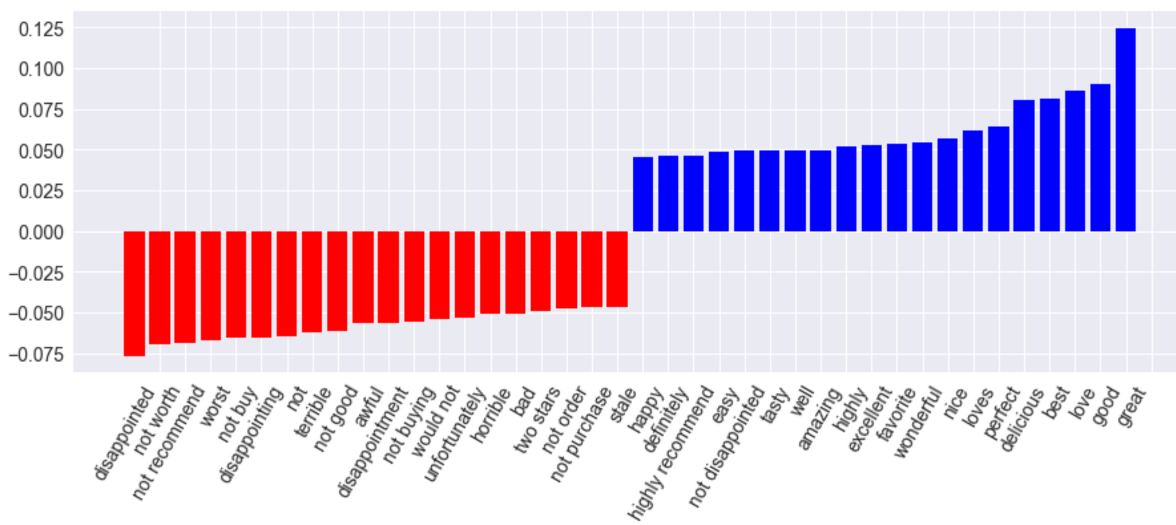

[4.3]Word2Vec

In [134]:
```
i=0
list_of_sentance_train=[]
for sentance in X_Train:
    list_of_sentance_train.append(sentance.split())

w2v_model=Word2Vec(list_of_sentance_train,min_count=5,size=50, workers=4)
w2v_words = list(w2v_model.wv.vocab)
print("number of words that occured minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])
```

```
number of words that occured minimum 5 times  14786
sample words  ['weekend', 'week', 'long', 'fast', 'using', 'rice', 'green',
'tea', 'works', 'wonders', 'one', 'energy', 'level', 'tasty', 'even', 'bit',
'salt', 'makes', 'much', 'pleasant', 'family', 'favorite', 'flavor', 'hanse
n', 'diet', 'sodas', 'clean', 'crisp', 'taste', 'enjoyable', 'meals', 'calm
s', 'upset', 'tummy', 'love', 'compared', 'ones', 'used', 'eat', 'like', 'nis
sin', 'maruchan', 'really', 'tell', 'difference', 'big', 'tub', 'spice', 'dro
ps', 'better']
```

[4.3.1] Computing avg w2v for train, test, and CV

In [135]:
```
%%time
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_train = []; # the avg-w2v for each sentence/review is stored in t
his list
for sent in tqdm(list_of_sentance_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might
need to change this to 300 if you use google's w2v
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_train.append(sent_vec)
sent_vectors_train = np.array(sent_vectors_train)
print(sent_vectors_train.shape)
print(sent_vectors_train[0])
```

```
100%|████████████████████████████████████████████████████████████████████████
███| 61441/61441 [03:46<00:00, 271.43it/s]

(61441, 50)
[-0.25580425  0.53885295 -0.79598692  0.37587689  0.0918958  -0.04543212
  0.13915103  0.19200098  0.26599771 -0.52361513 -0.02006662  0.32819876
  0.90590816 -0.05630648 -0.63506039  0.45599255 -0.07527782  0.0455145
 -0.48546116 -0.83688501  0.06410028 -0.52239371 -0.37882031  0.22823052
  0.48452098  0.22572852 -0.81385296  0.1630834   0.47067616 -0.02069531
 -0.20240183  0.1844469   0.54086916 -0.16328886  0.54146948 -0.01545888
 -0.34224481 -1.28831982  0.4634165  -0.43568516  0.31676882 -0.94641337
 -1.34183771 -0.6562375  -0.08998562  0.74762513  0.27968631  0.15942349
 -0.23538514 -0.68152205]
Wall time: 3min 46s
```

In [136]:
```
%%time
i=0
list_of_sentance_cv=[]
for sentance in X_cv:
    list_of_sentance_cv.append(sentance.split())
# average Word2Vec
# compute average word2vec for each review.
sent_vectors_cv = []; # the avg-w2v for each sentence/review is stored in this
list
for sent in tqdm(list_of_sentance_cv): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might
need to change this to 300 if you use google's w2v
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_cv.append(sent_vec)
sent_vectors_cv = np.array(sent_vectors_cv)
print(sent_vectors_cv.shape)
print(sent_vectors_cv[0])
```

```
100%|████████████████████████████████████████████████████████████████████████
███| 26332/26332 [01:48<00:00, 243.27it/s]

(26332, 50)
[-0.11573441  0.46256466 -1.50561294 -0.50534047  0.67303949 -0.03247621
  0.9521336  -1.07461442  0.59043062 -0.51276795  0.16654758  0.27134059
  0.38075584 -1.17755664 -0.36282212  0.80457097  0.41690589  0.79902533
 -0.41350211 -0.65742333  1.34620721 -0.2116666  -0.1831447  -0.21734889
 -0.09684955 -1.16451201  0.29688882 -0.48283021  1.6556564  -0.6978559
  1.13678383 -0.2415488   0.22276458 -0.93023639  0.3729921   0.26260851
 -0.32164465  0.67837226  0.28677943 -0.04293923  0.17995327  0.27606643
 -0.53558971  0.3690836  -0.25666084  1.2172374   0.93561432 -0.05929136
 -0.88247013 -0.71623034]
Wall time: 1min 48s
```

In [137]:
```
%%time

i=0
list_of_sentance_test=[]
for sentance in X_Test:
    list_of_sentance_test.append(sentance.split())

# average Word2Vec
# compute average word2vec for each review.
sent_vectors_test = []; # the avg-w2v for each sentence/review is stored in th
is list
for sent in tqdm(list_of_sentance_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might
need to change this to 300 if you use google's w2v
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors_test.append(sent_vec)
sent_vectors_test = np.array(sent_vectors_test)
print(sent_vectors_test.shape)
print(sent_vectors_test[0])
```

```
100%|████████████████████████████████████████████████████████████████████████
███| 26332/26332 [01:37<00:00, 270.13it/s]

(26332, 50)
[-0.11573441  0.46256466 -1.50561294 -0.50534047  0.67303949 -0.03247621
  0.9521336  -1.07461442  0.59043062 -0.51276795  0.16654758  0.27134059
  0.38075584 -1.17755664 -0.36282212  0.80457097  0.41690589  0.79902533
 -0.41350211 -0.65742333  1.34620721 -0.2116666  -0.1831447  -0.21734889
 -0.09684955 -1.16451201  0.29688882 -0.48283021  1.6556564  -0.6978559
  1.13678383 -0.2415488   0.22276458 -0.93023639  0.3729921   0.26260851
 -0.32164465  0.67837226  0.28677943 -0.04293923  0.17995327  0.27606643
 -0.53558971  0.3690836  -0.25666084  1.2172374   0.93561432 -0.05929136
 -0.88247013 -0.71623034]
Wall time: 1min 37s
```

[4.3.2] Hyperameter tuning with L1 Regulizer and AUC Plot

In [138]:
```
Optimal_Lamda_L1(sent_vectors_train, Y_Train, sent_vectors_cv,Y_cv)
```

```
100%|
| 9/9 [00:14<00:00,  1.62s/it]
```

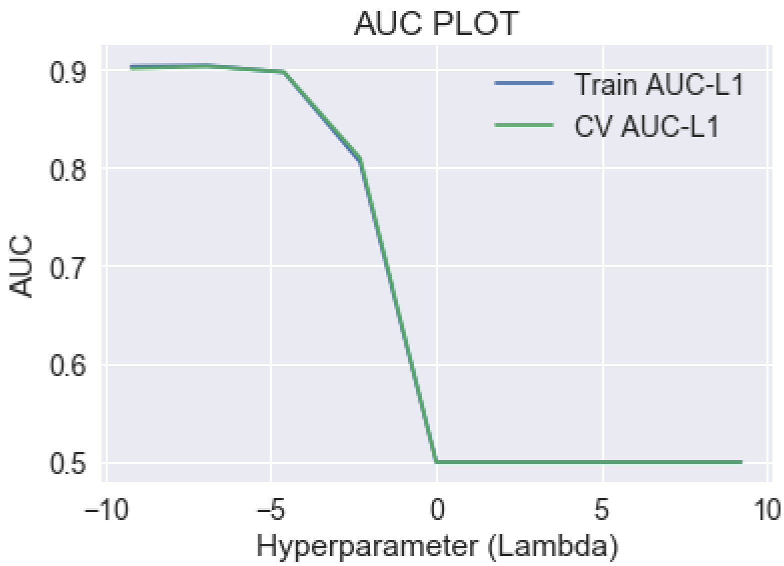


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L1
[0.9016006360296618, 0.9038676342464342, 0.8979925190886493, 0.80996416401315
85, 0.5, 0.5, 0.5, 0.5, 0.5]
Maximun AUC value : 0.9038676342464342
Index 1
```

[4.3.3] Hyperameter tuning with L2 Regulizer and AUC Plot

```
In [139]: Optimal_Lamda_L2(sent_vectors_train, Y_Train, sent_vectors_cv,Y_cv)
```

```
100%|
| 9/9 [00:08<00:00,  1.08it/s]
```

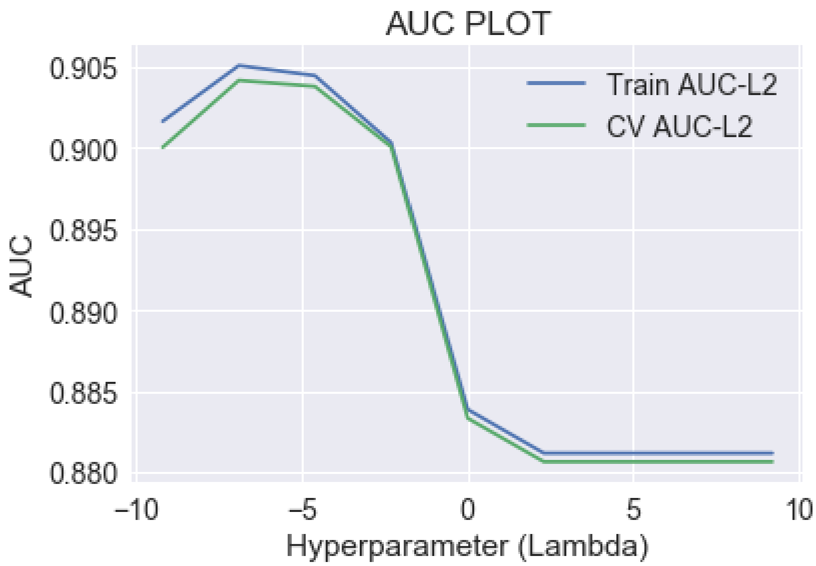


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L2
[0.9000430961270005, 0.9041814764847258, 0.9038064027322272, 0.90007921085670
98, 0.8833332521045835, 0.880647562434234, 0.8806469318426233, 0.880647081474
531, 0.880647081474531]
Maximun AUC value : 0.9041814764847258
Index 1
```

[4.3.4] ROC Curve of SVM

In [140]:
```
#Testing with test data
clf = SGDClassifier(alpha=0.1, class_weight='balanced', loss='hinge')
calibrated_clf = CalibratedClassifierCV(clf, cv=5)
calibrated_clf.fit(sent_vectors_train,Y_Train)
prediction = calibrated_clf.predict_proba(sent_vectors_test)[:,1]
print(prediction)
print(clf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, calibrated_clf.predict_proba(sent_vectors_train)[:,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, calibrated_clf.predict_proba(sent_vectors_test)[:,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of MNB')
plt.show()
```

```
[0.97856529 0.96440308 0.06950247 ... 0.61533381 0.90356187 0.07763442]
SGDClassifier(alpha=0.1, average=False, class_weight='balanced', epsilon=0.1,
       eta0=0.0, fit_intercept=True, l1_ratio=0.15,
       learning_rate='optimal', loss='hinge', max_iter=None, n_iter=None,
       n_jobs=1, penalty='l2', power_t=0.5, random_state=None,
       shuffle=True, tol=None, verbose=0, warm_start=False)
```

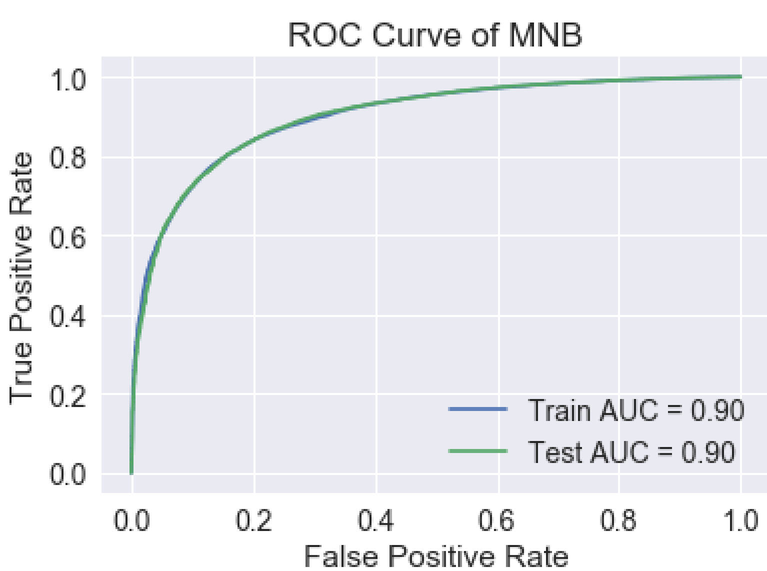

[4.3.5]Train and Test Accuracy

In [141]:
```
Training_Accuracy_w2v = calibrated_clf.score(sent_vectors_train, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_w2v)
Training_Error_w2v = 1 - Training_Accuracy_w2v
print('Training_Error=%0.3f'%Training_Error_w2v)

Test_Accuracy_w2v = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_w2v)
Test_Error_w2v = 1 - Test_Accuracy_w2v
print('Test_Error=%0.3f'%Test_Error_w2v)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))
```

```
Training_Accuracy=0.882
Training_Error=0.118
Test_Accuracy=0.884
Test_Error=0.116
```

[4.3.6]Confusion Matrix

In [142]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, calibrated_clf.predict(sent_vectors_tr
ain))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=n
p.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[142]: <matplotlib.axes._subplots.AxesSubplot at 0x23d2b08c8d0>

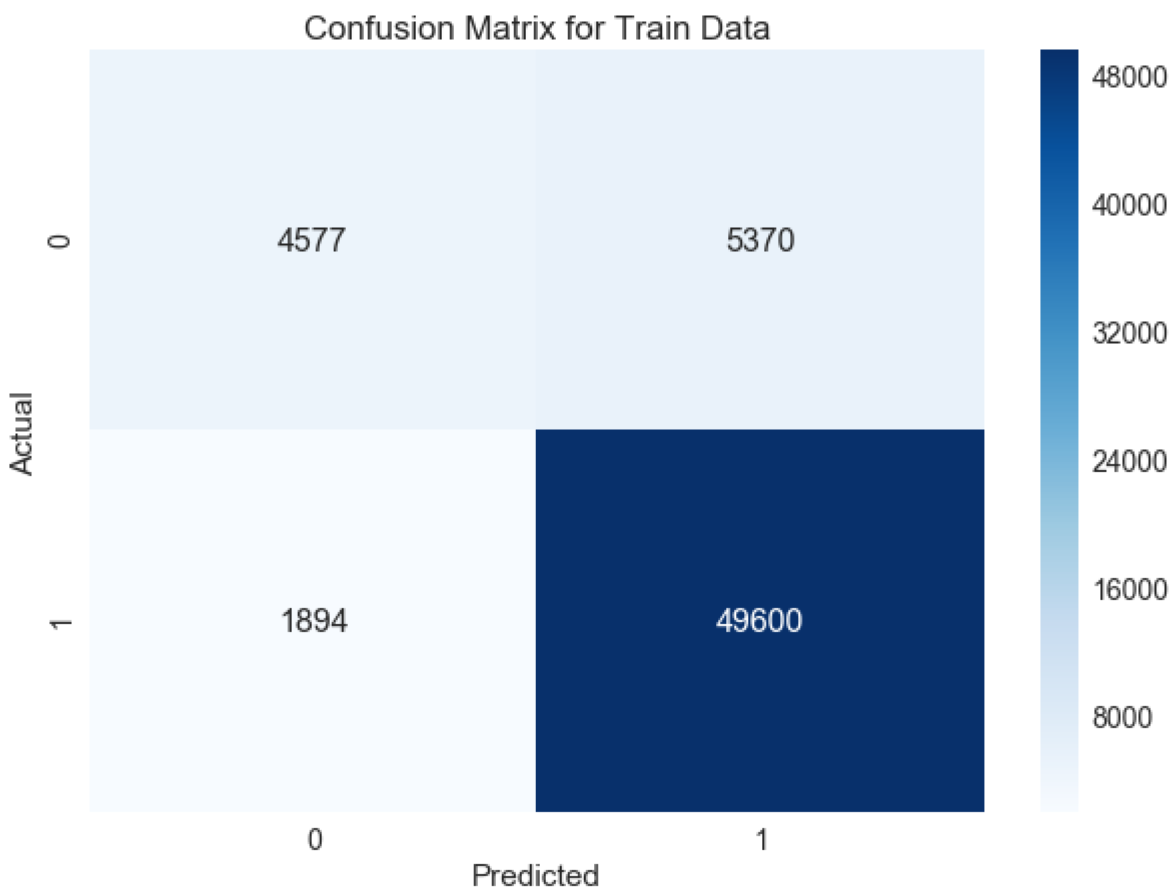

In [143]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, calibrated_clf.predict(sent_vectors_tes
t))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[143]: <matplotlib.axes._subplots.AxesSubplot at 0x23d18755f60>

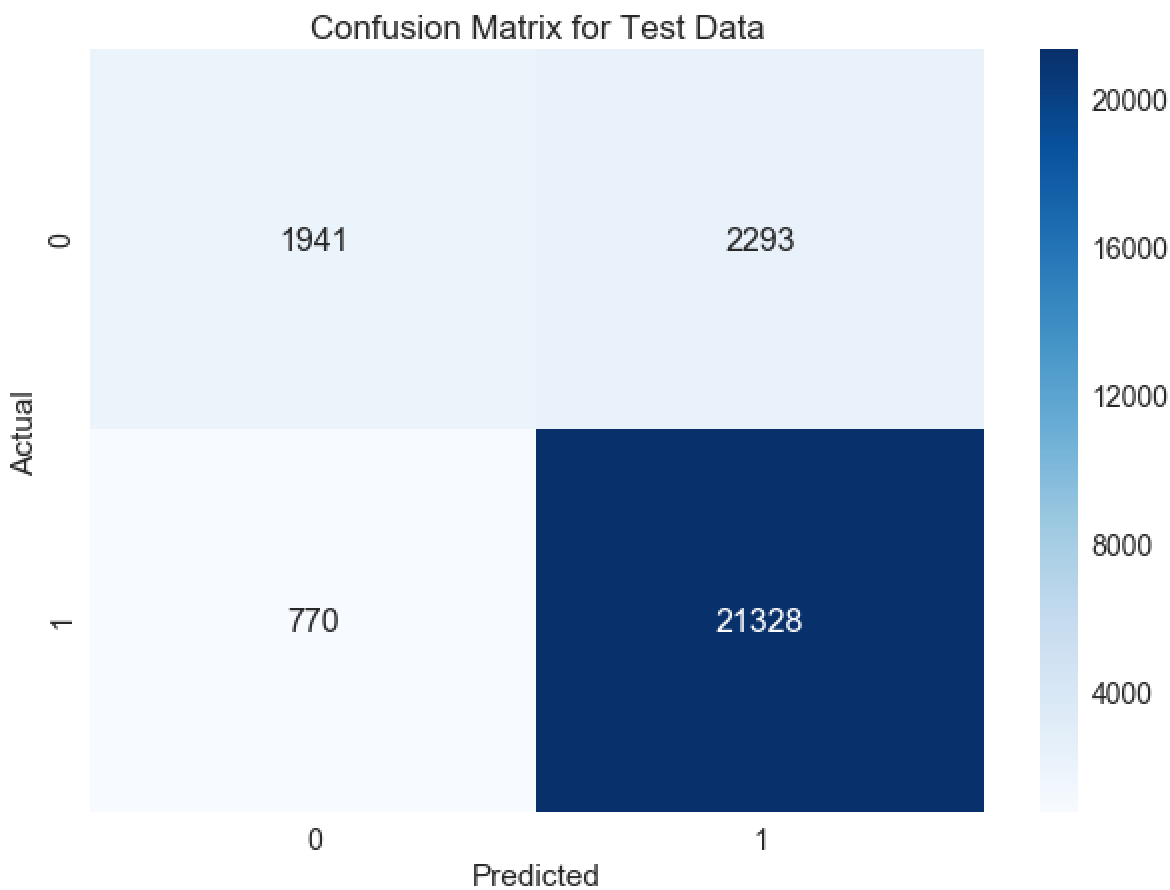

[4.3.7] Classification Report

In [144]:
```
from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

```
             precision    recall  f1-score   support

          0       0.72      0.46      0.56      4234
          1       0.90      0.97      0.93     22098

avg / total       0.87      0.88      0.87     26332
```

[4.2.7] Feature Importance

Feature Importance for Positive and Negetive Class

In [235]:
```
clf = SGDClassifier(alpha=0.1, class_weight='balanced')
clf.fit(X_Train_Bow, Y_Train)
plot_coefficients(clf, count_vect.get_feature_names())
```

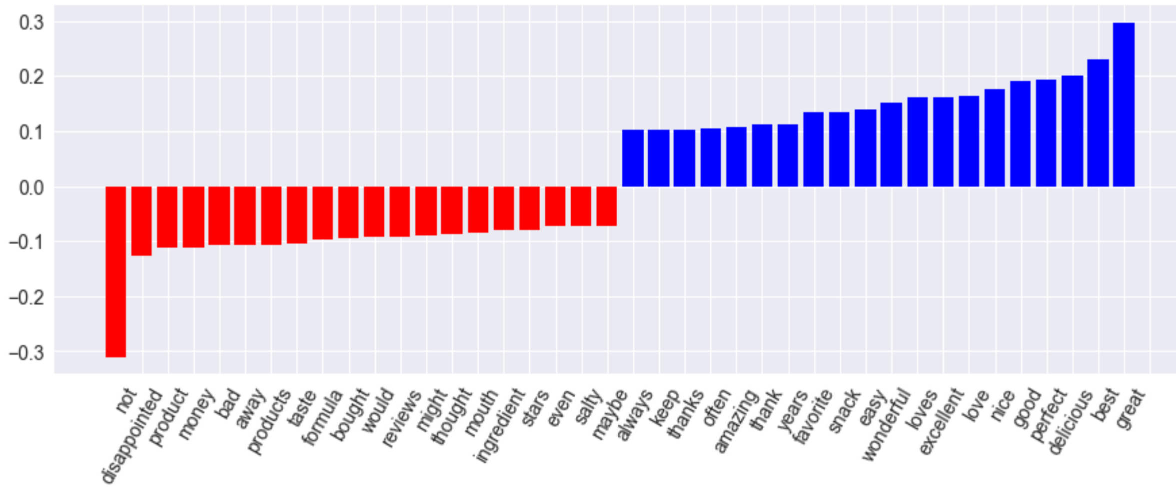

 [4.4] TFIDF weighted W2v

In [145]:
```
# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
model.fit(X_Train)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

[4.4.1] Compute TF-IDF weighted Word2Vec for Train, Test, and CV

In [146]:

```
i=0
list_of_sentance_train=[]
for sentance in X_Train:
    list_of_sentance_train.append(sentance.split())
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val
 = tfidf

tfidf_sent_vectors_train = []; # the tfidf-w2v for each sentence/review is sto
red in this list
row=0;
for sent in tqdm(list_of_sentance_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_train.append(sent_vec)
    row += 1
```

```
100%|████████████████████████████████████████████████████████████████████████
████| 61441/61441 [53:42<00:00, 19.07it/s]
```

In [147]:
```
i=0
list_of_sentance_test=[]
for sentance in X_Test:
    list_of_sentance_test.append(sentance.split())
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val
 = tfidf

tfidf_sent_vectors_test = []; # the tfidf-w2v for each sentence/review is stor
ed in this list
row=0;
for sent in tqdm(list_of_sentance_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_test.append(sent_vec)
    row += 1
```

```
100%|████████████████████████████████████████████████████████████████████
████| 26332/26332 [20:23<00:00, 17.00it/s]
```

In [148]:
```
i=0
list_of_sentance_cv=[]
for sentance in X_cv:
    list_of_sentance_cv.append(sentance.split())
# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val
 = tfidf

tfidf_sent_vectors_cv = []; # the tfidf-w2v for each sentence/review is stored
in this list
row=0;
for sent in tqdm(list_of_sentance_cv): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors_cv.append(sent_vec)
    row += 1
```

```
100%|██████████████████████████████████████████████████████████████████████
████| 26332/26332 [23:33<00:00, 18.63it/s]
```

[4.4.2] Hyperameter tuning with L1 Regulizer and AUC Plot

In [150]:
```
Optimal_Lamda_L1(tfidf_sent_vectors_train, Y_Train, tfidf_sent_vectors_cv,Y_cv
)
```

```
100%|████████████████████████████████████████| 9/9 [00:42<00:00,  3.80s/it]
```

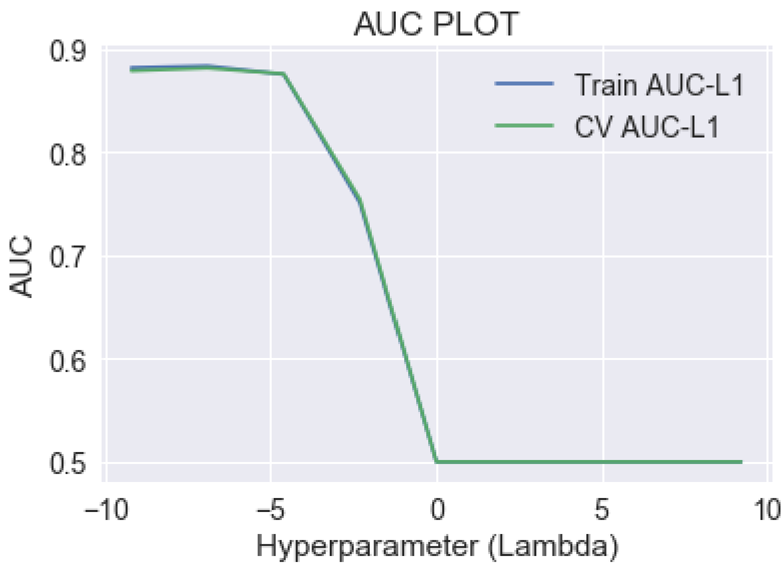


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L1
[0.8787021445629771, 0.8813609004899504, 0.8756963815541822, 0.75417574558266
31, 0.5, 0.5, 0.5, 0.5, 0.5]
Maximun AUC value : 0.8813609004899504
Index 1
```

[4.4.3] Hyperameter tuning with L2 Regulizer and AUC Plot

In [151]:
```
Optimal_Lamda_L2(tfidf_sent_vectors_train, Y_Train, tfidf_sent_vectors_cv,Y_cv
)
```

```
100%|██████████████████████████████████████████████████████████████████████████| 9/9 [00:10<00:00,  1.12s/it]
```

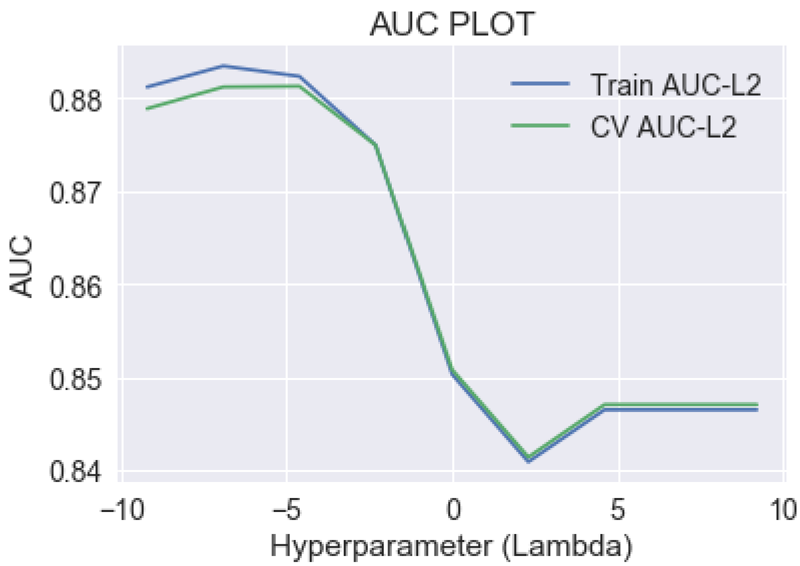


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L2
[0.8788924763495014, 0.8812327086970725, 0.8813056756280361, 0.87495876037745
37, 0.8508129266406487, 0.8414037623361355, 0.8470622959955979, 0.84706205017
17496, 0.8470620501717496]
Maximun AUC value : 0.8813056756280361
Index 2
```

[4.4.4] ROC Curve of SVM

In [158]:
```
#Testing with test data
clf = SGDClassifier(alpha=0.1, class_weight='balanced', loss='hinge')
calibrated_clf = CalibratedClassifierCV(clf, cv=5)
calibrated_clf.fit(tfidf_sent_vectors_train,Y_Train)
prediction = calibrated_clf.predict_proba(tfidf_sent_vectors_test)[:,1]
print(prediction)
print(clf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, calibrated_clf.predict_pr
oba(tfidf_sent_vectors_train)[:,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, calibrated_clf.predict_proba
(tfidf_sent_vectors_test)[:,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of MNB')
plt.show()
```

```
[0.93023877 0.9403652  0.10374825 ... 0.69068035 0.8233184  0.03266469]
SGDClassifier(alpha=0.1, average=False, class_weight='balanced', epsilon=0.1,
       eta0=0.0, fit_intercept=True, l1_ratio=0.15,
       learning_rate='optimal', loss='hinge', max_iter=None, n_iter=None,
       n_jobs=1, penalty='l2', power_t=0.5, random_state=None,
       shuffle=True, tol=None, verbose=0, warm_start=False)
```

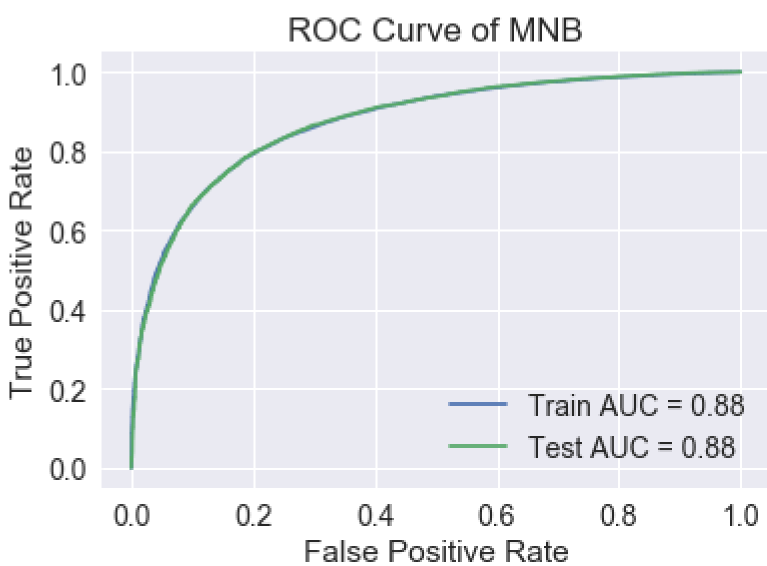

[4.4.5]Train and Test Accuracy

In [159]:
```
Training_Accuracy_tfidfw2v = calibrated_clf.score(tfidf_sent_vectors_train, Y_
Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_tfidfw2v)
Training_Error_tfidfw2v = 1 - Training_Accuracy_tfidfw2v
print('Training_Error=%0.3f'%Training_Error_tfidfw2v)

Test_Accuracy_tfidfw2v = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_tfidfw2v)
Test_Error_tfidfw2v = 1 - Test_Accuracy_tfidfw2v
print('Test_Error=%0.3f'%Test_Error_tfidfw2v)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))
```

```
Training_Accuracy=0.870
Training_Error=0.130
Test_Accuracy=0.871
Test_Error=0.129
```

[4.4.6]Confusion Matrix

In [160]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, calibrated_clf.predict(tfidf_sent_vect
ors_train))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=n
p.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[160]: <matplotlib.axes._subplots.AxesSubplot at 0x23d56f9ab00>

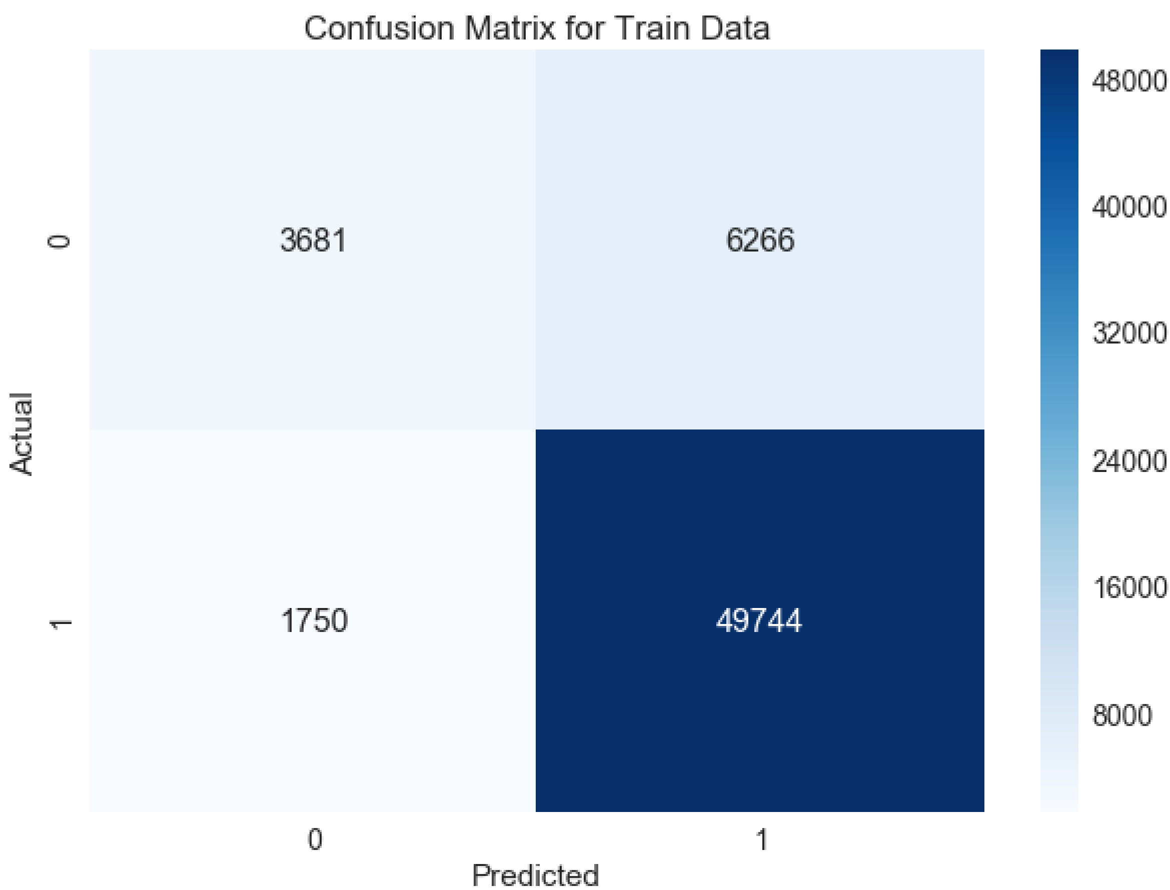

In [161]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, calibrated_clf.predict(tfidf_sent_vecto
rs_test))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[161]:
```
<matplotlib.axes._subplots.AxesSubplot at 0x23d3297d470>
```

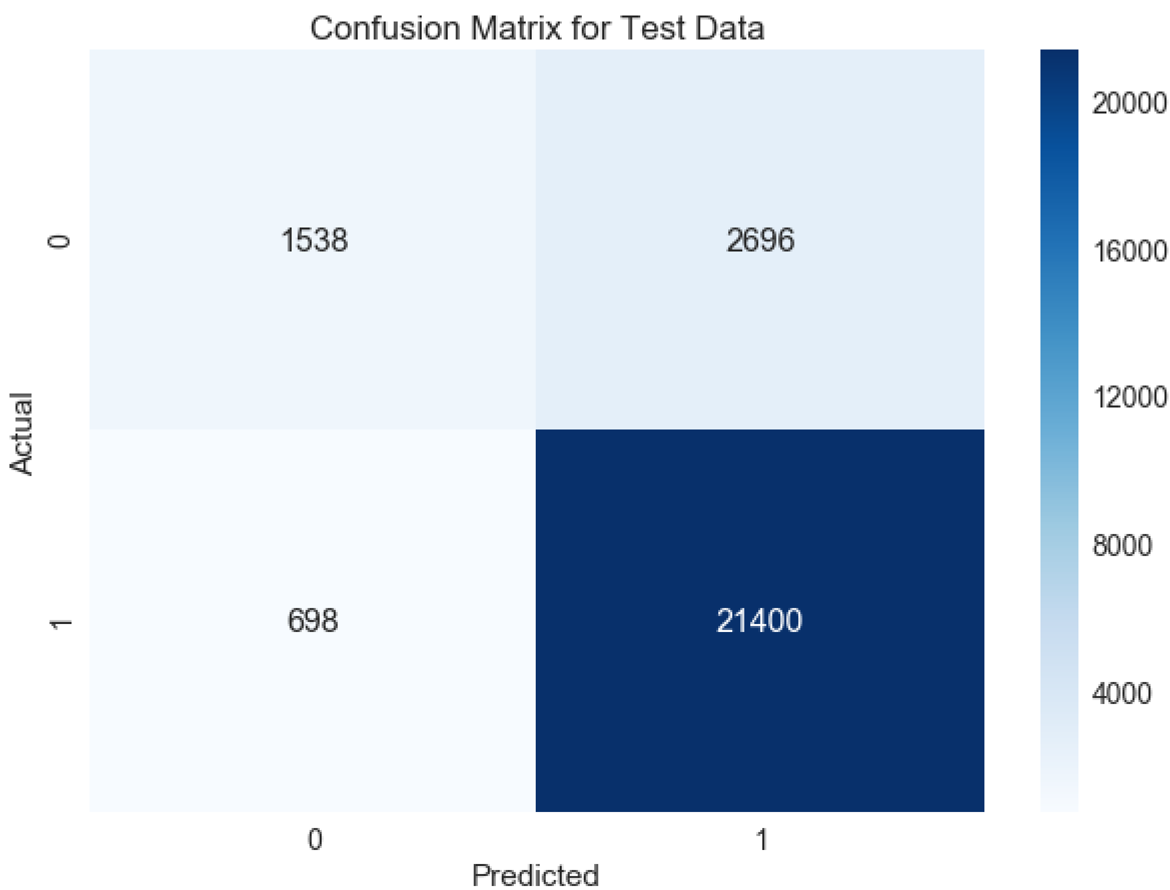

[4.4.7] Classification Report

In [162]:
```
from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

```
             precision    recall  f1-score   support

          0       0.69      0.36      0.48      4234
          1       0.89      0.97      0.93     22098

avg / total       0.86      0.87      0.85     26332
```

[4.4.8] Feature Importance

Feature Importance for Positive and Negetive Class

In [163]:
```
clf = SGDClassifier(alpha=0.1, class_weight='balanced')
clf.fit(tfidf_sent_vectors_train, Y_Train)
plot_coefficients(clf, model.get_feature_names())
```

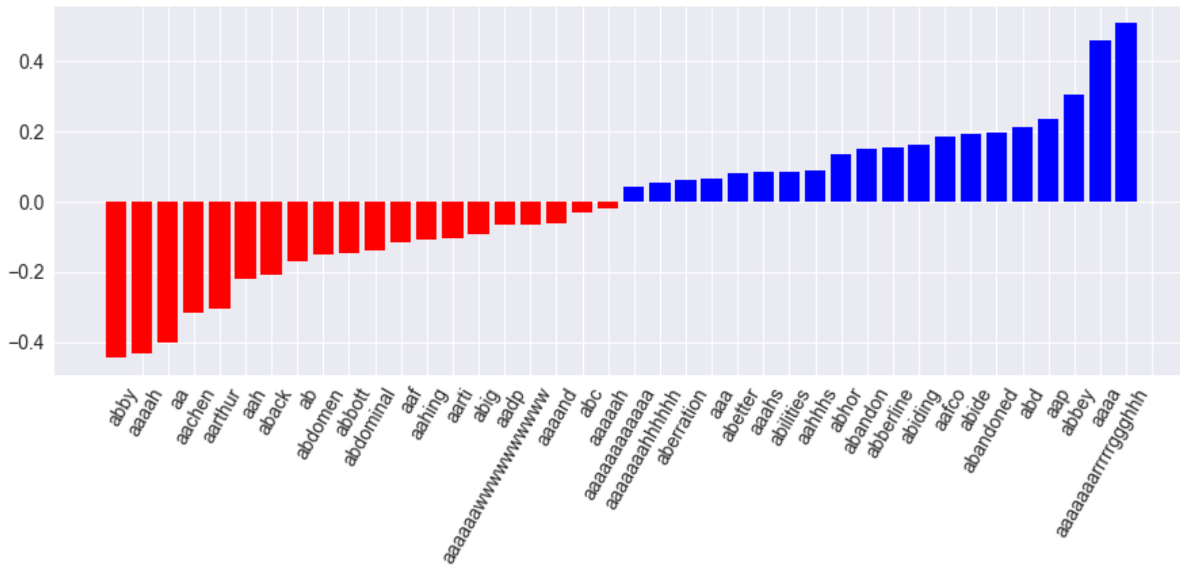

[5] RBF Kernel

In [171]:
```
Y = final['Score'].values
X = final['cleaned_text'].values

X=X[:20000]
Y=Y[:20000]

print(Y.shape)
print(type(Y))
print(X.shape)
print(type(X))
```

```
(20000,)
<class 'numpy.ndarray'>
(20000,)
<class 'numpy.ndarray'>
```

In [173]:
```
# split the data set into train and test
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X,Y,test_size=0.3, random_
state=12, shuffle = False)

# split the train data set into cross validation train and cross validation te
st
X_tr, X_cv, Y_tr, Y_cv = train_test_split(X,Y, test_size=0.3, random_state=12,
shuffle = False)
```

In [174]:
```
print('='*100)
print("After splitting")
print("X_Train Shape:",X_Train.shape, "Y_Train Shape:",Y_Train.shape)
print("X_cv Shape:",X_cv.shape,       "Y_cv Shape",Y_cv.shape)
print("X_Test Shape",X_Test.shape,    "Y_Test Shape",Y_Test.shape)
```

```
==============================================================================
======================
After splitting
X_Train Shape: (14000,) Y_Train Shape: (14000,)
X_cv Shape: (6000,) Y_cv Shape (6000,)
X_Test Shape (6000,) Y_Test Shape (6000,)
```

[5] Featurization

[5.1] BAG OF WORDS

In [238]:
```
import math
def Optimal_Lamda(X_Train,Y_Train,X_CV,Y_CV):
    train_AUC = []
    CV_AUC = []
    tuned_parameters=[10**-4, 10**-3, 10**-2, 10**-1, 1,10**1, 10**2, 10**3, 1
0**4]
    for j in  tqdm(tuned_parameters):
        clf = SVC(C=j, probability=True)
        clf.fit(X_Train, Y_Train)
        y_train_pred =  clf.predict_proba(X_Train)[:,1]
        y_cv_pred =  clf.predict_proba(X_CV)[:,1]
        train_AUC.append(roc_auc_score(Y_Train,y_train_pred))
        CV_AUC.append(roc_auc_score(Y_CV, y_cv_pred))

    #Error plots with penaly L1
    plt.plot(np.log(tuned_parameters), train_AUC, label='Train AUC')
    plt.plot(np.log(tuned_parameters), CV_AUC, label='CV AUC')
    plt.legend()
    plt.xlabel("Hyperparameter (Lambda)")
    plt.ylabel("AUC")
    plt.title("AUC PLOT")
    plt.show()

    #Cv auc scores
    print("CV AUS Scores with Penalty=? Cv auc scores")
    print(CV_AUC)
    print("Maximun AUC value :",max(CV_AUC))
    print("Index",CV_AUC.index(max(CV_AUC)))
```

[5.1.2] Hyperameter tuning and AUC Plot

In [239]:
```
Optimal_Lamda(X_Train_Bow, Y_Train, X_CV_Bow,Y_cv)
```

```
100%|██████████████████████████████████████████████████
████████| 9/9 [24:16<00:00, 164.61s/it]
```

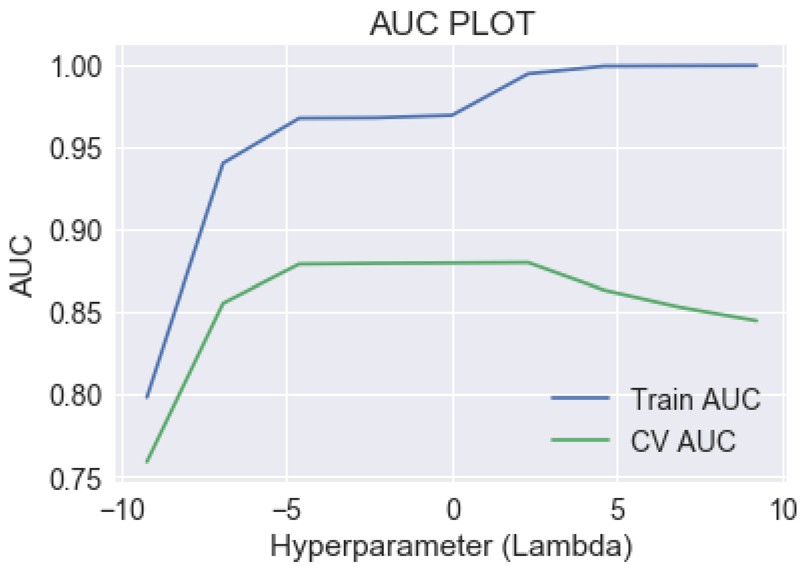


```
CV AUS Scores with Penalty=? Cv auc scores
[0.7589056933521543, 0.8550491810624903, 0.878669198539132, 0.879082588335461
9, 0.8793186350294638, 0.8796249243799152, 0.8629734937599427, 0.852581613676
5925, 0.8445151352199145]
Maximun AUC value : 0.8796249243799152
Index 5
```

[5.1.4] ROC Curve of SVM

In [179]:
```
#Testing with test data
clf = SVC(C=1, probability=True)
clf.fit(X_Train_Bow,Y_Train)
prediction = clf.predict_proba(X_Test_Bow)[:,1]
print(prediction)
print(clf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, clf.predict_proba(X_Train
_Bow)[:,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, clf.predict_proba(X_Test_Bow
)[:,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of MNB')
plt.show()
```

```
[0.10968408 0.84796084 0.99999183 ... 0.97645698 0.89394841 0.97582892]
SVC(C=1, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
  max_iter=-1, probability=True, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

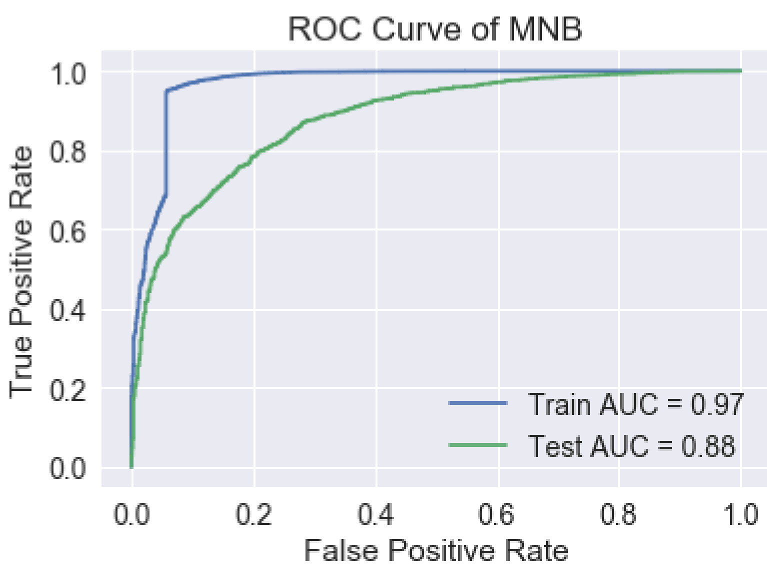

[5.1.5]Train and Test Accuracy

In [180]:
```
Training_Accuracy_Bow = clf.score(X_Train_Bow, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Bow)
Training_Error_Bow = 1 - Training_Accuracy_Bow
print('Training_Error=%0.3f'%Training_Error_Bow)

Test_Accuracy_Bow = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Bow)
Test_Error_Bow = 1 - Test_Accuracy_Bow
print('Test_Error=%0.3f'%Test_Error_Bow)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))
```

```
Training_Accuracy=0.940
Training_Error=0.060
Test_Accuracy=0.888
Test_Error=0.112
```

[5.1.6]Confusion Matrix

In [182]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, clf.predict(X_Train_Bow))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=n
p.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[182]: <matplotlib.axes._subplots.AxesSubplot at 0x23d3280ea90>

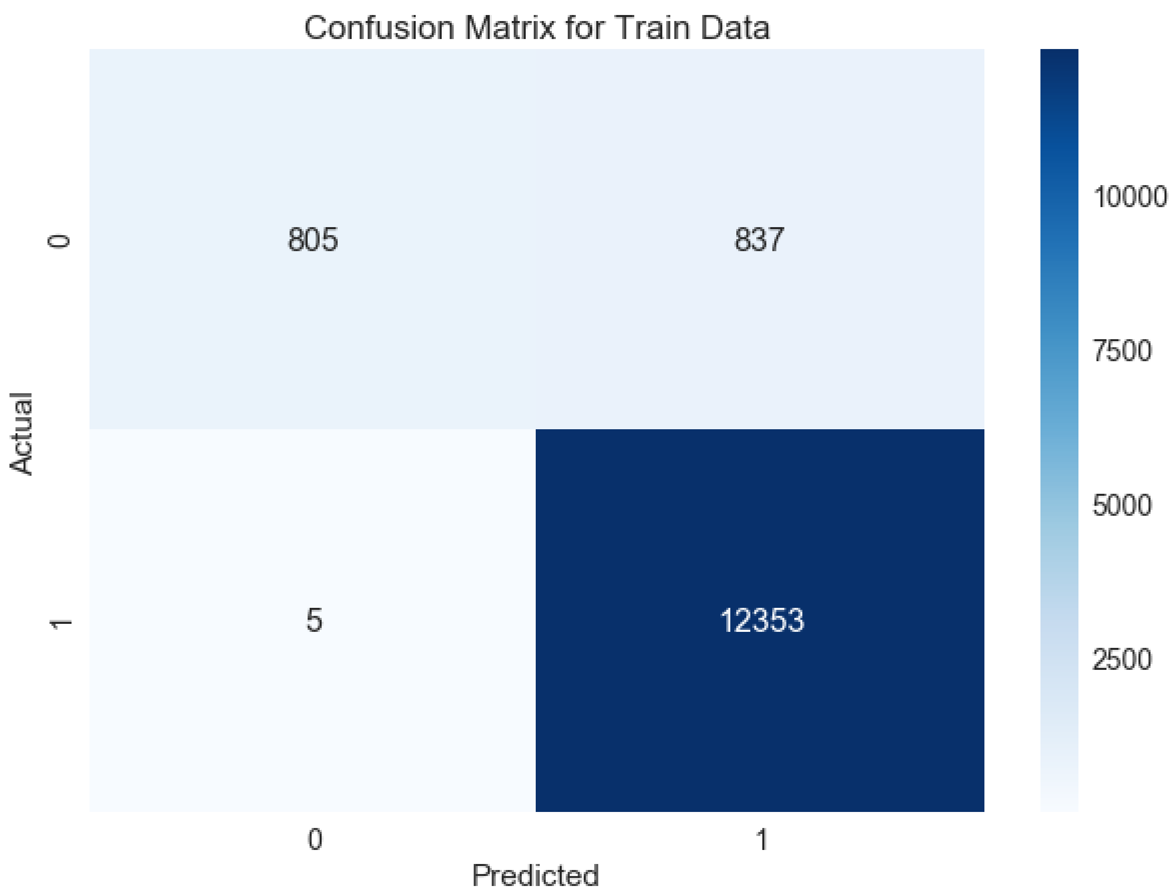

In [183]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, clf.predict(X_Test_Bow))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[183]: <matplotlib.axes._subplots.AxesSubplot at 0x23d18114668>

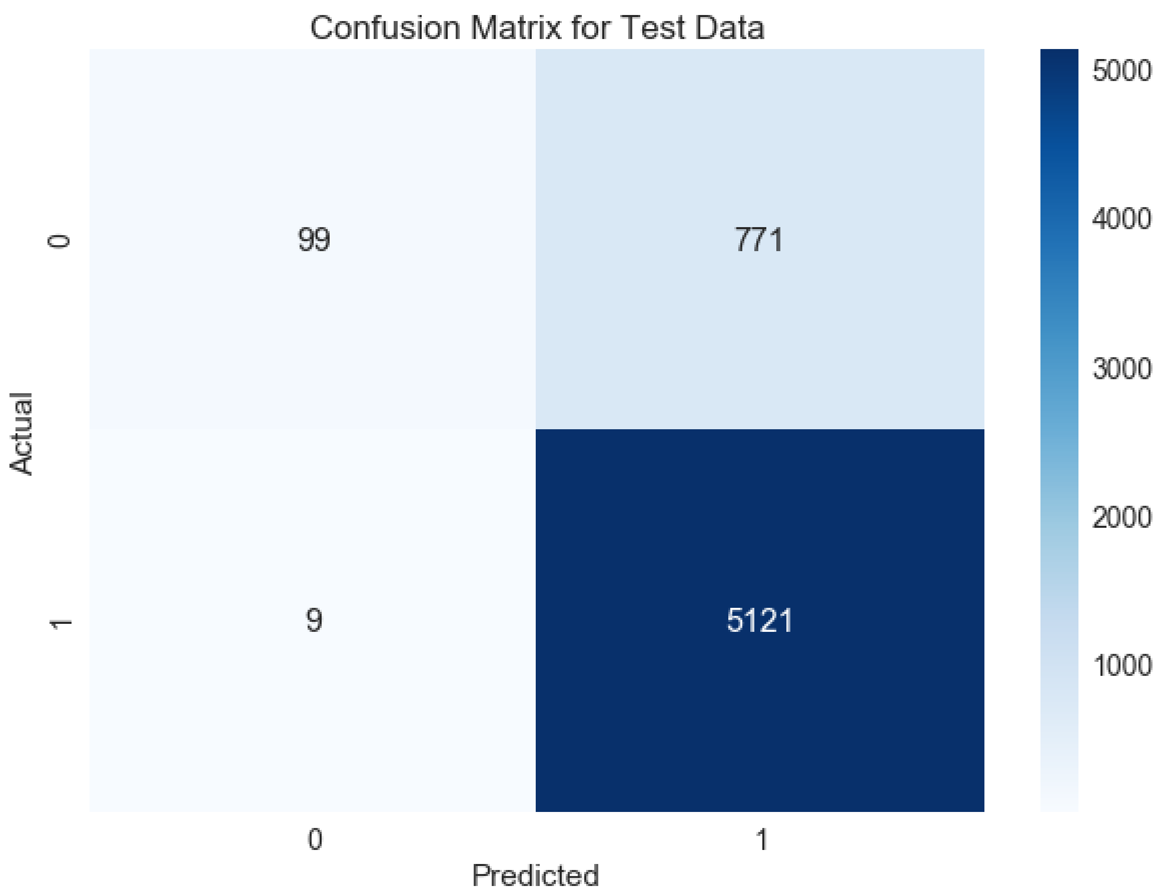

[5.1.7] Classification Report

In [184]:
```
from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

```
             precision    recall  f1-score   support

          0       0.75      0.33      0.46       870
          1       0.90      0.98      0.94      5130

avg / total       0.88      0.89      0.87      6000
```

[5.1.8] Feature Importance

Feature Importance for Positive and Negetive Class

In [234]:
```
clf = SVC(C=1, class_weight='balanced', kernel = 'linear')
clf.fit(X_Train_Bow, Y_Train)
plot_coefficients(clf, count_vect.get_feature_names())
```

```
---------------------------------------------------------------------------
AttributeError                            Traceback (most recent call last)
<ipython-input-234-145e40f58dcd> in <module>()
      1 clf = SVC(C=1, class_weight='balanced', kernel = 'linear')
      2 clf.fit(X_Train_Bow, Y_Train)
----> 3 plot_coefficients(clf, count_vect.get_feature_names())

<ipython-input-119-7beddd17e766> in plot_coefficients(classifier, feature_nam
es, top_features)
      1 #https://medium.com/@aneesha/visualising-top-features-in-linear-svm-w
ith-scikit-learn-and-matplotlib-3454ab18a14d
      2 def plot_coefficients(classifier, feature_names, top_features=20):
----> 3     coef = classifier.coef_.ravel()
      4     top_positive_coefficients = np.argsort(coef)[-top_features:]
      5     top_negative_coefficients = np.argsort(coef)[:top_features]

D:\Anaconda3\lib\site-packages\scipy\sparse\base.py in __getattr__(self, att
r)
    684             return self.getnnz()
    685         else:
--> 686             raise AttributeError(attr + " not found")
    687 
    688     def transpose(self, axes=None, copy=False):

AttributeError: ravel not found
```

[5.2] TF-IDF

[5.2.1] Hyperameter tuning with L1 Regulizer and AUC Plot

In [245]:
```
Optimal_Lamda(X_Train_TfIdf, Y_Train, X_CV_TfIdf,Y_cv)
```

```
100%|
| 9/9 [38:03<00:00, 338.87s/it]
```

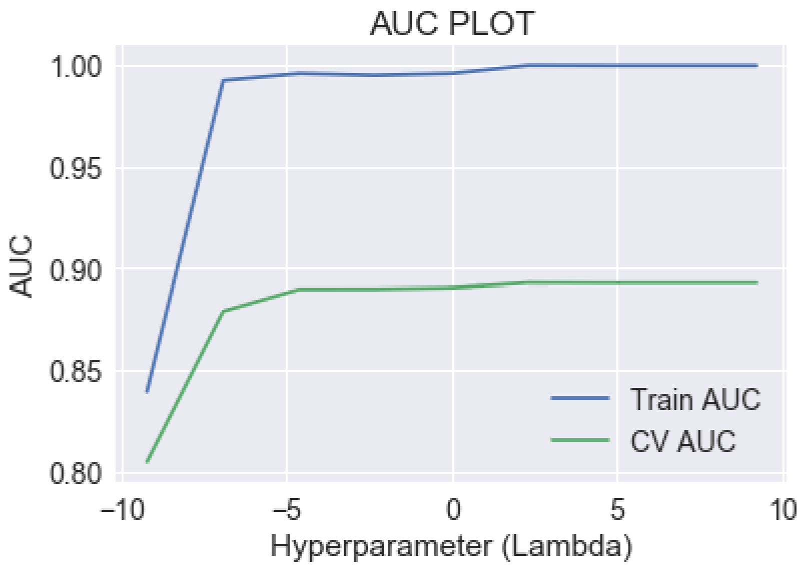


```
CV AUS Scores with Penalty=? Cv auc scores
[0.8045293629988126, 0.8787419058501938, 0.8894954179830162, 0.88952387354081
24, 0.8901895543456343, 0.8926871457058994, 0.8925605520826331, 0.89256615357
03883, 0.8925666016894086]
Maximun AUC value : 0.8926871457058994
Index 5
```

[5.2.2] ROC Curve of SVM

In [193]:
```
#Testing with test data
clf = SVC(C=1, probability=True)
clf.fit(X_Train_TfIdf,Y_Train)
prediction = clf.predict_proba(X_Test_TfIdf)[:,1]
print(prediction)
print(clf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, clf.predict_proba(X_Train
_TfIdf)[:,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, clf.predict_proba(X_Test_TfI
df)[:,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of MNB')
plt.show()
```

```
[0.18960522 0.89269833 0.99677773 ... 0.9488998  0.87401    0.96948229]
SVC(C=1, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
  max_iter=-1, probability=True, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

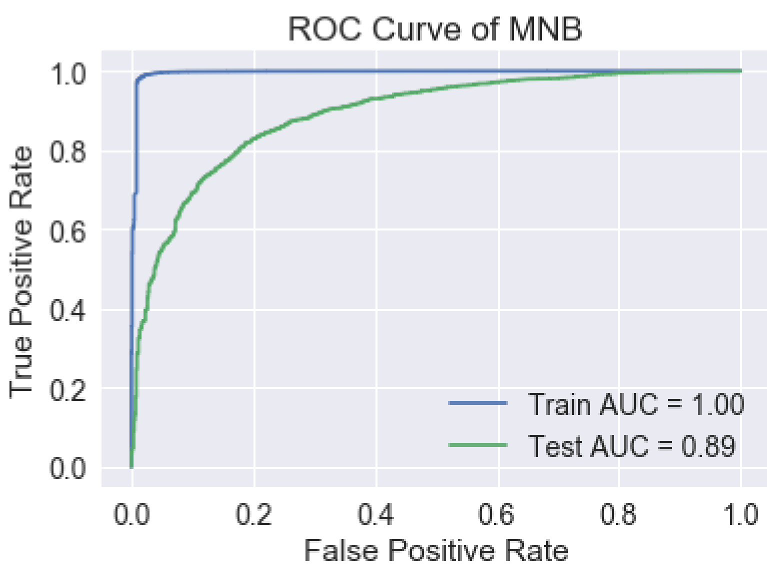

[5.2.3]Train and Test Accuracy

In [194]:
```
Training_Accuracy_Tfidf = clf.score(X_Train_TfIdf, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_Tfidf)
Training_Error_Tfidf = 1 - Training_Accuracy_Tfidf
print('Training_Error=%0.3f'%Training_Error_Tfidf)

Test_Accuracy_Tfidf = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_Tfidf)
Test_Error_Tfidf = 1 - Test_Accuracy_Tfidf
print('Test_Error=%0.3f'%Test_Error_Tfidf)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))
```

```
Training_Accuracy=0.952
Training_Error=0.048
Test_Accuracy=0.888
Test_Error=0.112
```

[5.2.4] Confusion Matrix

In [195]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, clf.predict(X_Train_TfIdf))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=n
p.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[195]: <matplotlib.axes._subplots.AxesSubplot at 0x23d18a22be0>

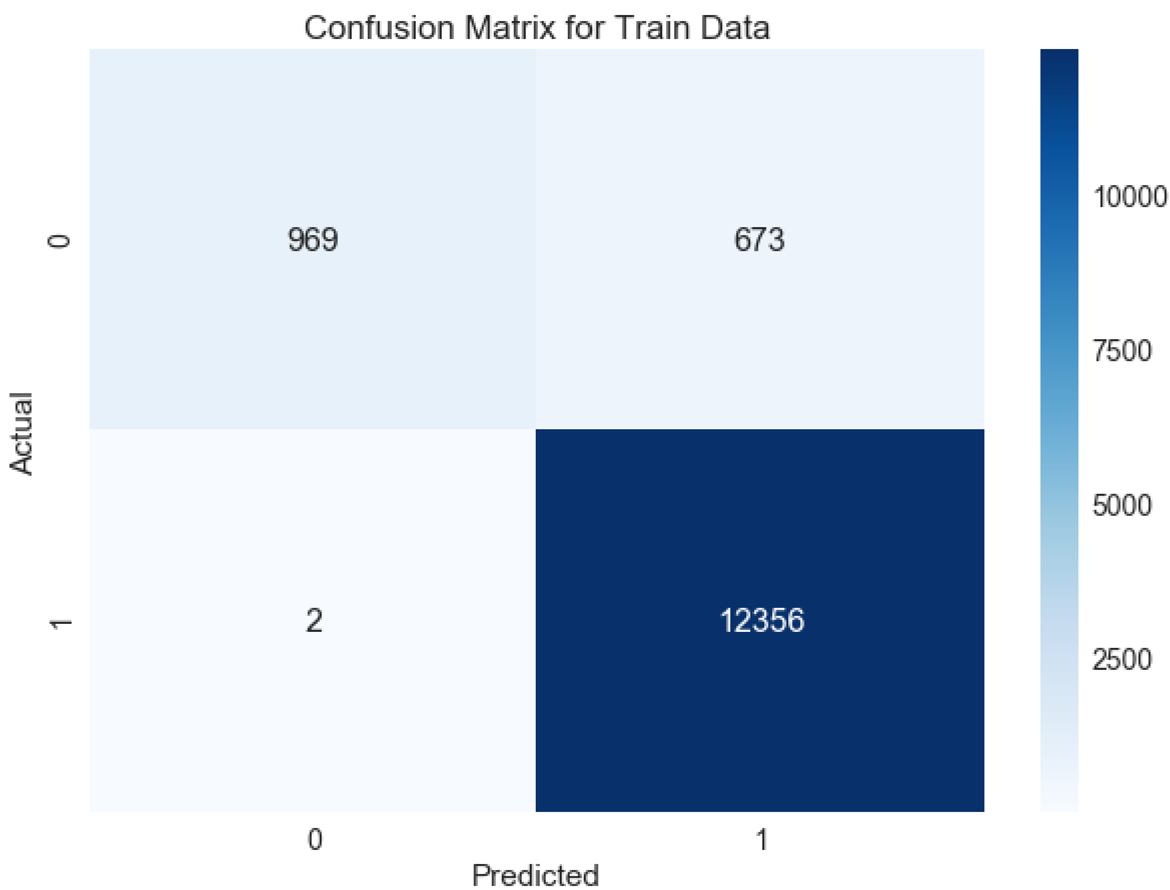

In [196]:
```
#With the reference of below link:
#https://www.kaggle.com/agungor2/various-confusion-matrix-plots
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, clf.predict(X_Test_TfIdf))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[196]: <matplotlib.axes._subplots.AxesSubplot at 0x23d10368978>

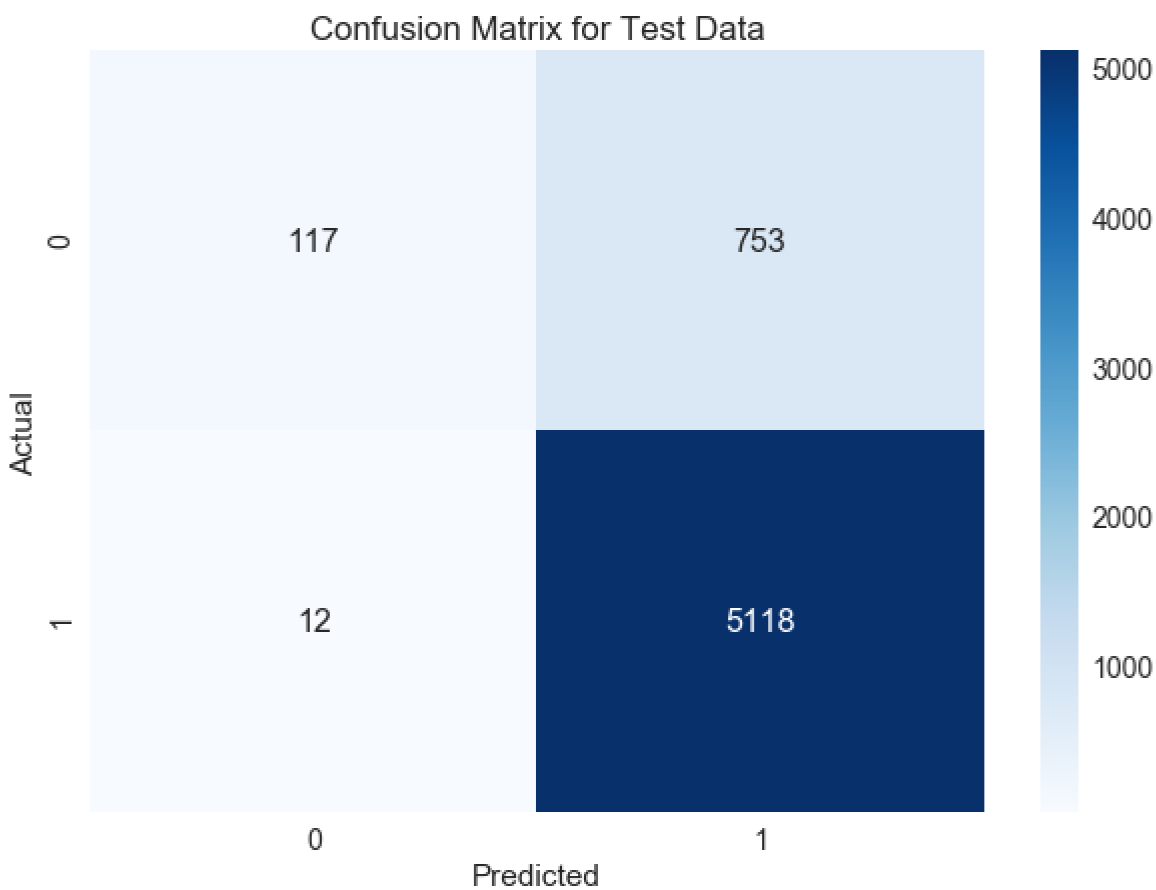

[5.2.5] Classification Report

In [197]:
```
from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

```
             precision    recall  f1-score   support

          0       0.73      0.36      0.48       870
          1       0.90      0.98      0.94      5130

avg / total       0.88      0.89      0.87      6000
```

[5.2.6] Feature Importance

Feature Importance for Positive and Negetive Class

In [233]:
```
clf = SVC(C=1, class_weight='balanced', kernel='linear')
clf.fit(X_Train_TfIdf, Y_Train)
plot_coefficients(clf, tf_idf_vect.get_feature_names())
```

```
---------------------------------------------------------------------------
AttributeError                            Traceback (most recent call last)
<ipython-input-233-2a1dd6b34b71> in <module>()
      1 clf = SVC(C=1, class_weight='balanced', kernel='linear')
      2 clf.fit(X_Train_TfIdf, Y_Train)
----> 3 plot_coefficients(clf, tf_idf_vect.get_feature_names())

<ipython-input-119-7beddd17e766> in plot_coefficients(classifier, feature_nam
es, top_features)
      1 #https://medium.com/@aneesha/visualising-top-features-in-linear-svm-w
ith-scikit-learn-and-matplotlib-3454ab18a14d
      2 def plot_coefficients(classifier, feature_names, top_features=20):
----> 3     coef = classifier.coef_.ravel()
      4     top_positive_coefficients = np.argsort(coef)[-top_features:]
      5     top_negative_coefficients = np.argsort(coef)[:top_features]

D:\Anaconda3\lib\site-packages\scipy\sparse\base.py in __getattr__(self, att
r)
    684             return self.getnnz()
    685         else:
--> 686             raise AttributeError(attr + " not found")
    687 
    688     def transpose(self, axes=None, copy=False):

AttributeError: ravel not found
```

[5.3]Word2Vec

[5.3.1] Hyperameter tuning and AUC Plot

In [246]:
```
Optimal_Lamda(sent_vectors_train, Y_Train, sent_vectors_cv,Y_cv)
```

```
100%|██████████████████████████████████████████████████| 9/9 [20:57<00:00, 271.94s/it]
```

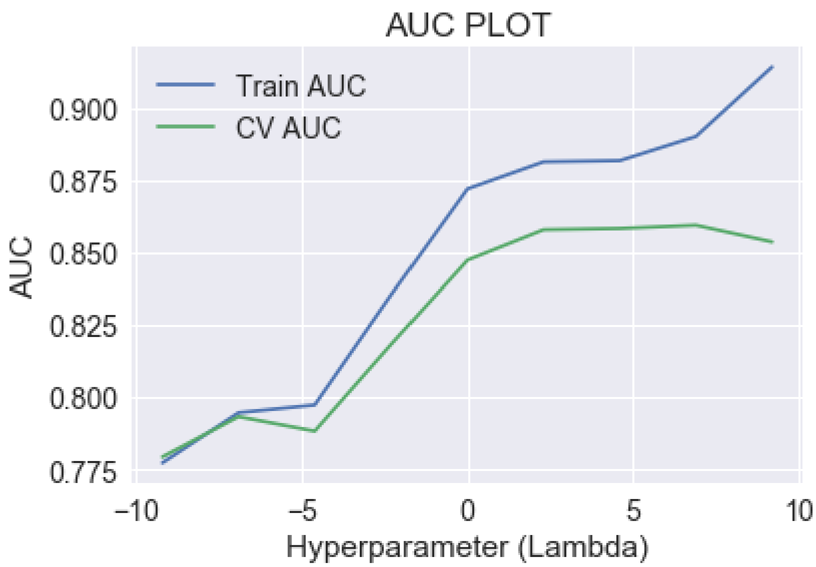


```
CV AUS Scores with Penalty=? Cv auc scores
[0.779341713159015, 0.7932784387533328, 0.7882664515695369, 0.818254016266720
6, 0.8474486343572855, 0.8579478389460241, 0.8583690708252112, 0.859500571351
7511, 0.8537265577737446]
Maximun AUC value : 0.8595005713517511
Index 7
```

[5.3.2] ROC Curve of SVM

In [211]:
```
#Testing with test data
clf = SVC(C=1, probability=True)
clf.fit(sent_vectors_train,Y_Train)
prediction = clf.predict_proba(sent_vectors_test)[:,1]
print(prediction)
print(clf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, clf.predict_proba(sent_ve
ctors_train)[:,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, clf.predict_proba(sent_vecto
rs_test)[:,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of MNB')
plt.show()
```

```
[0.65462583 0.64586442 0.9698258  ... 0.89672998 0.91763891 0.94664735]
SVC(C=1, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
  max_iter=-1, probability=True, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

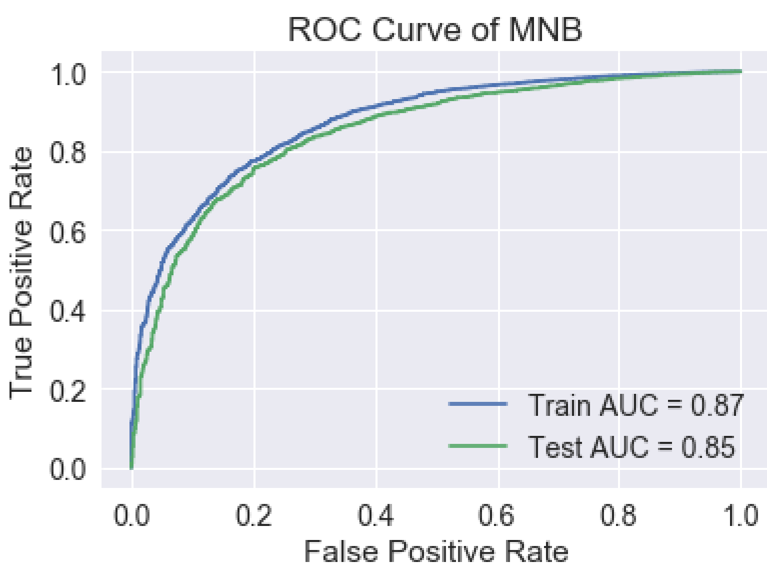

[5.3.4]Train and Test Accuracy

In [212]:
```
Training_Accuracy_w2v = clf.score(sent_vectors_train, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_w2v)
Training_Error_w2v = 1 - Training_Accuracy_w2v
print('Training_Error=%0.3f'%Training_Error_w2v)

Test_Accuracy_w2v = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_w2v)
Test_Error_w2v = 1 - Test_Accuracy_w2v
print('Test_Error=%0.3f'%Test_Error_w2v)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))
```

```
Training_Accuracy=0.883
Training_Error=0.117
Test_Accuracy=0.866
Test_Error=0.134
```

[5.3.5]Confusion Matrix

In [213]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, clf.predict(sent_vectors_train))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=n
p.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[213]: <matplotlib.axes._subplots.AxesSubplot at 0x23d0d1a4748>

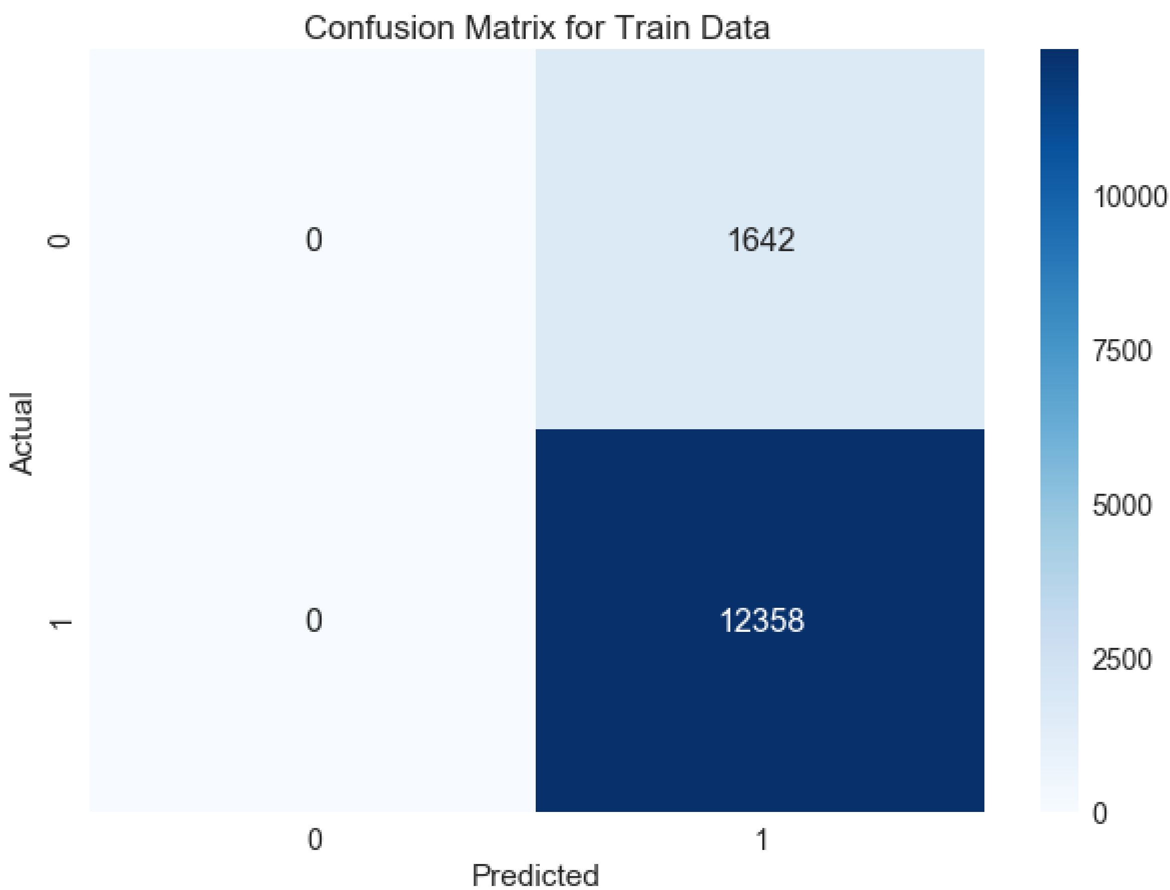

In [214]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, clf.predict(sent_vectors_test))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[214]: <matplotlib.axes._subplots.AxesSubplot at 0x23d2aea12e8>

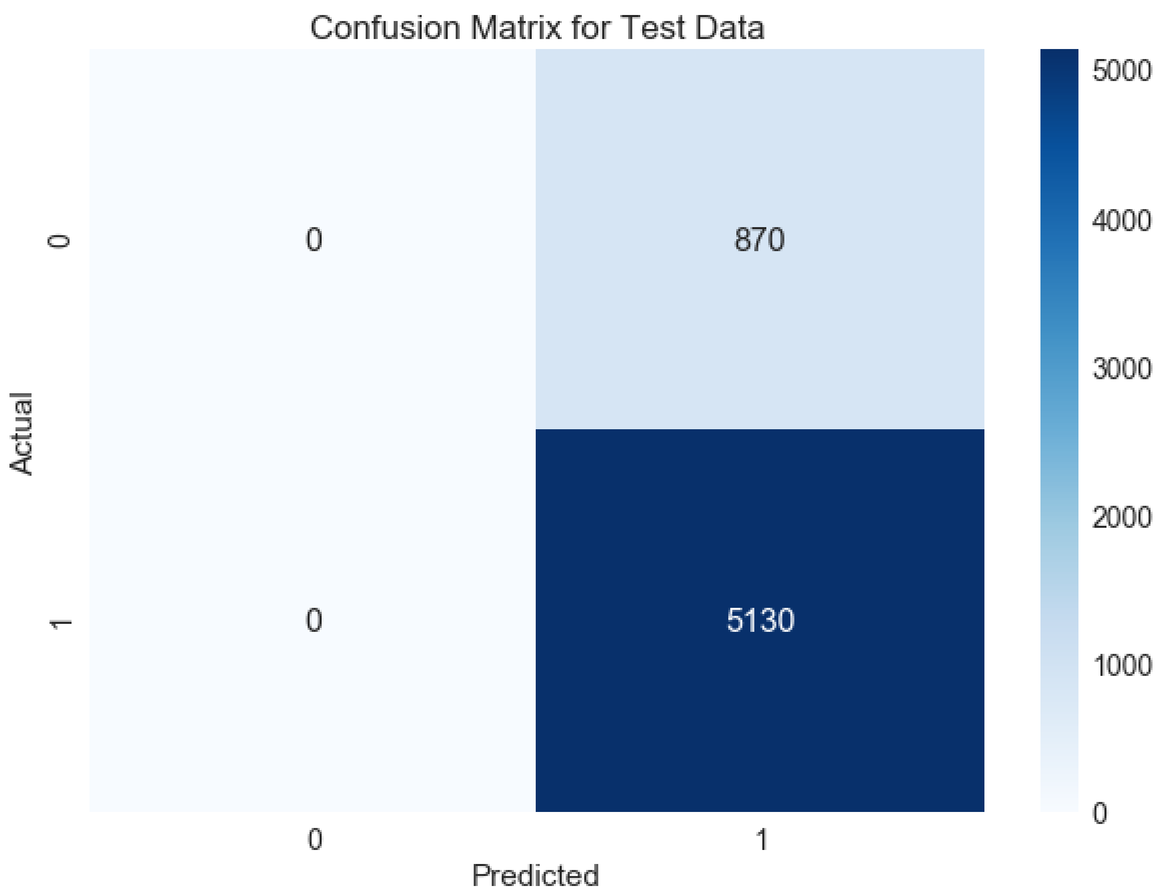

[5.3.6] Classification Report

In [215]:
```
from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

```
             precision    recall  f1-score   support

          0       0.68      0.14      0.23       870
          1       0.87      0.99      0.93      5130

avg / total       0.84      0.87      0.83      6000
```

[5.3.7] Feature Importance

Feature Importance for Positive and Negetive Class

In [231]:
```
clf = SVC(C=1, class_weight='balanced', kernel='linear')
clf.fit(sent_vectors_train, Y_Train)
plot_coefficients(clf, count_vect.get_feature_names())
```

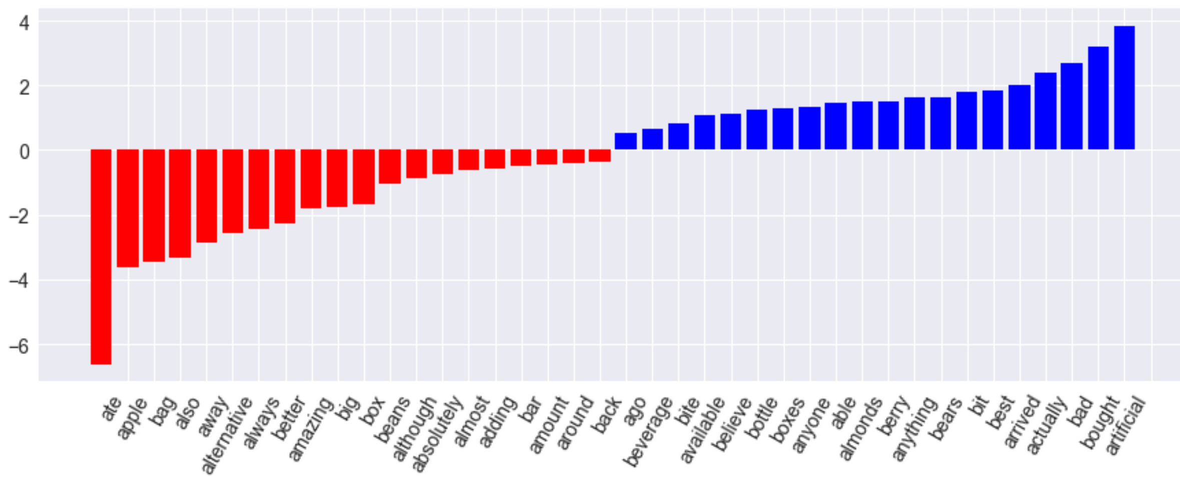

 [5.4] TFIDF weighted W2v

[5.4.1] Hyperameter tuning with L1 Regulizer and AUC Plot

In [251]:
```
Optimal_Lamda_L1(tfidf_sent_vectors_train, Y_Train, tfidf_sent_vectors_cv,Y_cv
)
```

```
100%|████████████████████████████████████████
██████| 9/9 [23:32<00:00, 301.39s/it]
```

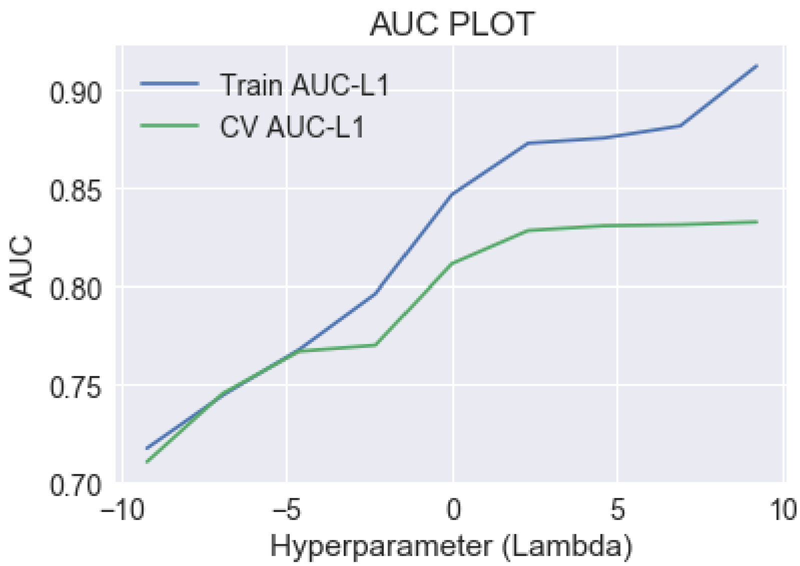


```
CV AUS Scores with Penalty=? Cv auc scores with penalty L1
[0.7105182496471063, 0.7455152248437185, 0.7668969998431583, 0.77004369160449
02, 0.8114655732562569, 0.8283070063408842, 0.8307577692635164, 0.83130234590
3072, 0.832652976630593]
Maximun AUC value : 0.832652976630593
Index 8
```

[5.4.2] ROC Curve of SVM

In [222]:
```
#Testing with test data
clf = SVC(C=1, probability=True)
clf.fit(tfidf_sent_vectors_train,Y_Train)
prediction = clf.predict_proba(tfidf_sent_vectors_test)[:,1]
print(prediction)
print(clf)

Train_FPR, Train_TPR, threshold = roc_curve(Y_Train, clf.predict_proba(tfidf_s
ent_vectors_train)[:,1])
Test_FPR, Test_TPR, threshold = roc_curve(Y_Test, clf.predict_proba(tfidf_sent
_vectors_test)[:,1])
roc_auc = auc(Train_FPR, Train_TPR)
roc_auc1 = auc(Test_FPR, Test_TPR)

plt.plot(Train_FPR, Train_TPR, label = 'Train AUC = %0.2f' % roc_auc)
plt.plot(Test_FPR, Test_TPR, label = 'Test AUC = %0.2f' % roc_auc1)
plt.legend()
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve of MNB')
plt.show()
```

```
[0.80022965 0.44810087 0.95012909 ... 0.90640622 0.93275641 0.92449651]
SVC(C=1, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
  max_iter=-1, probability=True, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

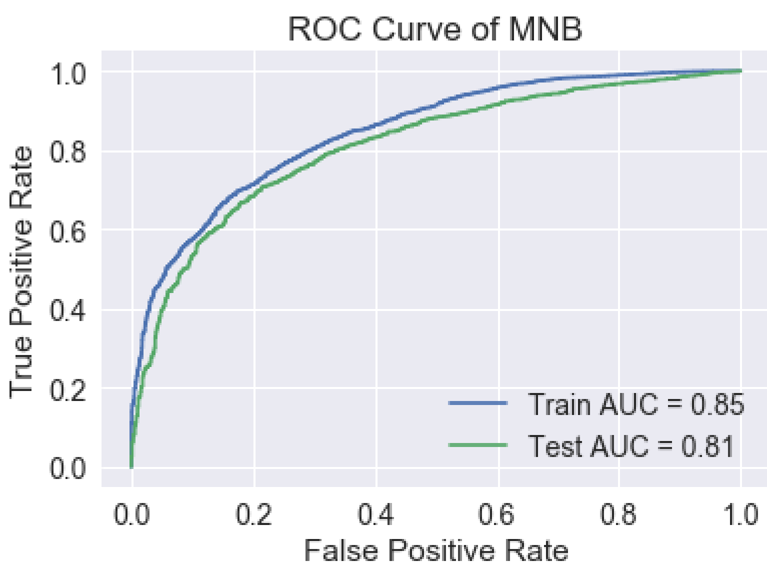

[5.4.3]Train and Test Accuracy

In [223]:
```
Training_Accuracy_tfidfw2v = clf.score(tfidf_sent_vectors_train, Y_Train)
print('Training_Accuracy=%0.3f'%Training_Accuracy_tfidfw2v)
Training_Error_tfidfw2v = 1 - Training_Accuracy_tfidfw2v
print('Training_Error=%0.3f'%Training_Error_tfidfw2v)

Test_Accuracy_tfidfw2v = accuracy_score(Y_Test, prediction.round())
print('Test_Accuracy=%0.3f'%Test_Accuracy_tfidfw2v)
Test_Error_tfidfw2v = 1 - Test_Accuracy_tfidfw2v
print('Test_Error=%0.3f'%Test_Error_tfidfw2v)
#print('\nThe accuracy of the MNB classifier for k = %d is %f%%' % (optimal_al
pha_bow, Test_Accuracy_Bow))
```

```
Training_Accuracy=0.883
Training_Error=0.117
Test_Accuracy=0.856
Test_Error=0.144
```

[5.4.4]Confusion Matrix

In [224]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Train, clf.predict(tfidf_sent_vectors_train))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Train), index=n
p.unique(Y_Train))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Train Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[224]: <matplotlib.axes._subplots.AxesSubplot at 0x23d2c91be48>

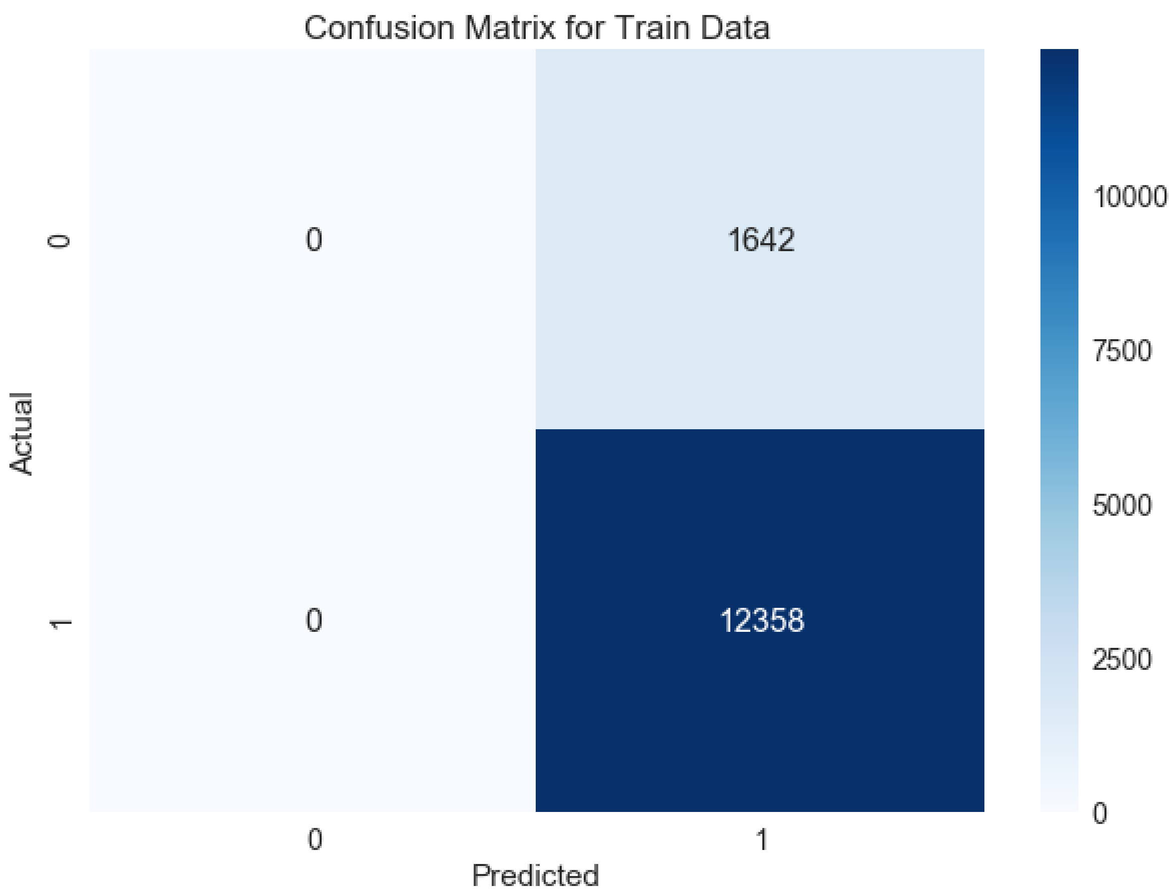

In [225]:
```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_Test, clf.predict(tfidf_sent_vectors_test))
df_conf_matrix = pd.DataFrame(conf_matrix, columns=np.unique(Y_Test), index=np
.unique(Y_Test))
df_conf_matrix.index.name = 'Actual'
df_conf_matrix.columns.name = 'Predicted'
plt.figure(figsize=(10,7))
plt.title("Confusion Matrix for Test Data")
sns.set(font_scale=1.4)
sns.heatmap(df_conf_matrix, cmap='Blues', annot=True, annot_kws={'size':16}, f
mt='d')
```

Out[225]: <matplotlib.axes._subplots.AxesSubplot at 0x23d18e02080>

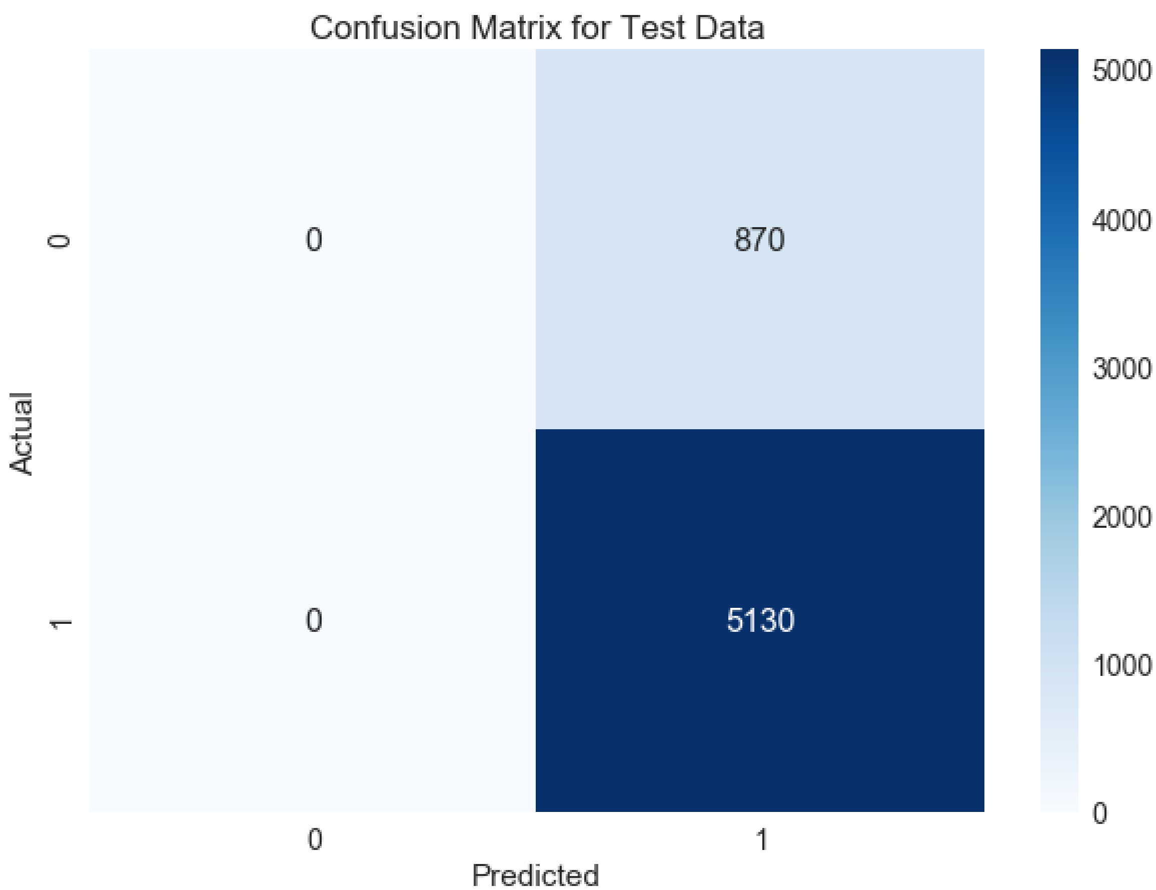

[5.4.5] Classification Report

In [226]:
```
from sklearn.metrics import classification_report
print(classification_report(Y_Test, prediction.round()))
```

```
             precision    recall  f1-score   support

          0       0.54      0.06      0.11       870
          1       0.86      0.99      0.92      5130

avg / total       0.82      0.86      0.80      6000
```

[5.4.6] Feature Importance

Feature Importance for Positive and Negetive Class

In [232]:
```
clf = SVC(C=1, class_weight='balanced', kernel='linear')
clf.fit(tfidf_sent_vectors_train, Y_Train)
plot_coefficients(clf, model.get_feature_names())
```

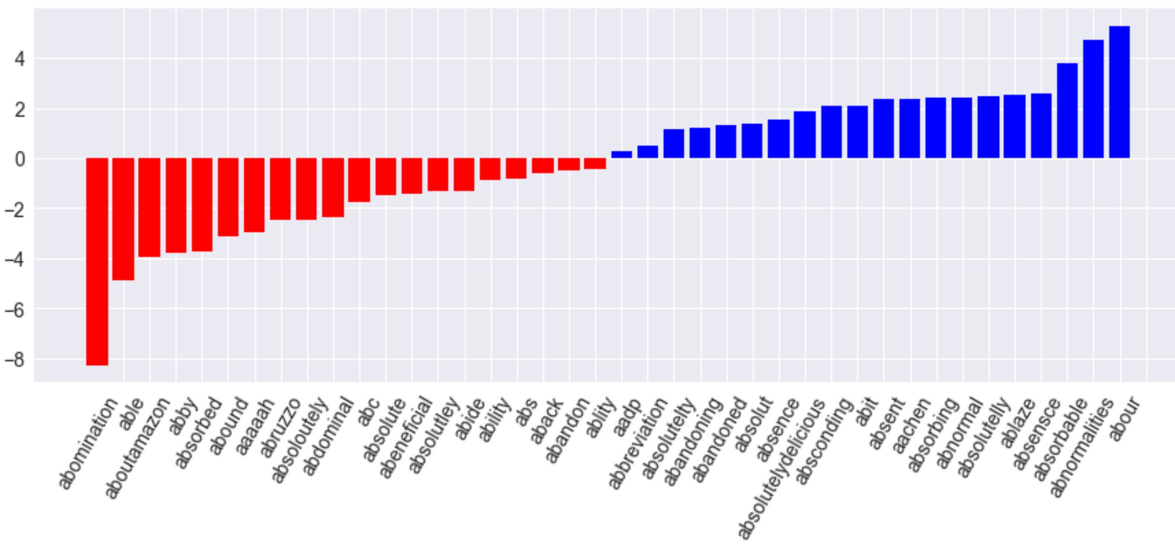

Pretty Table

In [258]:
```
from prettytable import PrettyTable
comparision = PrettyTable()
comparision.field_names = ["Vectorizer", "Kernel", "AUC - L1","AUC-L2", "Train
ing Error", "Test Error"]
comparision.add_row(["BOW","Linear",0.83,0.93,0.02,0.12])
comparision.add_row(["TF-IDF","Linear",0.904,0.95,0.01,0.09])
comparision.add_row(["Avg W2V","Linear",0.903,0.904,0.118,0.116])
comparision.add_row(["TF-IDFWeighted W2V","Linear",0.88,0.88,0.130,0.129])
comparision.add_row(["BoW","RBF","-",0.87,0.060,0.888])
comparision.add_row(["TF-IDF","RBF","-" ,0.89,0.048,0.888])
comparision.add_row(["Avg W2V","RBF","-" ,0.85,0.117,0.134])
comparision.add_row(["TF-IDFWeighted W2V","RBF","-" ,0.83,0.117,0.144])
print(comparision)
```

```
+--------------------+--------+----------+--------+----------------+---------
---+
|     Vectorizer     | Kernel | AUC - L1 | AUC-L2 | Training Error | Test Err
or |
+--------------------+--------+----------+--------+----------------+---------
---+
|        BOW         | Linear |   0.83   |  0.93  |      0.02      |    0.12
|
|       TF-IDF       | Linear |  0.904   |  0.95  |      0.01      |    0.09
|
|      Avg W2V       | Linear |  0.903   | 0.904  |     0.118      |   0.116
|
| TF-IDFWeighted W2V | Linear |   0.88   |  0.88  |      0.13      |   0.129
|
|        BoW         |  RBF   |    -     |  0.87  |      0.06      |   0.888
|
|       TF-IDF       |  RBF   |    -     |  0.89  |     0.048      |   0.888
|
|      Avg W2V       |  RBF   |    -     |  0.85  |     0.117      |   0.134
|
| TF-IDFWeighted W2V |  RBF   |    -     |  0.83  |     0.117      |   0.144
|
+--------------------+--------+----------+--------+----------------+---------
---+
```

Conclusion

1. Applied SVM on all the 4 vectorizers(BOW, TFIDF, AVG-W2V, TFIDF-AVG_W2V).
2. Sorted the data based on Time and Considered 100 K data points for Training set 70K, Test set: 30K. Worked on 2 version of SVM 1-Linear Kernel and 2-RBF Kernel
3. While working with linear kernel used 'SGDClassifier' with 'hinge loss' and used CalibratedClassifierCV. and Class_Weight has set to 'balanced'
4. While working with RBF Kernel, set min_df=10 and max_features=500 and considered 20K data points.
5. Used AUC as a metric for hyperparameter tuning. And took the range of lambda values between (10^-4 to 10^4).
6. Found the top 20 features of positive and negative class for the featurizations Bow and TF-IDF, AvgW2V and TFIDF Weighted vector using Linear Kernel.
7. Found the top 20 features of positive and negative class for the featurizations AvgW2V and TF-IDFWeighted W2V using SVC Kernel. Got error for BoW and TF-IDF.
8. With reference to the pretty table, here is my understanding: a. Linear kernel SVM by using TF-IDF featurization is having the best AUC score: 0.95.
9. RBF Kernel SVM by using TF-IDF having best AUC score: 0.89.
10. Plotted ROC Curve and Confusion Matrix for train and test data for each vectorizer.